

Exploiting Structural Synergy in post-SIRT imaging: Anatomically Guided Joint Variational Reconstruction of Yttrium-90 PET/SPECT/CT

Sam Porter

A dissertation submitted in partial fulfillment
of the requirements for the degree of
Doctor of Philosophy
of
University College London.

The Institute of Nuclear Medicine, Division of Medicine
University College London

June 14, 2026

I, Sam Porter, confirm that the work presented in this thesis is my own. Where information has been derived from other sources, I confirm that this has been indicated in the work.

Abstract

Selective Internal Radiotherapy (SIRT) uses Yttrium-90 (^{90}Y) microspheres to treat unresectable liver tumours. Accurate post-treatment imaging is required to verify microsphere distribution and perform dosimetric evaluations. However, imaging with ^{90}Y presents unique challenges. Bremsstrahlung SPECT offers high count statistics but suffers from poor spatial resolution and difficulties in correcting for scatter and other additive terms. ^{90}Y PET offers superior resolution, but suffers from high noise due to the extremely low internal pair branching ratio ($\sim 0.003\%$). Traditional reconstruction fails to exploit complementary information in these paired datasets.

This thesis proposes a synergistic framework that uses anatomically guided joint variational regularisation to exploit shared underlying structure. We introduce the Directional Total Nuclear Variation Prior, which promotes edge alignment between PET and SPECT without enforcing intensity correlation, while incorporating CT edge-guidance and suppressing unwanted noise.

Further to this, we introduce a simple method of bremsstrahlung range modelling approach and residual correction routine to improve quantification.

To ensure computational efficiency, we propose a synergy-aware preconditioner which, in conjunction with a variance reduced stochastic gradient algorithms, results in fast convergence.

We validate our method against the state-of-the-art sequential hybrid kernel expectation maximisation scheme (HKEM). We demonstrate that dTNV shows improved quantitative accuracy and increased reproducibility over noise realisations compared to existing methods.

Impact Statement

UCL theses now have to include an impact statement. (*I think for REF reasons?*)

The following text is the description from the guide linked from the formatting and submission website of what that involves. (Link to the guide: <http://www.grad.ucl.ac.uk/essinfo/docs/Impact-Statement-Guidance-Notes-for-Research-Students-and-Supervisors.pdf>)

The statement should describe, in no more than 500 words, how the expertise, knowledge, analysis, discovery or insight presented in your thesis could be put to a beneficial use. Consider benefits both inside and outside academia and the ways in which these benefits could be brought about.

The benefits inside academia could be to the discipline and future scholarship, research methods or methodology, the curriculum; they might be within your research area and potentially within other research areas.

The benefits outside academia could occur to commercial activity, social enterprise, professional practice, clinical use, public health, public policy design, public service delivery, laws, public discourse, culture, the quality of the environment or quality of life.

The impact could occur locally, regionally, nationally or internationally, to individuals, communities or organisations and could be immediate or occur incrementally, in the context of a broader field of research, over many years, decades or longer.

Impact could be brought about through disseminating outputs (either in scholarly journals or elsewhere such as specialist or mainstream media),

education, public engagement, translational research, commercial and social enterprise activity, engaging with public policy makers and public service delivery practitioners, influencing ministers, collaborating with academics and non-academics etc.

Further information including a searchable list of hundreds of examples of UCL impact outside of academia please see <https://www.ucl.ac.uk/impact/>. For thousands more examples, please see <http://results.ref.ac.uk/Results/SelectUoa>.

Acknowledgements

So, it turns out that a PhD is quite an undertaking. It's been a long old journey, but the end is in sight!

I've had an extraordinary four years and would like to thank some extraordinary people. Firstly, I'd like to thank my supervisors, Kris, Daniel, and Simon, who have been a near-bottomless source of knowledge, inspiration and ideas. Secondly, my colleagues in the INM, especially my fellow students. This would have been a very difficult (and much lonelier) few years without you all. I'm so grateful for our conversations in and out of the workplace. I'd like to also thank all of my colleagues in the scientific and medical spheres, whom I've found to be endlessly generous with their time and patience.

I'd like to thank my parents and my sister, Alice, for all your support over the past four years and the many years before. There's no way that any of this would have happened without you all. Thank you.

Most of all, I want to express my deep gratitude to my newly minted wife, Kate. Sorry for all the late evenings and weekends spent working. Thank you so much for going through this with me. I love you.

Contents

1	Introduction	20
	Introduction	20
2	Physics and Measurement Models in Emission Tomography	23
2.1	Principles of Emission Tomography	24
2.2	Positron Emission Tomography (positron emission tomography (PET))	24
2.2.1	Coincidence Detection and Data Acquisition	25
2.2.2	The PET model	27
2.3	Single Photon Emission Computed Tomography (single photon emission computed tomography (SPECT))	28
2.3.1	The SPECT model	29
2.4	Yttrium-90 Radioembolisation (selective internal radiotherapy (SIRT))	30
2.4.1	Decay Characteristics	30
2.4.2	Bremsstrahlung SPECT	30
2.4.3	Internal Pair Production PET	30
2.5	X-ray computed tomography (CT)	31
3	Penalised-Likelihood Reconstruction: Priors and Optimisation Algorithms	34
3.1	The Inverse Problem in Emission Tomography	35
3.2	Regularisation Strategies	37
3.2.1	Image-based Regularisation	37
3.2.2	Anatomically Guided Regularisation	39

3.2.3	Synergistic Regularisation	43
3.3	Optimisation Algorithms	46
3.4	Summary and Research Gap	54
4	Improvements to the ^{90}Y Bremsstrahlung System Model	57
4.1	Challenges in Bremsstrahlung Imaging	57
4.2	Optimisation Strategies and Monte Carlo Modeling	58
4.3	Additive Corrections	59
4.3.1	Statistical model with an additive background term	59
4.3.2	Initial scatter estimate via a pre-trained deep network	60
4.3.3	Iterative OSEM–SIMIND refinement loop	60
4.3.4	Offline versus online additive correction strategies	61
4.3.5	Full replacement versus damped scatter updates	61
4.4	Bremsstrahlung Range Blur	63
4.4.1	Projector point-spread function	63
4.4.2	Bremsstrahlung-specific blur	65
4.4.3	Validation against SIMIND and GATE forward projections	66
4.5	Residual Correction	67
4.5.1	Residual correction update	67
4.5.2	Regularised objective with a Relative Difference prior	68
4.5.3	Interpretation via matched versus unmatched projectors	69
4.5.4	Empirical observations	71
4.6	Putting It All Together	71
5	Anatomy-informed deconvolution for PET partial volume correction	72
5.1	Introduction and relationship to the thesis	72
5.2	Image-domain inverse problem and objective	73
5.3	Baseline: Richardson–Lucy deconvolution	74
5.3.1	Scaled-gradient interpretation and EM-type preconditioning	75
5.4	Kernelised Richardson–Lucy deconvolution	75
5.4.1	Kernel image model	75

5.4.2	KRL update	76
5.4.3	Hybrid KRL (HKRL)	76
5.5	Directional-TV regularised RL deconvolution	76
5.5.1	Penalised objective and update	77
5.6	Data and experimental setup	78
5.6.1	Hyperparameter tuning	80
5.7	Results	81
5.7.1	Spheres	81
5.7.2	τ -PET brain	82
5.8	Discussion	83
5.9	Summary	85
6	Introduction to Weighted Directional Total Nuclear Variation	86
6.1	Total Nuclear Variation Revisited	86
6.1.1	Mathematical properties and differentiability	87
6.2	Weighted Directional Total Nuclear Variation (w -dT _{NV})	90
6.2.1	Directional operators from guidance	91
6.2.2	Modality-wise weighting and the full w -dT _{NV} prior	92
6.3	Conclusion	93
7	Preconditioning and Subset Selection for Synergistic Reconstruction	94
7.1	Stochastic MAP Optimisation	94
7.2	Prior-Aware Preconditioning	95
7.2.1	Data Curvature	96
7.2.2	Current w -dT _{NV} Prior Curvature	96
7.2.3	Combining Data and Prior Curvature	98
7.2.4	Preconditioner Comparison Methods	98
7.3	Subset Selection	99
7.3.1	Data-Subset Organisation	99
7.3.2	Prior Sampling Strategy	100
7.3.3	Subset-Selection Experiment Methods	101

7.4	Multi-Bed PET Acquisition	101
7.4.1	Notation	101
7.4.2	Why Not Stitch Independent Reconstructions?	102
7.4.3	Use in the Subset Experiments	103
7.5	Conclusion	103
8	Weighted Directional Total Nuclear Variation for Joint ^{90}Y PET/SPECT Reconstruction with CTAC-derived Guidance	104
8.1	Motivation and clinical context	105
8.2	Joint MAP formulation for PET/SPECT	105
8.3	w - d TNV instantiated for ^{90}Y PET/SPECT with CTAC guidance . .	106
8.3.1	Directional guidance from the CTAC-derived μ -map	106
8.3.2	Global modality normalisation as a constant diagonal weight- ing	107
8.3.3	Directional finite-difference operator for PET/SPECT coupling	108
8.3.4	The w - d TNV Jacobian and regulariser (two modalities) . . .	108
8.3.5	Ablation comparators	109
8.4	Optimisation and implementation	110
8.4.1	Diagonal preconditioning	111
8.4.2	Comparator: sequential hybrid kernel EM	111
8.4.3	Software implementation	112
8.5	Forward models, corrections, and co-registration	112
8.5.1	PET forward model	112
8.5.2	SPECT forward model	113
8.5.3	Co-registration and warp operators	114
8.6	Data and evaluation protocol	114
8.6.1	Phantom data	115
8.6.2	Clinical data	117
8.7	Results	119
8.7.1	Hyperparameter selection	119
8.7.2	Phantom results	119

8.7.3 Clinical results	122
8.8 Discussion	124
8.9 Conclusion	128
Appendices	129
A simind-python-connector	129
A.1 Purpose and Scope	129
A.2 Connector Workflow	130
A.3 Geometry and Data Conventions	131
A.4 Validation and Reproducibility	132
B Colophon	133
References	134

List of Figures

- 2.1 A positron traverses tissue, losing kinetic energy, before annihilating with an electron, resulting in two (nearly) anti-parallel photons. . . . 26
- 2.2 Simple diagram of a PET scanner. Left, axial view showing detectors arranged in a ring. Right, sagittal/coronal view showing multiple rings forming the axial FoV. 26
- 2.3 Possible coincidences under PET electronic collimation. From left to right: (a) scatter (green dashed line shows effective line of response (LoR)), true; (b) random (green dashed line shows effective LoR). . 27
- 2.4 Left, non-parallel photon rejected by physical collimation. Right, photon parallel to the collimator can be detected, giving a line of response containing the emission source. 28
- 2.5 Representative diagram of bremsstrahlung production in tissue for an emitted β^- particle. 31
- 2.6 Bremsstrahlung spectra: (a) emitted spectrum from Geant4 Application for Tomographic Emission (GATE) FastY90 model [18, 19] and (b) detected (after attenuation by detector cover and backscattering) using AnyScan Trio GATE scanner model [20]. 32
- 2.7 Tangential profile of measured prompt counts and estimated additive terms (randoms+scatter), with the corresponding corrected prompts for an Yttrium-90 (^{90}Y) PET scan. The data are summed over all views and segments, leaving counts per tangential bin. 33

- 3.1 Visualisation of the Frobenius and nuclear norms (vertical axis) of a matrix X as functions of its (two) singular values (horizontal axis). The Frobenius norm is $\|X\|_F = \sqrt{\sigma_1(X)^2 + \sigma_2(X)^2}$ (red mesh), whereas the nuclear norm is $\|X\|_* = |\sigma_1(X)| + |\sigma_2(X)|$ (black mesh). For visualisation purposes we extend the singular-value axes to negative values. The nuclear norm exhibits sharp grooves along the coordinate axes, encouraging one singular value to vanish (i.e. promoting low-rank structure), while the Frobenius norm is smooth and does not preferentially drive a singular value exactly to zero. This behaviour can also be understood via subgradients: the nuclear-norm subgradient (cyan) points towards rank-1 grooves, preferentially shrinking the smaller singular value, whereas the Frobenius-norm subgradient (red) points towards the origin, shrinking both singular values equally. 47
- 3.2 Demonstration of limit cycle behaviour with an ordered subset gradient descent algorithm. Top row: 3 different simple sub-functions representing subset likelihoods. Bottom row: left, full function with sub-gradient steps; centre, zoomed version of the full function demonstrating limit cycle behaviour (green dots show approximate points of limit cycle); right, likelihood value per iteration, demonstrating limit cycle precession. 50
- 3.3 Demonstration of reduction in limit cycle behaviour with step size relaxation $\varsigma_k = \varsigma_0 / (1 + 0.05k)$. Left, full function with sub-gradient steps; centre, zoomed version of the full function demonstrating relaxed limit cycle behaviour; right, likelihood value per iteration, demonstrating relaxed limit cycle precession. 51

3.4	Demonstration of removal of limit cycle behaviour with an stochastic variance-reduced gradient (SVRG)-type (Ordered VRG) algorithm. Left, full function with sub-gradient steps; centre, zoomed version of the full function demonstrating removal of limit cycle behaviour; right, likelihood value per iteration, demonstrating removal of limit cycle precession	52
4.1	Plots showing probability of bremsstrahlung gamma emission versus the distance from ^{90}Y decay site. Data from Rault <i>et al.</i> [18]	58
4.2	Figures showing profile through projections for NEMA phantom. Summed over projections, axial slice is through NEMA spheres. Horizontal axis show tangential coordinate on gamma camera, Vertical axis shows counts. Top, replacing additive data; bottom, damped addition of additive data ($\alpha = 0.2$).	62
4.3	Example fits of an isotropic Gaussian proxy to the Monte Carlo source-to-emission displacement distribution for ^{90}Y Bremsstrahlung in water (clinical energy windows).	63
5.1	Left, ground truth phantom, centre blurred and noisy reconstruction and right MRI guidance image for the simulated sphere phantom. Units are arbitrary.	78
5.2	Left, reconstructed τ -PET brain, right registered MRI guidance image. Units are arbitrary.	79
5.3	Figures showing central slice of reconstructions with best NRMSE. Top, initialised by the blurred PET image and bottom, initialised by a uniform image.	81
5.4	Figures showing iteration versus global NRMSE for the best performing reconstructions. Top, initialised by the blurred PET image and bottom, initialised by a uniform image.	83

- 5.5 Figures showing RC versus CoV for all reconstructions. Top, initialised by the blurred PET image and bottom, initialised by a uniform image. KRL and HKRL are visually indistinguishable in these plots. 84
- 5.6 Figures showing visually best reconstructions for the different reconstructions. 84
- 8.1 NEMA phantom analysis regions. Left: warm-background (BG), lung insert (LI), and high-count (HC) VOIs. Right: CTAC-derived attenuation map showing sphere edges, the lung insert, and the under-filled sphere excluded from the HC VOI. 115
- 8.2 Training-patient VOIs. The five identified lesions and the background reference region are shown on the vendor PET reconstruction with registered attenuation maps. Top row: lesions 1–3. Bottom row: lesions 4–5 and background. 118
- 8.3 Training-patient hyperparameter sweep. Impact of varying PET weighting while holding SPECT regularisation constant (top row) and varying SPECT weighting while holding PET regularisation constant (bottom row) on TBR and CoV_{bkg} for lesions 1 and 2. The first and third columns show w -dTNV; the second and fourth show w -TNV. 120
- 8.4 Visual comparison of NEMA phantom reconstructions for a randomly chosen bootstrap. From left to right: w -dTNV, dTV, w -TNV, SHKEM (subiteration 27). PET reconstructions are shown on the top row and SPECT on the bottom row (SPECT guidance for SHKEM). Images are displayed in reconstruction intensity units, proportional to activity concentration. 120

- 8.5 Recovery coefficient (RC) for NEMA phantom spheres for w -dTNV, dTV, w -TNV and SHKEM (subiteration 27). Also shown are the equivalent RCs for a high-count, long-OSEM reconstruction. The grey line indicates perfect recovery. Statistical significance is reported for paired comparisons versus w -dTNV (Benjamini–Hochberg adjusted): * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$ 121
- 8.6 Mean voxel-wise coefficient of variation (CoV) across 20 bootstrapped PET realisations within each VOI. Statistical significance is reported for paired comparisons versus w -dTNV (BH-adjusted): * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$ 122
- 8.7 Visual comparison of reconstruction methods for training patient SIRT3. Columns show w -dTNV, dTV, w -TNV and SHKEM (subiteration 27). The upper two quadrants show PET reconstructions; the lower two show SPECT reconstructions or guidance images. Images are displayed in reconstruction intensity units, proportional to activity concentration. 123
- 8.8 Clinical cohort summary: lesion-wise mean TBR (left) and patient-wise background CoV (right). Statistical significance is reported for comparisons versus w -dTNV (BH-adjusted): * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$ 123
- 8.9 Training-patient CTAC-derived guidance. The CTAC-derived attenuation map μ (top row; cm^{-1}) and magnitude of the w -dTNV directional weighting field $\|\xi\|$ (bottom row) are shown for lesions 1 and 2. CTAC streak artefacts are visible in μ and are reproduced in $\|\xi\|$, providing a pathway for artefact transfer into the emission reconstruction. 126

- 8.10 NEMA phantom, central axial slice. CTAC-derived attenuation map μ (left; cm^{-1}) and magnitude of the w -dTNV directional weighting field $\|\xi\|$ (right). With a high-quality CTAC free of prominent streak artefacts, $\|\xi\|$ concentrates around true object boundaries and insert edges, enabling stable μ -map-guided directional smoothing. . . . 127

List of Tables

5.1	Hyperparameter sweeps for kernel-based deconvolution	81
5.2	Tables summarising the hyperparameters and metric values of the best (measured by NRMSE) reconstructions of the various methods for (top) a uniform initialisation of ones and (bottom) the blurred PET image.	82
7.1	Preconditioner settings used in the current preconditioner experiments.	98
8.1	Selected regularisation weights (tuned on a single training patient) used for all phantom and clinical experiments.	119
8.2	Effect ratios $\exp(\Delta_{r,s})$ (comparator vs. baseline w -dTNN) with 95% confidence intervals and BH-adjusted p -values, reported by sphere diameter s . Values below one indicate lower recovery coefficient than w -dTNN.	121
8.3	Effect ratios $\exp(\Delta)$ with 95% confidence intervals and BH-adjusted p -values for mean tumour-to-background ratio (TBR), comparing each method to w -dTNN. Values below one indicate lower TBR than w -dTNN.	124
8.4	Effect ratios $\exp(\Delta)$ with 95% confidence intervals and BH-adjusted p -values for background CoV, comparing each method to w -dTNN. Values above one indicate higher background CoV than w -dTNN. .	124
A.1	Main <code>simind-python-connector</code> interfaces relevant to thesis workflows.	131

Chapter 1

Introduction

Yttrium-90 (^{90}Y) selective internal radiotherapy (SIRT), also referred to as transarterial radioembolisation (TARE), is an established treatment for primary and secondary liver malignancies [1]. In this procedure, ^{90}Y -labelled microspheres are administered into the hepatic artery to preferentially irradiate tumour tissue while sparing healthy liver parenchyma. The administered activities are high (typically of the order 2 GBq), and clinically relevant absorbed-dose heterogeneity can arise at multiple spatial scales due to vascular delivery, microsphere clustering, and patient-specific hepatic anatomy [2]. Tumour doses of clinical interest often exceed 200 Gy [3], and recent evidence has shown potential benefit from additional subsequent treatment, which could result in more dose to healthy liver tissue without proper treatment planning and verification [4]. As a consequence, post-therapy quantitative imaging and dosimetry are central to treatment verification, outcome modelling, and the optimisation of any subsequent therapy.

Despite its importance, routine post-SIRT dosimetry is not universally performed in clinical practice, primarily due to practical workflow constraints and limitations in quantitative accuracy [5, 6]. Post-therapy imaging of ^{90}Y is inherently challenging: image quality is often degraded by low effective count statistics and by physics that is less favourable than standard diagnostic radionuclide imaging, leading to increased noise and modelling error [5]. More broadly, emission tomography is vulnerable to substantial degradation from Poisson counting statistics as well as inaccuracies in system modelling and hardware limitations [7]. Improving

reconstruction quality in this low-count, model-mismatch regime is therefore a key enabling step for accurate absorbed-dose estimation and for robust patient-specific outcome prediction [8].

In post-SIRT imaging, the two clinically relevant emission modalities are complementary. ^{90}Y positron emission tomography (PET) offers comparatively high spatial resolution, but is severely count-limited because of the extremely low internal pair-production branch. Bremsstrahlung single photon emission computed tomography (SPECT), by contrast, provides higher count statistics but suffers from poor spatial resolution and more severe modelling challenges, particularly for scatter, septal penetration, and additive corrections [5]. In contemporary workflows, these emission images are routinely paired with computed tomography (CT), which provides attenuation correction and anatomical context. This naturally motivates reconstruction approaches in which information is shared across modalities, with CT contributing structural guidance, and PET and SPECT contributing complementary functional information. Synergistic methods, where such information is incorporated during reconstruction rather than only after image formation, have gained prominence in related hybrid imaging contexts [9]. Within post- ^{90}Y SIRT imaging, these approaches have the potential to improve lesion-level quantification and absorbed-dose estimation without increasing acquisition time or administered activity.

The central aim of this thesis is therefore to develop and evaluate a synergistic reconstruction framework tailored to post- ^{90}Y SIRT, with a particular focus on improving the reliability of quantitative emission imaging and therefore downstream dosimetry under clinically realistic count levels and modelling limitations. To this end, the thesis develops practical improvements to the bremsstrahlung SPECT system model, introduces a weighted directional total nuclear variation prior for joint PET/SPECT reconstruction with CT-derived directional guidance, and investigates optimisation and preconditioning strategies for this coupled penalised-likelihood problem.

The thesis is organised as follows:

- **Chapter 1** introduces the PET/SPECT/CT imaging problem, outlining the fundamental principles of each modality before describing the specific challenges of post- ^{90}Y imaging.
- **Chapter 2** describes the inverse problem in emission tomographic reconstruction and reviews regularisation strategies and efficient optimisation methods.
- **Chapter 3** addresses the specific challenges of post-SIRT bremsstrahlung imaging and develops practical improvements to the reconstruction system model.
- **Chapter 4** studies anatomy-informed regularisation and optimisation in a controlled post-reconstruction deconvolution setting.
- **Chapter 5** introduces the weighted directional total nuclear variation prior and its differentiable formulation.
- **Chapter 6** develops the optimisation machinery used with weighted directional total nuclear variation (w -dT_{NTV}), including prior-aware preconditioning, subset selection, and multi-bed PET handling.
- **Chapter 7** evaluates w -dT_{NTV} for joint ^{90}Y PET/SPECT reconstruction, compares it with the current state-of-the-art method sequential hybrid kernelised expectation maximisation (SHKEM), and discusses the main limitations and future directions.
- **The Appendix** documents additional software and reproducibility material that supports the thesis workflow.

Chapter 2

Physics and Measurement Models in Emission Tomography

This chapter introduces the physical principles and measurement models underpinning emission tomography, with emphasis on the effects that dominate reconstruction quality and quantitative reliability. In both positron emission tomography (PET) and single photon emission computed tomography (SPECT), the reconstruction task is to infer a non-negative radiotracer activity distribution from projection data generated by stochastic radioactive decay and corrupted by attenuation, scatter, and blurring effects. These are naturally captured by a forward model that maps the image to expected measurement counts, together with additive components representing background contributions.

We first summarise the detection physics of PET and SPECT and present convenient factorisations of the system matrix that separate geometry, attenuation, and detector effects. We then focus on the particularly challenging case of post-SIRT ^{90}Y imaging, where quantification is limited by extreme count starvation in internal pair production (IPP) PET and by spectral and resolution degradations in bremsstrahlung SPECT. Finally, we describe the role of computed tomography (CT) in this thesis as the source of attenuation information required for physically accurate emission modelling and a high-resolution anatomical reference for reconstruction regularisation in low-count regimes.

2.1 Principles of Emission Tomography

emission tomography (ET) fundamentally differs from functional modalities such as CT. Rather than mapping the structural properties, ET measures the spatio-temporal distribution of a radiolabelled substance introduced into the physiological system. A representative example is ^{18}F -fluoro-2-deoxyglucose (FDG)[10]. FDG is a glucose-analogue, meaning energy-hungry cells take up an increased amount of the molecules, and therefore radionuclide. Cells are unable to metabolise FDG, meaning it accumulates. The distribution of FDG can be imaged using a PET scanner, giving a functional view of glucose metabolism and highlighting any unusual uptake caused by overactive cells such as cancer.

While the instrumentation differs between Single Photon Emission Computed Tomography (SPECT) and PET, both solve a similar inverse problem: estimating the activity concentration from a set of projection measurements that are corrupted by Poisson noise, attenuation, and scatter. In practice, the continuous activity distribution is represented on a finite voxel grid, producing an unknown non-negative image vector, $x \in \mathbb{R}_+^J$. This discretisation allows the acquisition physics to be expressed as a linear operator mapping the image to expected measurement counts, $A : \mathbb{R}_+^J \rightarrow \mathbb{R}_+^I$. The precise discretisation and statistical modelling are introduced later. For now, we present a simple model of the physical acquisition process, including scanner geometry, photon attenuation, and unwanted background contributions. To do this, we use an affine forward model,

$$\bar{y} = Ax + b, \quad (2.1)$$

where $\bar{y} = \mathbb{E}[y]$ denotes the expected measurement counts, A is the system matrix encoding the measurement physics we introduced earlier, and b represents additive contributions not explained by the primary emission model.

2.2 Positron Emission Tomography (PET)

PET uses neutron-deficient isotopes that stabilise via β^+ decay. Within the nucleus, a proton converts into a neutron, emitting a positron (e^+) and an electron neutrino

(ν_e) to conserve lepton number:



The decay energy is shared between the positron and the neutrino, resulting in a continuous positron energy spectrum. This positron traverses the tissue losing kinetic energy via Coulomb interactions (a distance known as the positron range) before thermalising. This range imposes an effective blurring limit on spatial resolution, with an effective positron range of approximately 0.5–0.7 mm for ^{18}F and several millimetres for high-energy emitters such as ^{82}Rb [11].

Upon thermalisation, the positron interacts with an electron to form positronium, which rapidly decays via annihilation:



The rest mass energy of the pair (2×511 keV) is converted into two¹ (near) anti-parallel² photons (Figure 2.1).

2.2.1 Coincidence Detection and Data Acquisition

PET detectors are typically assembled into modular detector blocks, which are then arranged circumferentially to form a ring around the patient. Modern clinical systems comprise multiple such rings stacked along the axial direction, increasing axial coverage and sensitivity by enlarging the axial field of view (FoV). Figure 2.2 illustrates this geometry showing a single ring in the axial view, and multiple rings in the sagittal view.

PET imaging relies on a technique known as electronic collimation. A valid event is recorded only if two detectors register photons within a coincidence timing window 2τ (typically 6–12 ns) [11]. The measured coincidences (or counts) along a LoR, connecting detectors 1 and 2, denoted as $y_{i1,2}$, is a superposition of three

¹Positron/electron pairs can also annihilate to produce more than two photons but this is much less likely ($p(3\gamma)/p(2\gamma) \approx 1/378$) [12]

²In order to conserve momentum, a certain a-collinearity will exist between the photons, introducing additional blurring [13].




Figure 2.1: A positron traverses tissue, losing kinetic energy, before annihilating with an electron, resulting in two (nearly) anti-parallel photons.




Figure 2.2: Simple diagram of a PET scanner. Left, axial view showing detectors arranged in a ring. Right, sagittal/coronal view showing multiple rings forming the axial FoV.

components:

$$y_{i,2} = t_{i,2} + s_{i,2} + r_{i,2} \quad (2.4)$$




Figure 2.3: Possible coincidences under PET electronic collimation. From left to right: (a) scatter (green dashed line shows effective LoR), true; (b) random (green dashed line shows effective LoR).

where t_{ij} represents true coincidences (Figure 2.3a, right), $s_{i,2}$ represents scatter (where one or both photons have undergone at least one Compton interaction; Figure 2.3a, left)), and $r_{i,2}$ represents random coincidences (Figure 2.3b).

Random coincidences occur when two unrelated annihilation events are detected within the timing window. The rate of randoms on a specific LoR is governed by the singles rates of the individual detectors (R_1, R_2), often approximated as:

$$r_{i,2} = 2\tau R_1 R_2. \quad (2.5)$$

2.2.2 The PET model

In PET, the system matrix may be factorised as

$$A = A_{\text{det,sens}} A_{\text{det,blur}} A_{\text{att}} A_{\text{geom}} A_{\text{ToF}} A_{\text{positron}}. \quad (2.6)$$

Here, A_{positron} models positron-range blurring, A_{ToF} models the timing localisation accuracy for time of flight (ToF) PET (and is omitted for non-ToF PET), A_{geom} is the geometric projection operator mapping voxels to lines of response in the absence of attenuation, A_{att} is a diagonal operator containing attenuation factors along each line




Figure 2.4: Left, non-parallel photon rejected by physical collimation. Right, photon parallel to the collimator can be detected, giving a line of response containing the emission source.

of response, $A_{\text{det,blur}}$ models detector-related resolution effects in reporting the true line-of-response positions, and $A_{\text{det,sens}}$ is a diagonal sensitivity factor incorporating multiplicative effects such as detector efficiencies and normalisation. The additive term is commonly written as $b = s + r$, corresponding to scatter and randoms, respectively.

2.3 Single Photon Emission Computed Tomography (SPECT)

SPECT imaging relies on isotropic single gamma emissions. As there is no coincident pair to define the LoR, physical collimation is required. A lead collimator, typically a parallel-hole geometry, allows only photons travelling through the collimator holes. A single (or multiple, most often a dual-head system) collimated gamma camera is rotated around the patient (Figure 2.4).

The sensitivity (g) of a parallel-hole collimator depends on the hole diameter (d), septal thickness (t), and the effective hole length (L_e). A commonly used geometric-efficiency scaling is

$$g \propto \frac{d^4}{L_e^2 (d + t)^2}, \quad (2.7)$$

which makes explicit the trade-off between resolution and sensitivity: reducing d improves geometric resolution but decreases sensitivity. In practice, the geometric efficiency of a (high-resolution) parallel-hole collimator can be very low (often $\lesssim 0.05\%$) [14]. In the context of bremsstrahlung imaging (Section 2.4.2), this low efficiency is compounded by the continuous energy spectrum of the emissions, which limits the effectiveness of standard photopeak-based energy windowing for scatter rejection.

2.3.1 The SPECT model

For SPECT, a convenient factorisation is

$$A = A_{\text{det,sens}} A_{\text{det,blur}} A_{\text{geom,att}} A_{\text{bremss}}, \quad (2.8)$$

where $A_{\text{det,sens}}$ collects multiplicative sensitivity factors (e.g. detector efficiency and calibration/normalisation terms), $A_{\text{det,blur}}$ represents intrinsic resolution effects within the gamma camera, and $A_{\text{geom,att}}$ is the geometric projection operator that encapsulates collimator effects (including depth-dependent resolution) together with depth- and view-dependent attenuation. The operator A_{bremss} is specific to bremsstrahlung SPECT and models the additional blurring associated with bremsstrahlung range (discussed in Section 2.4.2).

For a parallel-hole collimator, septal penetration can be incorporated via an energy-dependent effective hole length,

$$L_e(E) = l - \frac{2}{\mu_{\text{Pb}}(E)}, \quad (2.9)$$

where l is the physical hole length and $\mu_{\text{Pb}}(E)$ is the linear attenuation coefficient of lead at photon energy E . The combined intrinsic detector blur and collimator blur is commonly approximated by a depth-dependent Gaussian point-spread function embedded within $A_{\text{geom,att}}$, with standard deviation $\sigma(z, E)$ increasing with source-to-collimator distance z (and depending on E through penetration).

2.4 Yttrium-90 Radioembolisation (SIRT)

SIRT uses ^{90}Y microspheres for the locoregional treatment of hepatic malignancies. It is an established treatment for unresectable hepatocellular carcinoma (HCC) and liver metastases [1], recommended by NICE for unresectable HCC and used for neuroendocrine and colorectal liver metastases in the UK [15, 16]. The reconstruction of post-SIRT microsphere distribution presents a severe inverse problem, as ^{90}Y is designed for therapy, not imaging.

2.4.1 Decay Characteristics

^{90}Y is an almost pure β^- emitter decaying to the stable ground state of Zirconium-90 (^{90}Zr) with a half-life of 64.1 hours:




The electrons have a high endpoint energy of 2.28 MeV. The primary challenge for quantitative imaging is the absence of primary gamma photons. Imaging must therefore exploit two secondary, low-probability physical interactions.

2.4.2 Bremsstrahlung SPECT

As high-energy β^- particles decelerate in tissue through Coulomb interactions with charged particles, they emit Bremsstrahlung photons. These photons can be imaged with SPECT systems. However, the continuous spectrum of Bremsstrahlung photons (Figure 2.6a) hugely reduces the suitability of traditional energy window scatter correction methods and Bremsstrahlung range blurring (Figure 2.5 reduces resolution. Consequently, Bremsstrahlung SPECT is often characterised by poor spatial resolution and limited quantitative accuracy [17].

2.4.3 Internal Pair Production PET

A minor decay branch (31.86×10^{-6}) allows the nucleus to transition to an excited 0^+ state of ^{90}Zr , which de-excites via internal pair production (IPP):






Figure 2.5: Representative diagram of bremsstrahlung production in tissue for an emitted β^- particle.


The resulting positron annihilates to produce standard 511 keV photon pairs [21]. While this enables quantitative PET imaging, the acquisition is “count-starved” due to the extremely low branching ratio and noise-dominated due to the Bremsstrahlung background (Figure 2.7).

2.5 X-ray CT


In this work, CT is used primarily to provide anatomical context and attenuation information for emission tomography reconstruction. X-ray CT measures the transmitted intensity of a photon beam after passage through an object or patient. In an ideal monoenergetic setting, transmission follows the Beer–Lambert law, so that the log-normalised measurements are proportional to line integrals of the linear attenuation coefficient. In practice, polychromaticity introduces effects such as beam hardening, which are mitigated by scanner calibration and correction steps to yield images that approximate an attenuation map.

Reconstructed CT images are typically reported in Hounsfield units (HU), calibrated such that water is 0 HU and air is -1000 HU. The high spatial resolution of CT provides structurally reliable edge information (e.g. liver boundaries and lesion margins) that can be used to guide regularisation in the ill-posed reconstruction of post-SIRT ^{90}Y PET and SPECT.

In addition, CT is routinely used for computed tomography attenuation correction (CTAC) in both PET and SPECT. The reconstructed HU image is first converted



(a) Bremsstrahlung emission spectrum.



(b) Detected spectrum (10 cm, AnyScan).

Figure 2.6: Bremsstrahlung spectra: (a) emitted spectrum from GATE FastY90 model [18, 19] and (b) detected (after attenuation by detector cover and backscattering) using AnyScan Trio GATE scanner model [20].

to an estimate of the linear attenuation coefficient map $\mu(E) \in \mathbb{R}_+^J$ at the relevant photon energy E (e.g. $E = 511$ keV for PET, or the photopeak energy for SPECT). For a given line of response (PET) or collimator ray (SPECT) indexed by i , the corresponding attenuation factor is then modelled as

$$a_i(E) = \exp\left(-\int_{\text{LoR}_i} \mu(\mathbf{x}, E) ds\right) \approx \exp\left(-\sum_{j=1}^J l_{ij} \mu_j(E)\right), \quad (2.12)$$

where $\mu_j(E)$ denotes the attenuation coefficient in voxel j and $l_{ij} \geq 0$ is the intersec-




Figure 2.7: Tangential profile of measured prompt counts and estimated additive terms (randoms+scatter), with the corresponding corrected prompts for an ^{90}Y PET scan. The data are summed over all views and segments, leaving counts per tangential bin.

tion length of LoR_i with voxel j . This is incorporated either directly into the system matrix (as part of $A_{\text{geom,att}}$) or applied as a multiplicative correction in projection space (A_{att}). Misregistration (e.g. respiratory motion), metal artefacts, and beam-hardening residuals can all propagate into CTAC and therefore bias quantitative PET/SPECT activity estimates [22, 23].

Chapter 3

Penalised-Likelihood Reconstruction: Priors and Optimisation Algorithms

Emission tomography reconstruction is an ill-posed inverse problem: the unknown activity must be inferred from noisy Poisson count data through a highly smoothing and poorly conditioned forward operator. While maximum-likelihood (ML) approaches remain the clinical workhorse [24, 25], their practical success relies heavily on early stopping, which acts as an implicit regulariser but can introduce iteration-dependent bias and compromise quantification [26, 27]. For quantitative, low-count, and multi-modality settings, precisely those that arise in modern PET/SPECT practice, explicit regularisation and fast, robust optimisation become essential.

This chapter reviews the modelling and algorithmic foundations that underpin the remainder of the thesis. We first formalise the Poisson statistical model and the penalised-likelihood reconstruction objective (often termed maximum a posteriori (MAP)), highlighting where ill-conditioning and noise amplification enter the problem. We then present some regularisation strategies used in emission tomography: from classical quadratic smoothness priors and total variation (TV) [28], through clinically deployed edge-preserving models such as the relative difference prior (RDP) [29], to anatomically guided methods that use aligned structural images (e.g. Bowsher [30], parallel level sets (PLS) [31], and kernelised representations [32, 33]. We then focus on synergistic regularisation, where multiple co-registered modalities

are coupled during reconstruction, and motivate Jacobian-based vectorial priors such as joint total variation (JTV) and total nuclear variation (TNV) [34, 35, 36], which will form the basis of the synergistic models developed later in the thesis.

Finally, we review optimisation algorithms suitable for large-scale maximum a posteriori (MAP) reconstruction. After discussing EM-type updates and their subset-accelerated variants, we highlight why subset methods can exhibit limit-cycle behaviour and why convergent penalised algorithms typically require relaxation (e.g. block-sequential regularised expectation maximisation (BSREM) [37]). We then introduce variance-reduced subset methods (e.g. stochastic average gradient amélioré (SAGA) and stochastic variance-reduced gradient (SVRG)) as a mechanism to retain low per-iteration cost while suppressing subset-induced stochasticity, and discuss the role of preconditioning, including “prior-aware” diagonal curvature approximations, in accelerating and stabilising convergence [38, 39].

3.1 The Inverse Problem in Emission Tomography

To discretely represent the spatial distribution of tracer concentration within a patient, the continuous activity distribution $f(r)$ is approximated using a basis expansion $f(r) \approx \sum_{j=1}^J x_j h_j(r)$. Typically, $\{h_j\}$ are chosen to be indicator functions on uniform cuboid voxels. In that case, it is customary to use a non-negative activity image vector $x = (x_1, \dots, x_J)^\top \in \mathbb{R}_+^J$.

The measured data consist of coincidence counts along lines of response (PET) or projection bins (SPECT). We denote the observed count in bin i as y_i , and model it as a realisation of a Poisson counting process,

$$y_i \sim \text{Poisson}(Y_i), \quad i = 1, \dots, I, \quad (3.1)$$

where $Y_i = \mathbb{E}[y_i]$ is the (unknown) mean count, which we model as $Y_i \approx \bar{y}_i(x)$ with

$$\bar{y}_i(x) = \sum_{j=1}^J a_{ij} x_j + b_i, \quad i = 1, \dots, I. \quad (3.2)$$

For notational simplicity, we will use the vector form

$$\bar{y}(x) = Ax + b, \quad \bar{y}_i(x) = (Ax + b)_i, \quad (3.3)$$

where $A = [a_{ij}] \in \mathbb{R}_+^{I \times J}$ and $b \in \mathbb{R}_+^I$. Assuming independence of measurement bins, we obtain

$$p(y | x) = \prod_{i=1}^I \frac{\bar{y}_i(x)^{y_i} \exp(-\bar{y}_i(x))}{y_i!}. \quad (3.4)$$

Taking logarithms yields the Poisson log-likelihood

$$\log p(y | x) = \sum_{i=1}^I (y_i \log \bar{y}_i(x) - \bar{y}_i(x) - \log(y_i!)). \quad (3.5)$$

Since the term $\sum_{i=1}^I \log(y_i!)$ does not depend on x , it can be dropped for optimisation purposes. Thus, maximising the likelihood w.r.t. x is equivalent to minimising

$$\mathcal{L}(x) := -\log p(y | x) + \text{const} = \sum_{i=1}^I (\bar{y}_i(x) - y_i \log \bar{y}_i(x)). \quad (3.6)$$

Unconstrained maximum-likelihood (ML) solutions are dominated by noise due to the poor conditioning of A and the Poisson noise inherent in y . While iterative expectation–maximisation (EM) methods such as MLEM and OSEM [24, 25] are the clinical standard, they rely on early stopping to act as an implicit regulariser. This yields clinically acceptable images but introduces solution-dependent bias and often fails to converge to a quantitative solution [26]. In addition, finding the ML solution for quantification can be prohibitively slow [27].

A robust alternative to early stopping is the Bayesian formulation. By Bayes' rule, we have

$$p(x | y) = \frac{p(y | x) p(x)}{p(y)}, \quad (3.7)$$

where $p(y | x)$ is the Poisson likelihood induced by the forward model, $p(x)$ encodes prior information about the unknown image, and the evidence

$$p(y) = \int p(y | x) p(x) dx \quad (3.8)$$

acts as a normalising constant independent of x .

A maximum a posteriori (MAP) estimate is then defined as

$$\hat{x}_{\text{MAP}} \in \arg \max_{x \in \mathcal{C}} p(x | y) = \arg \max_{x \in \mathcal{C}} p(y | x) p(x), \quad (3.9)$$

where $\mathcal{C}$ denotes a feasible set (typically enforcing non-negativity). Taking negative logarithms, as above, gives an equivalent minimisation problem:

$$\hat{x}_{\text{MAP}} \in \arg \min_{x \in \mathcal{C}} \{-\log p(y | x) - \log p(x)\}, \quad (3.10)$$

since $-\log p(y)$ is constant with respect to x .

It is common to express the prior in Gibbs form,

$$p(x) \propto \exp(-\beta \mathcal{R}(x)), \quad (3.11)$$

where $\mathcal{R}(x)$ is a regularisation functional and $\beta > 0$ controls the strength of the prior. Substituting this into the MAP objective yields the penalised-likelihood form

$$\hat{x}_{\text{MAP}} \in \arg \min_{x \in \mathcal{C}} \{\mathcal{L}(x) + \beta \mathcal{R}(x)\}. \quad (3.12)$$

3.2 Regularisation Strategies

The choice of $\mathcal{R}(x)$ determines the texture and quantitative characteristics of the reconstructed image. Historically, this has evolved from simple smoothing constraints to sophisticated, anatomically guided models.

3.2.1 Image-based Regularisation

One of the simplest forms of regularisation is Tikhonov regularisation, which imposes an ℓ_2 -norm penalty on local voxel-wise differences. This yields the quadratic prior

$$\mathcal{R}_{\text{Quad}}(x) = \sum_{j=1}^J \sum_{k \in \mathcal{N}_j} w_{jk} (x_j - x_k)^2, \quad (3.13)$$

where $\mathcal{N}_j$ denotes a local neighbourhood of voxel j (e.g. 6, 18, or 26-connected in 3D) and $w_{jk} \geq 0$ are weights (often distance-based and symmetric). While strongly convex and computationally efficient, the quadratic penalty strongly discourages large local differences, which tends to oversmooth edges and suppress high-contrast structures [40].

To better preserve edges, total variation (TV) [28] replaces the squared penalty with an ℓ_1 -type penalty on local differences. A neighbourhood-difference form that makes the connection to $\mathcal{R}_{\text{Quad}}$ explicit is

$$\mathcal{R}_{\text{TV}}(x) = \sum_{j=1}^J \sum_{k \in \mathcal{N}_j} w_{jk} |x_j - x_k|. \quad (3.14)$$

Relative to $\mathcal{R}_{\text{Quad}}$, the linear growth of $|x_j - x_k|$ reduces the marginal cost of large voxel-wise jumps in intensity, thereby promoting piecewise-constant reconstructions with sharper edges.

In much of the imaging literature, TV is instead written in terms of a discrete gradient operator. Defining $(Dx)_j \in \mathbb{R}^d$ to collect finite differences at voxel j along the d coordinate directions, a common choice on a regular grid is

$$((Dx)_j)^{(\ell)} = \frac{x_{j+\delta_\ell} - x_j}{\Delta_\ell}, \quad \ell = 1, \dots, d, \quad (3.15)$$

where Δ_ℓ denotes the Euclidean distance between voxels $j + \delta_\ell$ and j . The discrete TV functional is then

$$\mathcal{R}_{\text{TV}}(x) = \sum_{j=1}^J |(Dx)_j|, \quad (3.16)$$

with an isotropic alternative, more common in emission tomography [41, 42],

$$\mathcal{R}_{\text{TV}}(x) = \sum_{j=1}^J \|(Dx)_j\|_2. \quad (3.17)$$

For optimisation, it is common to use a smooth approximation to avoid nondif-

ferentiability at zero, for example

$$\tilde{\mathcal{R}}_{\text{TV}}(x) = \sum_{j=1}^J \sqrt{\|(Dx)_j\|_2^2 + \varepsilon^2}, \quad \varepsilon > 0. \quad (3.18)$$

In clinical practice, adaptable edge-preserving priors are often preferred. The relative difference prior (RDP) [29], the core component of the GE “Q.Clear” algorithm, modifies the penalty to depend on the absolute intensity sums, penalising relative differences, as opposed to absolute ones.

$$\mathcal{R}_{\text{RDP}}(x) = \sum_{j=1}^J \sum_{k \in \mathcal{N}_j} w_{jk} \frac{(x_j - x_k)^2}{(x_j + x_k) + \gamma|x_j - x_k|}, \quad (3.19)$$

with γ controlling the amount of edge preservation. Note that if $(x_j + x_k) \gg \gamma|x_j - x_k|$ then $\mathcal{R}_{\text{RDP}}$ behaves like a relative quadratic prior, whereas if $\gamma|x_j - x_k| \gg (x_j + x_k)$ it behaves like an anisotropic TV penalty (up to the factor $1/\gamma$).

3.2.2 Anatomically Guided Regularisation

In low-count regimes, single-modality priors struggle to distinguish between noise and legitimate small structures. Anatomically guided reconstruction uses a high-resolution structural image u (e.g., CT) to modulate the regularisation of the emission image x .

Geometric Alignment methods assume that if an edge exists in the emission image, it should be supported by corresponding boundaries in a registered anatomical guidance image. The seminal Bowsher prior [30] implements this idea by defining, for each voxel j , a data-adaptive and dynamic neighbourhood $\mathcal{B}_j(u)$ using the guidance image u . Concretely, from a fixed candidate neighbourhood $\mathcal{N}_j$, the Bowsher construction selects the B voxels whose intensities in u are most similar to u_j , i.e. those minimising a similarity metric such as

$$d_{jk}(u) := |u_j - u_k| \quad (k \in \mathcal{N}_j), \quad (3.20)$$

and sets $\mathcal{B}_j(u)$ to be the indices of the B smallest values of $d_{jk}(u)$. Regularisation

is then applied only over this anatomically consistent subset, for example via a quadratic penalty

$$\mathcal{R}_{\text{Bow}}(x; u) = \sum_{j=1}^J \sum_{k \in \mathcal{B}_j(u)} w_{jk} (x_j - x_k)^2, \quad (3.21)$$

(or an alternative edge-preserving potential). By restricting the penalty to neighbours that are locally homogeneous in u , the prior encourages smoothing within anatomical regions while reducing smoothing across anatomical edges, thereby allowing sharp boundaries to form in x when they are supported by the guidance image.

An even more geometric framework is parallel level sets (PLS) [31]. These methods encourage structural consistency by encouraging the gradient vectors of the emission image and the guidance image to be aligned, i.e. $Dx_j \parallel Du_j$. A common formulation (often termed PLS1) penalises the magnitude of the component of Dx_j orthogonal to Du_j , which is proportional to the sine of the enclosed angle θ_j :

$$R_{\text{PLS1}}(x, u) = \sum_{j=1}^J \|Dx_j\|_2 \|Du_j\|_2 |\sin \theta_j|. \quad (3.22)$$

This term vanishes when Dx_j and Du_j are parallel (or when either gradient is zero), and increases with the component of Dx_j orthogonal to Du_j . While numerically stable, PLS1 scales with $\|Du_j\|_2$ and therefore inherits the spatially varying gradient magnitude of the anatomical image, which can lead to over-weighting at strong anatomical edges and under-weighting in low-contrast regions.

To decouple the directional information from the anatomical gradient magnitude, directional total variation (DTV) [43] introduces a normalised direction field

$$\xi_j = \frac{Du_j}{\sqrt{\|Du_j\|_2^2 + \eta^2}}, \quad (3.23)$$

where $\eta > 0$ is a small stabilisation parameter that prevents division by zero in homogeneous regions ($Du_j \approx 0$) and suppresses unreliable directions induced by noise or very fine-scale anatomical texture. Using ξ_j , one defines a local anisotropic

projection operator

$$\Pi_j := I - \gamma_j \xi_j \xi_j^\top, \quad (3.24)$$

where I is the $d \times d$ identity matrix and $\gamma_j \in [0, 1]$ controls the strength of directional guidance. The resulting DTV prior is

$$R_{\text{DTV}}(x; u) = \sum_{j=1}^J \|\Pi_j D x_j\|_2, \quad (3.25)$$

which penalises predominantly the gradient component orthogonal to the anatomical level-set direction, thereby encouraging edges in x to align with edges in u . In regions where u is locally flat (so $\xi_j \approx 0$), $\Pi_j \approx I$ and DTV reduces to standard isotropic TV.

Regularisation strategies also exist outside the standard Bayesian (penalised-likelihood) framework. A prominent example is the kernelised expectation maximisation (KEM), in which prior information is encoded directly in the image representation rather than through an explicit penalty term. Following Wang and Qi [32] and its anatomically-aided extension [33], the emission image is modelled as a linear combination of kernel-derived basis functions:

$$x = K \alpha, \quad (3.26)$$

where $\alpha \in \mathbb{R}^J$ is a coefficient image and $K \in \mathbb{R}^{J \times J}$ is a (typically sparse) kernel matrix whose entries encode similarity between voxels based on features extracted from a guidance image [33]. Each voxel j is assigned a feature vector f_j (commonly a patch of intensities in the high-resolution guidance image), and a radial Gaussian kernel is often used,

$$\kappa(f_j, f_k) = \exp\left(-\frac{\|f_j - f_k\|_2^2}{2\sigma^2}\right), \quad (3.27)$$

with σ controlling edge sensitivity. The kernel matrix is then defined entrywise by $K_{jk} = \kappa(f_j, f_k)$, with additional sparsification for tractability: for each voxel j , only a small set of the most similar neighbours (e.g. via a local k -nearest-neighbour search within a window) are allowed to contribute, yielding a sparse K . To preserve overall intensity scaling (counts), K is typically row-normalised.

Substituting $x = K\alpha$ into the Poisson forward model gives a kernelised expectation for the data,

$$\bar{y}(\alpha) = A(K\alpha) + b, \quad (3.28)$$

so that reconstruction reduces to maximum-likelihood estimation of α under the same Poisson model, followed by $x = K\hat{\alpha}$. Importantly, because the model remains linear in the unknown coefficients, an EM-type update can be derived with essentially the same structure as conventional MLEM/OSEM, but with K embedded within the projector/backprojector operations. This yields a procedure that is straightforward to implement in existing EM pipelines, while often delivering improved bias–variance trade-offs at low counts by enforcing anatomically-informed, spatially adaptive smoothing in the image representation [32, 33].

A known limitation of purely anatomy-derived kernels is PET-unique feature suppression: structures present in the emission image but absent (or weak) in the guidance image can be over-smoothed or biased towards anatomically plausible but functionally incorrect solutions [44]. To mitigate this, hybrid kernelised expectation maximisation (HKEM) augments the kernel with synergistic information from the evolving PET estimate itself, yielding an iteration-dependent kernel that combines anatomical and emission-driven similarity. In the list-mode HKEM formulation of Deidda *et al.*, the hybrid kernel entry factorises as

$$k_{fj}^{(n)} = k_m(v_f, v_j) k_p(z_f^{(n)}, z_j^{(n)}), \quad (3.29)$$

where v_j denotes the (fixed) anatomical feature vector at voxel j and $z_j^{(n)}$ is a feature derived from the current PET update (updated each iteration) [44]. Both factors are commonly chosen as Gaussian similarities modulated by spatial proximity,

$$k_m(v_j, v_k) = \exp\left(-\frac{\|v_j - v_k\|_2^2}{2\sigma_a^2}\right) \exp\left(-\frac{\|x_j - x_k\|_2^2}{2\sigma_{da}^2}\right), \quad (3.30)$$

$$k_p(z_j^{(n)}, z_k^{(n)}) = \exp\left(-\frac{\|z_j^{(n)} - z_k^{(n)}\|_2^2}{2\sigma_p^2}\right) \exp\left(-\frac{\|x_j - x_k\|_2^2}{2\sigma_{df}^2}\right), \quad (3.31)$$

where x_j denotes the voxel coordinates and the parameters $\sigma_a, \sigma_f, \sigma_{da}, \sigma_{df}$ control the relative influence of anatomical similarity, functional-driven (emission) similarity, and spatial locality. The functional term helps address emission–guidance mismatch and, crucially, reduces the tendency of anatomy-only kernels to suppress PET-unique lesions by allowing the neighbourhood coupling to adapt to features emerging in the emission estimate. Empirically, Deidda *et al.* report that this hybridisation can reduce negative bias under low-count conditions and improve lesion contrast while retaining the denoising benefits of kernelised reconstruction [44].

Kernel approaches extend naturally to multiple guidance sources by enlarging the feature vectors (e.g. concatenating features from CT and/or additional functional channels) and defining the kernel on the resulting joint feature space; this provides a flexible mechanism for sequential or multi-modality guidance without introducing an explicit Bayesian penalty term.

3.2.3 Synergistic Regularisation

If multiple modalities are available that depict the same underlying anatomy or functional distribution, it can be advantageous to reconstruct them synergistically, allowing information to be shared across modalities throughout the reconstruction process. A prerequisite for such coupling is that the images are defined in a common spatial coordinate system. In practice, this requires geometric registration between modalities and resampling onto a common reconstruction grid, so that voxel index n refers to the same physical location across all channels.

When the native modalities are acquired on different grids or coordinate frames, we introduce spatial transforms W_m (registration + resampling operators) and define the co-registered images on a shared grid, with $W_m : \mathbb{R}_+^{J_m} \rightarrow \mathbb{R}_+^N$:

$$z_m := W_m x_m \in \mathbb{R}_+^N, \quad m = 1, \dots, M, \quad (3.32)$$

and collect them as the modality tuple

$$\mathbf{z} := (z_1, \dots, z_M). \quad (3.33)$$

Coupling priors are then naturally expressed in terms of $\mathbf{z}$, since neighbourhoods, gradients, and voxel-wise comparisons are only meaningful when images are spatially aligned.

In the Bayesian setting, a generic multi-modality reconstruction is given by

$$\hat{\mathbf{x}} \in \arg \min_{\mathbf{x} \in \mathcal{C}} \left\{ \sum_{m=1}^M \mathcal{L}_m(x_m) + \beta \mathcal{R}(\mathbf{z}) \right\}, \quad (3.34)$$

where $\mathbf{x} := (x_1, \dots, x_M)$, $\mathbf{z} = \mathbf{W}\mathbf{x}$ with $\mathbf{W} := \text{diag}(W_1, \dots, W_M)$, and $\mathcal{R}(\mathbf{z})$ promotes cross-modality consistency (e.g. co-located edges or shared support) in the co-registered domain.

A practical way to exploit multiple modalities without solving a fully joint inverse problem is sequential reconstruction, in which one modality is first reconstructed, and the resulting image is then used as side-information to guide the reconstruction of another modality. In the Bayesian regime, this is

$$\hat{x}_m \in \arg \min_{x_m \in \mathcal{C}} \left\{ \mathcal{L}_m(x_m) + \beta \mathcal{R}(x_m; W_{m \leftarrow (m-1)} x_{m-1}) \right\}, \quad (3.35)$$

where $W_{m \leftarrow (m-1)}$ maps the previous-stage estimate $\hat{x}_{m-1}$ into the spatial grid of modality m .

Deidda *et al.* have previously investigated a similar method for ^{90}Y bremsstrahlung SPECT-informed PET reconstructions [45]. Bremsstrahlung SPECT images are first reconstructed using HKEM with CT and SPECT image-update information to obtain an enhanced SPECT estimate. This SPECT image is then resampled to the PET grid and used as functional guidance for PET reconstruction.

$$\hat{\alpha}_m \in \arg \min_{\alpha_m \in \mathcal{C}} \left\{ \mathcal{L}_m(W_{m \leftarrow (m-1)} K_m \alpha_m) \right\}. \quad (3.36)$$

The paper reports that this sequential design is particularly beneficial for low-count ^{90}Y PET: SPECT provides complementary denoising/uptake support within PET-SPECT-consistent regions, while CT contributes primarily through resolution/edge information injected at the SPECT stage, and the PET image-update features reduce

suppression of PET-unique structures and mitigate inter-modality mismatch.

Sequential guidance is limited by one-way dependency: errors in u may propagate directly to x and the shared information may not be fully utilised. Truly synergistic reconstruction estimates both images simultaneously, allowing information to flow bi-directionally. Perhaps the most obvious way to extend single modality regularisation to achieve a synergistic regulariser is joint total variation (JTV) [46], a simple extension of TV which sums gradient magnitudes:

$$\mathcal{R}_{\text{JTV}}(\mathbf{z}) = \sum_{n=1}^N \sqrt{\|Dz_{1,n}\|^2 + \|Dz_{2,n}\|^2 + \cdots + \|Dz_{M,n}\|^2}. \quad (3.37)$$

This encourages joint sparsity, encouraging co-located edges. A convenient way to formalise multi-modality edge co-localisation is to view the set of images as a vector-valued field. For each voxel n , the voxel-wise Jacobian is defined by stacking image gradients (Equation 3.15) in columns:

$$\mathcal{G}_n(\mathbf{z}) := \begin{bmatrix} Dz_{1,n} & Dz_{2,n} & \cdots & Dz_{M,n} \end{bmatrix} \in \mathbb{R}^{d \times M}. \quad (3.38)$$

When viewed in this way, JTV can be redefined as a Schatten 2-norm (also known as a Frobenius norm) applied to this jacobian [34],

$$\mathcal{R}_{\text{JTV}}(\mathbf{z}) = \sum_{n=1}^N \|\mathcal{G}_n(\mathbf{z})\|_F \quad (3.39)$$

where the Schatten p -norm is defined as

$$\|Y\|_{S_p} := \left(\sum_{r=1}^R \sigma_r(Y)^p \right)^{1/p} \quad (3.40)$$

for singular values $\sigma_\ell(Y)$ of a rank- R matrix Y . Note that $R = \min\{d, M\}$. The Schatten 1-norm, also known as the nuclear norm is the sum of singular values

$$\|Y\|_* = \sum_{r=1}^R \sigma_r(Y). \quad (3.41)$$

The nuclear norm is the convex envelope of the Rank function. Applied to the voxel-wise jacobians, it promotes singular value sparsity and has the effect of either driving functionally misaligned gradients to zero or parallel/anti-parallel, promoting edge-alignment. Applying this nuclear norm to the voxel-wise jacobians is known as total nuclear variation (TNV) [34, 35].

$$\mathcal{R}_{\text{TNV}}(\mathbf{z}) = \sum_{n=1}^N \|\mathcal{G}_n(\mathbf{z})\|_* \quad (3.42)$$

It has been applied to synergistic PET/MR reconstruction through its extended variant, total generalised nuclear variation [36].

The benefits of using the nuclear norm over the Frobenius norm can be visualised in Figure 3.1, where the nuclear norm exhibits sharp grooves along the coordinate axes that encourage one singular value to vanish (promoting low-rank structure), whereas the Frobenius norm is smooth and tends to shrink singular values jointly without explicitly favouring rank reduction.

Total Nuclear variation will form the basis for the synergistic prior demonstrated later on in this thesis.

3.3 Optimisation Algorithms

The reconstruction of MAP (and, even more so, synergistic images) requires more nuanced algorithms. maximum-likelihood expectation–maximisation (MLEM) can be relatively simply extended for penalised reconstruction via Green’s one-step-late-method [47]. The one-step-late method constructs an EM-type multiplicative update by evaluating the penalty derivative at the current iterate $x^{(n)}$ rather than at $x^{(n+1)}$. The resulting OSL-MAP update is

$$x_j^{(n+1)} = \frac{x_j^{(n)}}{\sum_{i=1}^I a_{ij} + \beta \left. \frac{\partial R(x)}{\partial x_j} \right|_{x=x^{(n)}}} \sum_{i=1}^I a_{ij} \frac{y_i}{\bar{y}_i(x^{(n)})}, \quad j = 1, \dots, J, \quad (3.43)$$

The term $\left. \frac{\partial R}{\partial x_j} \right|_{x^{(n)}}$ acts as an iteration-dependent modification of the MLEM denominator, yielding an algorithm that is typically simple to implement for a broad




Figure 3.1: Visualisation of the Frobenius and nuclear norms (vertical axis) of a matrix X as functions of its (two) singular values (horizontal axis). The Frobenius norm is $\|X\|_F = \sqrt{\sigma_1(X)^2 + \sigma_2(X)^2}$ (red mesh), whereas the nuclear norm is $\|X\|_* = |\sigma_1(X)| + |\sigma_2(X)|$ (black mesh). For visualisation purposes we extend the singular-value axes to negative values. The nuclear norm exhibits sharp grooves along the coordinate axes, encouraging one singular value to vanish (i.e. promoting low-rank structure), while the Frobenius norm is smooth and does not preferentially drive a singular value exactly to zero. This behaviour can also be understood via subgradients: the nuclear-norm subgradient (cyan) points towards rank-1 grooves, preferentially shrinking the smaller singular value, whereas the Frobenius-norm subgradient (red) points towards the origin, shrinking both singular values equally.

class of differentiable priors. Unfortunately, the one-step-late method is not convergent and special attention must be paid to initialisation conditions for a reasonable reconstruction [48].

Although too fragile for low count / synergistic reconstruction, the one-step-late algorithm has some useful takeaways. Primarily the joint update of data and prior functions with some kind of voxel-wise scaling. This is otherwise known as a preconditioner. MLEM has the multiplicative update equation for iteration k :

$$x^{(k+1)} = \left(x^{(k)} \oslash A^T \mathbf{1}_I \right) \odot A^T (y \oslash \bar{y}), \quad (3.44)$$

where $\odot$ and $\oslash$ denote elementwise multiplication and division, respectively, $\mathbf{1}_I \in \mathbb{R}^I$

is an all-ones vector and k denotes the iteration number. For notational simplicity, we drop the explicit elementwise operators; all divisions are understood componentwise.

Then

$$x^{(k+1)} = \text{diag}\left(\frac{x^{(k)}}{A^\top \mathbf{1}_I}\right) A^\top \left(\frac{y}{\bar{y}}\right). \quad (3.45)$$

This can be rearranged to form the additive update equation

$$\begin{aligned} x^{(k+1)} &= x^{(k)} + \text{diag}\left(\frac{x^{(k)}}{A^\top \mathbf{1}_I}\right) \left(A^\top \left(\frac{y}{\bar{y}(x^{(k)})}\right) - \mathbf{1}_J\right) \\ &= x^{(k)} + \text{diag}\left(\frac{x^{(k)}}{A^\top \mathbf{1}_I}\right) \left(A^\top \left(\frac{y}{\bar{y}(x^{(k)})}\right) - A^\top \mathbf{1}_I\right) \\ &= x^{(k)} + \text{diag}\left(\frac{x^{(k)}}{A^\top \mathbf{1}_I}\right) A^\top \left(\frac{y}{\bar{y}(x^{(k)})} - \mathbf{1}_I\right) \\ &= x^{(k)} - \text{diag}\left(\frac{x^{(k)}}{A^\top \mathbf{1}_I}\right) D\mathcal{L}(x^{(k)}). \end{aligned} \quad (3.46)$$

where $\mathbf{1}_I \in \mathbb{R}^J$ is an all-ones vector. This form highlights that the MLEM update can be interpreted as a preconditioned gradient descent algorithm with preconditioner $P_{EM}^{(k)} = \text{diag}\left(\frac{x^{(k)}}{A^\top \mathbf{1}_I}\right)$, i.e.

$$x^{(k+1)} = x^{(k)} - \varsigma P_{EM}^{(k)} D_x \mathcal{L}(x^{(k)}), \quad (3.47)$$

with a step size $\varsigma = 1$.

A preconditioner aims to recondition the reconstruction space so that gradient steps are appropriately scaled across voxels, improving numerical conditioning and convergence. Intuitively, we seek larger steps in poorly sensitive directions and smaller steps in highly sensitive directions, thereby approximating curvature information. Since curvature is governed by the inverse Hessian, it is often approximated in a diagonal form for computational efficiency.

For the Poisson negative log-likelihood, the expected Hessian (Fisher information) is

$$\mathcal{I}(x) := \mathbb{E}[\nabla_x^2 \mathcal{D}(y | x)] = A^\top \text{diag}\left(\frac{1}{\bar{y}(x)}\right) A. \quad (3.48)$$

Its diagonal entries satisfy, for each voxel index j ,

$$\mathcal{I}_{jj}(x) = \sum_i \frac{a_{ij}^2}{\bar{y}_i(x)} \leq \sum_i \frac{a_{ij}^2}{a_{ij}x_j} = \frac{(A^\top \mathbf{1})_j}{x_j}, \quad (3.49)$$

where we used $\bar{y}_i(x) = \sum_\ell a_{i\ell}x_\ell + r_i \geq a_{ij}x_j$ and $r_i \geq 0$. Consequently,

$$\mathcal{I}_{jj}(x) \leq \frac{(A^\top \mathbf{1})_j}{x_j}. \quad (3.50)$$

This motivates the diagonal scaling

$$P_{\text{EM}}^{(k)} := \text{diag} \left(\frac{x^{(k)}}{A^\top \mathbf{1}} \right), \quad (3.51)$$

which can be interpreted as the inverse of a conservative diagonal upper bound on the per-voxel curvature of the Poisson negative log-likelihood (in expectation).

Subsets are commonly used to accelerate early convergence by partitioning the data into S subsets and updating the image using only the contribution from a single subset (or a small batch of subsets) at each iteration. Writing the negative log-likelihood as a sum of subset terms,

$$\mathcal{L}(x) = \sum_{s=1}^S \mathcal{L}_s(x), \quad (3.52)$$

a generic subset (stochastic / incremental) scaled-gradient update takes the form

$$x^{(k+1)} = x^{(k)} - \varsigma \mathbf{P}^{(k)} D_x \mathcal{L}_s(x^{(k)}). \quad (3.53)$$

If multiple subsets are used per iteration, $D_x \mathcal{L}_s$ is replaced by an average (or sum) over the chosen subset indices. The most common algorithm found in clinical systems is the ordered subsets expectation–maximisation (OSEM) algorithm. For subset s at iteration k , a standard OSEM update can be written as

$$x^{(k+1)} = \text{diag} \left(\frac{x^{(k)}}{A_s^\top \mathbf{1}_I} \right) A_s^\top \left(\frac{y_s}{\bar{y}_s(x^{(k)})} \right), \quad (3.54)$$

where A_s denotes the system matrix and data restricted to subset s .

OSEM typically shows a marked improvement in early-stage convergence by providing a cheaper, higher-variance approximation of the full-data EM update. However, because the subset gradient is only an approximation to the full gradient, OSEM is not, in general, a convergent algorithm for the maximum-likelihood problem: as the iterates approach a solution, the stochasticity introduced by cycling through subsets can dominate, leading to limit-cycle behaviour in which successive subset updates fail to settle to a fixed point (Figure 3.2).




Figure 3.2: Demonstration of limit cycle behaviour with an ordered subset gradient descent algorithm. Top row: 3 different simple sub-functions representing subset likelihoods. Bottom row: left, full function with sub-gradient steps; centre, zoomed version of the full function demonstrating limit cycle behaviour (green dots show approximate points of limit cycle); right, likelihood value per iteration, demonstrating limit cycle precession.

We seek a fast and convergent algorithm for MAP reconstruction. A widely used approach is the (modified) block-sequential regularised expectation maximisation (BSREM) framework [37], which combines EM-type preconditioning and subset acceleration with a relaxation (diminishing step-size) to recover convergence while incorporating an explicit penalty. In a scaled-gradient form, a generic BSREM-type update can be written as

$$x^{(k+1)} = x^{(k)} - \varsigma_k P_{EM}^{(k)} \left(D_x \mathcal{L}_s(x^{(k)}) + \frac{\beta}{S} D_x \mathcal{R}(x^{(k)}) \right), \quad (3.55)$$

where $\varsigma_k > 0$ is an iteration dependent step size. Convergence is promoted by choosing $\{\varsigma_k\}$ to diminish appropriately (e.g. $\varsigma_k \rightarrow 0$ with $\sum_k \varsigma_k = \infty$), which suppresses the subset-induced limit-cycle behaviour while retaining rapid early progress (see Figure 3.3).




Figure 3.3: Demonstration of reduction in limit cycle behaviour with step size relaxation $\varsigma_k = \varsigma_0 / (1 + 0.05k)$. Left, full function with sub-gradient steps; centre, zoomed version of the full function demonstrating relaxed limit cycle behaviour; right, likelihood value per iteration, demonstrating relaxed limit cycle precession.

However, we can do better. Variance-reduced subset methods have become increasingly popular in the PET reconstruction literature because they retain the low per-iteration cost of subset updates while substantially reducing the subset-induced stochasticity that drives OSEM limit cycles. The key idea is to replace the raw subset gradient, which is a high-variance estimator of the full gradient, with a variance-reduced estimator that is (approximately) unbiased and whose variance decreases as the iterates approach a solution. This allows the use of larger (or non-vanishing) relaxation parameters than classical diminishing-step BSREM-type schemes, while maintaining stable convergence behaviour in practice [38, 39].

A generic subset (stochastic) scaled-gradient step uses $g_{s_k}(x^{(k)})$ in place of $\nabla \mathcal{L}(x^{(k)}) = \sum_s g_s(x^{(k)})$, i.e.

$$x^{(k+1)} = x^{(k)} - \varsigma_k \mathbf{P}^{(k)} \left(g_{s_k}(x^{(k)}) + \beta r(x^{(k)}) \right), \quad (3.56)$$

which is cheap but noisy. Variance reduction augments this step by incorporating additional information (stored gradients or periodic full-gradient computations) to form a corrected estimator $\tilde{g}_k$ with reduced variance. Two widely used examples are

stochastic average gradient amélioré (SAGA) and SVRG:

SAGA maintains a table of the most recently computed subset gradients $\{g_s(\cdot)\}_{s=1}^S$ and updates using

$$\tilde{g}_k = g_{s_k}(x^{(k)}) - g_{s_k}(\phi_{s_k}) + \frac{1}{S} \sum_{s=1}^S g_s(\phi_s), \quad (3.57)$$

where ϕ_s denotes the iterate at which the stored gradient for subset s was last evaluated.

SVRG periodically computes a reference full gradient at a snapshot $\tilde{x}$ and then performs inner iterations using

$$\tilde{g}_k = g_{s_k}(x^{(k)}) - g_{s_k}(\tilde{x}) + \nabla_x \mathcal{L}(\tilde{x}). \quad (3.58)$$

In both cases, $\tilde{g}_k$ is designed so that $\mathbb{E}[\tilde{g}_k] \approx \nabla_x \mathcal{L}(x^{(k)})$ while its variance decreases as $x^{(k)}$ approaches a stationary point, thereby mitigating the oscillatory behaviour seen in OSEM and improving stability at high iteration counts (see Figure 3.4).




Figure 3.4: Demonstration of removal of limit cycle behaviour with an SVRG-type (Ordered VRG) algorithm. Left, full function with sub-gradient steps; centre, zoomed version of the full function demonstrating removal of limit cycle behaviour; right, likelihood value per iteration, demonstrating removal of limit cycle precession

A variance-reduced, prior-regularised update can then be written compactly as

$$x^{(k+1)} = x^{(k)} - \varsigma P^{(k)} \left(\tilde{g}_k + \beta r(x^{(k)}) \right). \quad (3.59)$$

Recent work indicates that incorporating the effect of the prior into $P^{(k)}$ (“prior-aware preconditioning”) can be crucial for fast and robust convergence. Ehrhardt *et al.*

[38] showed substantially improved convergence rates with RDP-regularised PET reconstruction when using harmonic averaging of the EM-type preconditioner and a scaled diagonal estimate of the inverse Hessian of the prior. Using only the diagonal elements of the RDP Hessian

$$\text{diag}(\mathcal{H}_{\text{RDP}}(x_j)) = \sum_{k \in \mathcal{N}_j} \frac{x_k^2}{(x_j + x_k + \gamma|x_j - x_k|)^3}, \quad (3.60)$$

they constructed a preconditioner as

$$P^{(k)} = \text{diag} \left(\frac{x^{(k)} + \delta}{A^T 1 + \alpha \text{diag}(\beta \mathcal{H}_{\text{RDP}}(x^{(k)})) (x^{(k)} + \delta)} \right) \quad (3.61)$$

with scalar α tuned to correct for the loss of off-diagonal information. Although this is not emphasised explicitly in [38], the harmonic averaging viewpoint can also be interpreted as a majorisation-preserving way to combine curvature information from the likelihood and the prior. In particular, the voxel-wise harmonic blend can be rearranged as

$$\left(\frac{1}{a} + \frac{1}{b} \right)^{-1}, \quad (3.62)$$

so that combining two diagonal scalings a and b harmonically is equivalent to inverting the sum of the corresponding inverse-scalings (i.e. two diagonal Hessian surrogates). Consequently, when these inverse-scalings are chosen as upper-bounding approximations of the respective Hessians, their sum remains a valid upper bound, and the resulting harmonic combination inherits this majorising property while still incorporating both data-fidelity and prior-induced curvature.

Preconditioning and stochastic variance reduction for synergistic MAP reconstruction has not been well investigated in the literature so far and will form a part of this thesis.

A widely used alternative to EM-type and subset-gradient schemes is the limited-memory BFGS algorithm with bound constraints (L-BFGS-B) [49, 50, 51, 52, 53, 54], a quasi-Newton method that targets $\min_{x \in \mathcal{C}} \{\mathcal{L}(x) + \beta \mathcal{R}(x)\}$ by iteratively building a low-rank approximation to the inverse Hessian from successive iterate and gradient

differences. The “limited-memory” construction stores only a small number of correction pairs, making it viable for high-dimensional images, while the “B” variant enforces simple bounds (e.g. non-negativity) via an active-set strategy and projected search directions. In practice, L-BFGS-B can exhibit rapid local convergence on smooth penalised objectives, but each iteration requires one or more full objective and gradient evaluations (coupled with a line search), which is computationally burdensome in large-scale emission tomography when the likelihood is evaluated over all subsets. For this reason it is not used in the main reconstruction experiments in this thesis; nevertheless, it provides a useful reference optimiser and is employed later for a small demonstrative deconvolution problem in Chapter 5.

A practical limitation of vanilla L-BFGS-B in emission tomography is its sensitivity to poor variable scaling: the line search and the quasi-Newton curvature updates can be dominated by a small subset of high-sensitivity voxels, leading to slow progress. Tsai *et al.* propose a diagonally preconditioned variant (L-BFGS-B-PC) that embeds analytic curvature information into the optimisation via a change of variables [55]. We can choose a diagonal preconditioner D from a cheap curvature surrogate (e.g. an EM-type diagonal approximation of the Hessian), and run L-BFGS-B on the transformed problem Dx . This rescales the objective so that gradients are more isotropic in the transformed coordinates, while bound constraints are preserved by transforming the box constraints accordingly. The resulting method combines low-cost, problem-specific diagonal curvature (from D) with the low-rank curvature learned by L-BFGS-B, typically yielding substantially faster and more robust convergence than unpreconditioned L-BFGS-B in penalised reconstructions.

3.4 Summary and Research Gap

This chapter reviewed the statistical and algorithmic foundations of emission tomography reconstruction under a Poisson measurement model. Starting from maximum-likelihood estimation, we motivated penalised-likelihood (MAP) formulations as a principled alternative to early stopping, particularly in low-count and quantitatively demanding settings. We then surveyed regularisation strategies ranging from

quadratic smoothness and total variation to clinically deployed edge-preserving models such as the relative difference prior, before focusing on anatomically guided approaches (e.g. Bowsher, PLS/DTV) and kernelised representations that incorporate structural side-information through the image model.

A central theme in modern PET/SPECT practice is the availability of multiple co-registered modalities. While sequential (one-way) guidance can be effective, it is inherently limited by asymmetric information flow and by the risk of propagating errors or mismatches from the guidance image into the emission reconstruction. This motivates synergistic reconstruction, in which modalities are coupled through joint priors. In this setting, Jacobian-based penalties such as JTV and TNV provide a natural mechanism to promote co-located, directionally consistent edges across modalities while allowing modality-specific contrast.

From an optimisation perspective, clinical subset-accelerated EM methods achieve fast early progress but can exhibit limit-cycle behaviour and lack convergence guarantees in penalised problems without relaxation. Convergent frameworks such as BSREM recover stability via diminishing step sizes, but may converge slowly when the curvature induced by strong, spatially varying priors dominates the data term. Variance-reduced subset methods (e.g. SAGA/SVRG) offer a promising compromise, retaining low per-iteration cost while suppressing subset-induced stochasticity. However, their practical performance depends critically on effective preconditioning.

Research gap. Despite substantial progress in both synergistic regularisation and variance-reduced optimisation, two gaps remain directly relevant to quantitative ^{90}Y imaging:

1. **Synergistic, anatomy-informed coupling:** existing Jacobian-based priors (e.g. TNV) are typically formulated to promote cross-modality edge consistency, but are not routinely combined with explicit directional guidance derived from high-resolution CT edges in a way that remains robust in low-count regimes.
2. **Synergy-aware optimisation:** variance-reduced subset algorithms and prior-

aware diagonal preconditioners have mainly been studied in single-modality PET with clinically popular priors (e.g. RDP). Their extension to synergistic MAP reconstruction, where cross-modality coupling induces strong and spatially heterogeneous curvature, is comparatively underexplored.

Thesis direction. Motivated by these gaps, the remainder of this thesis develops a directionally guided, differentiable total nuclear variation prior that uses CT-derived edge information while coupling functional modalities through voxel-wise Jacobians. We then design and evaluate a variance-reduced subset optimisation scheme (SVRG-type) equipped with a synergy-aware preconditioner that accounts for both data fidelity and coupled prior curvature, targeting fast, stable convergence in low-count ^{90}Y reconstruction.

Chapter 4

Improvements to the ^{90}Y Bremsstrahlung System Model

Before tackling the problem of synergistic reconstruction with ^{90}Y , we need to properly model the physics of out

4.1 Challenges in Bremsstrahlung Imaging

As introduced in Section 2.4, unlike diagnostic radionuclides such as ^{99m}Tc that emit gamma photons at discrete photopeaks, ^{90}Y has no practically useful prompt gamma emission for SPECT. Imaging therefore relies on detecting the continuous spectrum of Bremsstrahlung photons generated as emitted β -particles decelerate in tissue [56]. This emission physics gives rise to several fundamental challenges for quantitative reconstruction:

- **Continuous energy spectrum:** The emitted photon energies span a continuum from (near) zero up to the maximum β -energy of 2.28 MeV [57] (see Figure 2.6). In conventional SPECT, scatter may be estimated using energy windows adjacent to a photopeak; for ^{90}Y , the absence of a photopeak makes it difficult to separate primary photons from those that have undergone Compton scatter within the patient. As a result, more sophisticated scatter-modelling and correction strategies are typically required [58].
- **Septal penetration and collimator design:** The high-energy tail of the Bremsstrahlung spectrum increases the probability of septal penetration, which

degrades spatial resolution and contrast. Collimator choice and optimisation are therefore critical to limit penetration while maintaining adequate sensitivity [59].

- **Bremsstrahlung-induced blur:** Bremsstrahlung photons are produced along the electron track rather than at a single point, so the effective photon emission location is distributed relative to the ^{90}Y decay site. This introduces an additional range-dependent blurring component that is not present in photopeak imaging and can further degrade reconstructed resolution [58] (see Figure.




Figure 4.1: Plots showing probability of bremsstrahlung gamma emission versus the distance from ^{90}Y decay site. Data from Rault *et al.*[18]

4.2 Optimisation Strategies and Monte Carlo

Modeling

To overcome these limitations, optimisation must occur at both the acquisition and reconstruction stages. While hardware choices, such as the preference for Medium Energy General Purpose (MEGP) or High Energy General Purpose (HEGP) collimators, can mitigate septal penetration [60, 61, 62], substantial gains in quantitative accuracy are achieved through advanced computational modeling [63].

Monte Carlo (MC) simulation codes, such as SIMIND [64] and GATE [19, 18], play a pivotal role in this optimisation. These tools allow for the generation of

accurate scatter kernels, as well as other effects that are specific to the continuous ^{90}Y spectrum. Some methods have integrated modelling of the spatial offset between the β -decay location and the point of Bremsstrahlung emission into MC forward projectors to refine the system matrix [58].

Furthermore, novel reconstruction techniques such as Joint Spectral Reconstruction (JSR) have been proposed to exploit the information contained across multiple energy windows, rather than relying on a single wide integration window [65]. By combining multi-band acquisition with accurate Monte Carlo-based forward modelling, it is possible to improve the precision of activity estimation and contrast recovery in post-SIRT dosimetry **Can I fit this in somehow? Probably not... No time.**

This chapter explores these advanced reconstruction methodologies. We also introduce a simple Gaussian image-based blurring operation, in order to account for bremsstrahlung range blur. In addition, we investigate the suitability of Fu & Qi's residual correction method [66], originally designed for correcting for positron range blur, on bremsstrahlung reconstruction.

4.3 Additive Corrections

A key difficulty in quantitative ^{90}Y Bremsstrahlung SPECT is that a substantial fraction of the measured counts arise from photons that are not well modelled by a purely “primary-only” forward projector. In this work we treat these contributions as *additive terms* in the Poisson measurement model, with a particular focus on (patient) scatter.

4.3.1 Statistical model with an additive background term

Let $y \in \mathbb{N}^I$ denote the measured projection data (all bins across all views) and let $x \in \mathbb{R}_+^J$ denote the unknown activity image. We use the standard Poisson model

$$y_i \sim \text{Poisson}((Ax)_i + b_i), \quad i = 1, \dots, I, \quad (4.1)$$

where A is the system matrix modelling the (primary) detection process (including attenuation and detector response as implemented in our reconstruction pipeline),

and $b \in \mathbb{R}_+^I$ is an *additive* component. In Bremsstrahlung SPECT, b is typically dominated by Compton-scattered photons, with additional contributions possible from other backgrounds (e.g. stray radiation or modelling mismatch). The role of additive correction is therefore to provide an estimate of b that is consistent with the acquisition protocol and reconstruction model.

Given a *fixed* b , OSEM yields an approximate maximum-likelihood estimate of x under this Poisson model. However, in ^{90}Y Bremsstrahlung imaging, b is large and difficult to estimate from energy-window tricks, motivating model-based strategies.

4.3.2 Initial scatter estimate via a pre-trained deep network

To initialise b we use a deep-learning scatter estimator that predicts scatter projections directly from the measured emission projections (and CT-derived attenuation information), following the general approach of training a convolutional neural network to regress MC-derived “true” scatter [67].

In our case this initialisation is explicitly heuristic: the available pre-trained network was developed and trained under a different hardware/protocol configuration (scanner, collimation, energy windowing, and number of views) than the Siemens and Mediso datasets used in this thesis [67]. In particular, the network was trained on data acquired on a Symbia Intevo SPECT/CT system with high-energy collimation and a 105–195 keV window with 80 projections, whereas our reconstructions were performed on Siemens or Mediso systems with different energy windows and either MEGP (Siemens) or MLEGP (Mediso) collimation. Despite this domain shift, we found that the predicted scatter term fits the measured data sufficiently well to provide a sensible starting point for subsequent model-based refinement (Figure 4.2).

4.3.3 Iterative OSEM–SIMIND refinement loop

Starting from the deep-learning initialisation $b^{(0)}$, we then refine the additive term using an alternating procedure:

1. **Reconstruction step:** perform an OSEM update of x using the current additive estimate $b^{(k)}$. We used 12 subsets and 5 epochs (total 60 subiterations)
2. **Regularisation of the intermediate image:** apply an isotropic Gaussian post-

smoothing of 5 mm full width at half maximum (FWHM) to the intermediate reconstruction (used only for the subsequent scatter simulation step).

3. **Scatter simulation step:** run a SIMIND-based Monte Carlo scatter simulation using the (smoothed) current activity estimate to produce a new scatter estimate $\hat{b}^{(k+1)}$.

We repeat this OSEM $\rightarrow$ SIMIND $\rightarrow$ OSEM+5 mm smoothing $\rightarrow \dots$ loop for 10 outer iterations. The 5 mm smoothing is used to suppress high-frequency noise that would otherwise be amplified by repeatedly re-simulating scatter from noisy intermediate reconstructions.

4.3.4 Offline versus online additive correction strategies

We consider two practical strategies for how the scatter term is used throughout the full reconstruction:

Offline (fixed) additive correction. In the *offline* strategy, we compute a scatter estimate once (after the initialisation/refinement loop) and then keep it fixed for the remainder of the reconstruction. Concretely, we set $b \leftarrow b^{(*)}$ and run OSEM to convergence (or for a fixed number of iterations) with a constant additive term. This approach is typically stable, but may be biased if the fixed scatter estimate is inconsistent with the final activity estimate.

Online (intermittently updated) additive correction. In the *online* strategy, we update the scatter estimate during reconstruction. That is, after some number of OSEM iterations, we re-run the SIMIND scatter simulation using the current image estimate and update b . This is conceptually aligned with hybrid Monte-Carlo/analytical reconstruction approaches in which the scatter estimate is updated a small number of times until it stabilises [63, 68]. In practice, the update cadence is a tunable choice: too frequent updates increase computational cost and can introduce instability; too infrequent updates reduce the benefit relative to the offline approach.

4.3.5 Full replacement versus damped scatter updates

For the online strategy, we investigate two update rules.

(1) **Full replacement.** We set

$$b^{(k+1)} \leftarrow \hat{b}^{(k+1)}. \quad (4.2)$$

Although appealingly direct, we observed that full replacement can be unstable in some cases, with the behaviour varying across phantoms. A practical interpretation is that once b is changed, the likelihood model being optimised by OSEM is also changed; therefore the overall alternating procedure is not guaranteed to behave monotonically. An

(2) **Damped (relaxed) updates.** To improve stability, we use a convex combination of the old and new scatter estimates:

$$b^{(k+1)} \leftarrow (1 - \alpha) b^{(k)} + \alpha \hat{b}^{(k+1)}, \quad \alpha \in (0, 1]. \quad (4.3)$$

This can be viewed as a relaxation of the fixed-point iteration $b = \hat{b}(x)$, reducing sensitivity to noise and modelling mismatch at the cost of slower adaptation. In all experiments here we set $\alpha = 0.2$. This value was chosen heuristically; we did not perform a dedicated study to optimise α , but empirically it provided a good stability–convergence trade-off across datasets.

Overall, these offline/online and replacement/damping variants provide a practical framework for controlling the bias–variance and stability properties of Bremsstrahlung additive correction, while keeping the reconstruction pipeline compatible with Monte Carlo scatter modelling.




Figure 4.2: Figures showing profile through projections for NEMA phantom. Summed over projections, axial slice is through NEMA spheres. Horizontal axis show tangential coordinate on gamma camera, Vertical axis shows counts. Top, replacing additive data; bottom, damped addition of additive data ($\alpha = 0.2$).

4.4 Bremsstrahlung Range Blur




Figure 4.3: Example fits of an isotropic Gaussian proxy to the Monte Carlo source-to-emission displacement distribution for ^{90}Y Bremsstrahlung in water (clinical energy windows).

In ^{90}Y Bremsstrahlung SPECT, detected photons are not emitted at the ^{90}Y decay site: Bremsstrahlung photons are generated along the β -particle track, inducing an additional *range blur* that is absent in photopeak imaging. Several MC frameworks can model this effect using detailed, energy-dependent probability density functions (PDFs) for the displacement between the decay location and the photon emission location. However, in practice, most general-purpose reconstruction toolkits are not designed to ingest arbitrary, precomputed displacement PDFs within the forward projector.

Here we adopt a pragmatic approximation: for clinically relevant energy windows and in water-equivalent tissue, the marginal displacement distribution is observed to be close to Gaussian. We therefore model the Bremsstrahlung range blur as an *isotropic image-space Gaussian* whose variance is selected to match MC-derived displacement statistics (Figure 4.3). This is simple to implement (a single 3D Gaussian smoothing) and is compatible with standard projector architectures.

4.4.1 Projector point-spread function

We model the collimator–detector response (CDR) by an approximately Gaussian point spread function (PSF) whose width increases with the source–collimator

distance h and depends on energy E . In the projector we represent the depth-dependent blur by a spatially variant isotropic Gaussian with standard deviation

$$\sigma_{\text{lin}}(h; E) = c_0(E) + c_1(E)h. \quad (4.4)$$

This linear model approximates the standard quadrature combination of intrinsic detector resolution and geometric collimator resolution,

$$\sigma_{\text{true}}(h; E) = \left[\sigma_{\text{intr}}(E)^2 + \sigma_{\text{geom}}(h; E)^2 \right]^{1/2}, \quad (4.5)$$

while avoiding repeated evaluation of a square root inside the projector. The coefficients $c_0(E)$ and $c_1(E)$ are obtained by a least-squares fit of $\sigma_{\text{lin}}(h_i; E)$ to $\sigma_{\text{true}}(h_i; E)$ over a clinically relevant distance range ($h \in [2, 20]$ cm). Writing

$$\bar{h} = \frac{1}{N} \sum_{i=1}^N h_i, \quad \bar{\sigma} = \frac{1}{N} \sum_{i=1}^N \sigma_{\text{true}}(h_i; E), \quad (4.6)$$

the closed-form solution (assuming $\sum_i (h_i - \bar{h})^2 > 0$) is

$$c_1(E) = \frac{\sum_{i=1}^N (h_i - \bar{h}) (\sigma_{\text{true}}(h_i; E) - \bar{\sigma})}{\sum_{i=1}^N (h_i - \bar{h})^2}, \quad c_0(E) = \bar{\sigma} - c_1(E) \bar{h}. \quad (4.7)$$

Intrinsic detector spatial response. Intrinsic resolution is measured using a standard source [20] and modelled as an isotropic Gaussian,

$$\text{PSF}_{\text{intr}}(\mathbf{r}; E) = \mathcal{N}(\mathbf{r}; \mathbf{0}, \sigma_{\text{intr}}^2(E) \mathbf{I}). \quad (4.8)$$

When intrinsic resolution is reported as $\text{FWHM}_{\text{intr}}(E)$, we convert using

$$\sigma_{\text{intr}}(E) = \frac{\text{FWHM}_{\text{intr}}(E)}{2\sqrt{2\ln 2}}. \quad (4.9)$$

Geometric collimator response with penetration via effective length. For a parallel-hole collimator, the (FWHM) geometric resolution at distance h is commonly

approximated by

$$\text{FWHM}_{\text{geom}}(h; E) \approx d_{\text{eff}} \left(1 + \frac{h}{l_{\text{eff}}(E)} \right), \quad (4.10)$$

where d_{eff} is the effective hole diameter and $l_{\text{eff}}(E)$ is an *effective* hole length used to account for septal penetration. We use

$$l_{\text{eff}}(E) = l - \frac{2}{\mu_{\text{Pb}}(E)}, \quad \mu_{\text{Pb}}(E) \text{ interpolated from NIST XCOM tables,} \quad (4.11)$$

where l is the physical hole length and $\mu_{\text{Pb}}(E)$ is the linear attenuation coefficient of lead at energy E [69]. Converting from FWHM to standard deviation gives

$$\sigma_{\text{geom}}(h; E) = \frac{\text{FWHM}_{\text{geom}}(h; E)}{2\sqrt{2 \ln 2}}. \quad (4.12)$$

4.4.2 Bremsstrahlung-specific blur

Gaussian proxy for range blur. Let Δ denote the displacement between the ^{90}Y decay location and the Bremsstrahlung photon emission location. Rather than using a detailed displacement PDF inside the projector, we approximate the induced range blur by an isotropic Gaussian kernel

$$\text{PSF}_{\text{brem}}(\mathbf{r}; E) = \mathcal{N}(\mathbf{r}; \mathbf{0}, \sigma_{\text{brem}}^2(E) \mathbf{I}). \quad (4.13)$$

For each clinical energy window, $\sigma_{\text{brem}}(E)$ is obtained from a coarse look-up table (LUT) by fitting a zero-mean Gaussian proxy to MC displacement curves for ^{90}Y Bremsstrahlung in water from GATE fast-simulation studies [18] (Figure 4.3).

Composite energy windows. For a set of energy windows $\mathcal{S}$ combined into a single projection dataset, we collapse the corresponding Gaussian mixture by second-moment matching:

$$\sigma_{\text{clin}}^2 = \frac{\sum_{i \in \mathcal{S}} w_i \sigma_i^2}{\sum_{i \in \mathcal{S}} w_i}, \quad w_i \propto \text{per-window counts.} \quad (4.14)$$

Implementation in image space. Because the convolution of isotropic Gaussians is Gaussian, incorporating Bremsstrahlung range blur is equivalent to adding variances. In practice we implement this as an image-space convolution by PSF_{brem} (i.e. a 3D Gaussian smoothing) applied to the current activity estimate before forward projection. This keeps the forward model compatible with standard projector implementations while capturing the dominant range-broadening effect in water-equivalent tissue.

4.4.3 Validation against SIMIND and GATE forward projections

We assessed how well the isotropic Gaussian proxy reproduces the broader displacement PDFs (including angular effects in GATE) by comparing projected point-source responses produced by:

1. STIR forward projections with geometric CDR modelling only;
2. STIR forward projections with geometric CDR plus the proposed isotropic Bremsstrahlung range blur;
3. SIMIND simulations (which use an approximate isotropic Bremsstrahlung emission model); and
4. GATE simulations (used here as a higher-fidelity reference that includes angular dependence of Bremsstrahlung emission).

Point-source-in-sphere configuration. A true ^{90}Y point source is not physically realisable at clinical activity levels because the β -particles have a finite range. To approximate a point source for spatial response measurements, we embedded the activity within a small water-equivalent sphere of radius 1.2 cm, consistent with the maximum β -particle range in soft tissue/water (approximately 11 mm) [70]. For each configuration we measured the detector-plane FWHM at source–collimator distances of 2, 5, 10, 15, and 20 cm.

Clinical energy windows. Comparisons were performed for three representative clinical windows:

$$[38.24, 126.16] \text{ keV}, \quad [50.0, 150.0] \text{ keV}, \quad [75.0, 225.0] \text{ keV}. \quad (4.15)$$

Scope and limitations. This validation focuses on water-equivalent media, which is appropriate for the liver (the primary application in post-radioembolization imaging). Differences in range and scatter properties in lung or cortical bone are not explicitly modelled here and are not expected to dominate performance for liver-focused dosimetry.

4.5 Residual Correction

Fu and Qi proposed a residual-correction (RC) strategy to accelerate reconstructions when an *accurate* forward model is too expensive to apply at every iteration [66]. In their positron-range setting, the accurate system model can be expressed as a fast projector composed with an additional blurring operator. RC then intermittently estimates the modelling discrepancy and folds it into an additive term, allowing subsequent iterations to proceed using the cheaper projector.

In this work we investigate whether RC can be applied analogously to ^{90}Y Bremsstrahlung SPECT reconstruction. Empirically, we find that RC behaves well for high-count, low-background phantom data, but can become unreliable in low-count clinical-like regimes where regularisation is necessarily strong. We attribute this behaviour to the fact that RC corrects the *forward prediction* but does not, in general, restore a *matched* forward/adjoint pair inside a regularised reconstruction.

4.5.1 Residual correction update

Let $A_\star$ denote the fast (computationally cheap) system matrix used in the iterative updates and let A_{MC} denote a more accurate (expensive) system model (here, Monte Carlo based). We consider the Poisson model

$$y_i \sim \text{Poisson}((A_\star x)_i + b_{E,i}), \quad (4.16)$$

where b_E is an additive correction term (e.g. scatter). At spaced iterations, we estimate the discrepancy between the accurate and fast forward predictions at the current iterate by approximately solving

$$\hat{x} \approx \arg \min_{x \geq 0} \mathcal{O}(x \mid y; b_E, A_\star), \quad (4.17)$$

and then forming the residual

$$(\Delta A)\hat{x} = A_{\text{MC}}\hat{x} - A_\star\hat{x}. \quad (4.18)$$

This residual is incorporated into the additive term,

$$b_E^{(\text{new})} = b_E^{(\text{old})} + (\Delta A)\hat{x}, \quad (4.19)$$

after which reconstruction proceeds again under the updated additive correction. In favourable regimes, this produces forward predictions closer to those of A_{MC} while retaining the per-iteration cost of $A_\star$.

4.5.2 Regularised objective with a Relative Difference prior

In Bremsstrahlung SPECT, low-count data frequently necessitate comparatively strong regularisation. In all RC experiments reported here, we reconstruct using a Relative Difference prior (RDP), i.e. we solve a MAP problem of the form

$$\hat{x} = \arg \min_{x \geq 0} \{ \Phi_{A_\star}(x) + \beta R_{\text{RDP}}(x) \}, \quad (4.20)$$

where $\Phi_{A_\star}(x)$ is the negative Poisson log-likelihood (up to constants),

$$\Phi_{A_\star}(x) = \sum_{i=1}^I [(A_\star x)_i + b_{E,i} - y_i \log((A_\star x)_i + b_{E,i})], \quad (4.21)$$

and the Relative Difference prior is defined over a neighbourhood system $\mathcal{N}$ by

$$R_{\text{RDP}}(x) = \sum_{(j,k) \in \mathcal{N}} w_{jk} \frac{(x_j - x_k)^2}{(x_j + x_k)^2 + \eta^2}. \quad (4.22)$$

Here, x_j and x_k denote neighbouring voxels, $w_{jk} \geq 0$ are neighbourhood weights, and $\eta > 0$ is a small smoothing constant that prevents numerical instability when $x_j + x_k$ is small. The hyperparameter $\beta > 0$ controls the overall strength of regularisation.

4.5.3 Interpretation via matched versus unmatched projectors

To understand why RC can be fragile in low-count, strongly regularised settings, it is helpful to view the method through the lens of *matched* versus *unmatched* projector pairs.

Set-up. Suppose the accurate forward model can be written as a fast projector composed with an additional blur (or resolution) operator K ,

$$P := A_{\text{fast}}K, \quad \bar{y}(x) = Px + r = A_{\text{fast}}Kx + r, \quad (4.23)$$

with r an additive background term and Poisson data $y \sim \text{Poisson}(\bar{y}(x))$.

Maximum likelihood (no regularisation). The negative log-likelihood for forward model P has gradient

$$\nabla \Phi_P(x) = P^T \left(1 - \frac{y}{Px + r} \right). \quad (4.24)$$

For $P = A_{\text{fast}}K$, the *matched* stationarity condition is

$$K^T A_{\text{fast}}^T \left(1 - \frac{y}{A_{\text{fast}}Kx + r} \right) = 0. \quad (4.25)$$

If instead one uses the fast backprojector A_{fast}^T with the same forward model (an *unmatched* pair), the stationarity condition would be

$$A_{\text{fast}}^T \left(1 - \frac{y}{A_{\text{fast}}Kx + r} \right) = 0. \quad (4.26)$$

In the special case where K^T is injective (and numerically well-conditioned on the relevant subspace), (4.25) and (4.26) share the same stationary points, since $K^T z = 0 \Rightarrow z = 0$. This explains why unmatched backprojection can appear benign in some unregularised, high-count regimes.

How RC fits. Residual correction aims to make the *forward prediction* behave like the accurate model without applying K at every iteration, i.e.

$$A_{\text{fast}}Kx \approx A_{\text{fast}}x + \text{RC}(x), \quad (4.27)$$

where $\text{RC}(x)$ approximates $(A_{\text{fast}}K - A_{\text{fast}})x$ (often using lagged/intermittent updates). If the correction is consistent at a limit point x^* , then the ratio term $y/(A_{\text{fast}}Kx + r)$ is effectively “accurate-forward”. However, the backprojection used by the iterative update typically remains A_{fast}^T , not $(A_{\text{fast}}K)^T = K^T A_{\text{fast}}^T$. In this sense, RC can be interpreted as correcting the forward model while still behaving like an *unmatched* method on the adjoint side.

Why low-count + RDP regularisation is high-risk. With an RDP prior, the MAP stationarity condition (ignoring the non-negativity constraint) is

$$P^T \left(1 - \frac{y}{Px + r} \right) + \beta \nabla R_{\text{RDP}}(x) = 0. \quad (4.28)$$

For $P = A_{\text{fast}}K$, the matched condition becomes

$$K^T A_{\text{fast}}^T \left(1 - \frac{y}{A_{\text{fast}}Kx + r} \right) + \beta \nabla R_{\text{RDP}}(x) = 0, \quad (4.29)$$

whereas an unmatched analogue would be

$$A_{\text{fast}}^T \left(1 - \frac{y}{A_{\text{fast}}Kx + r} \right) + \beta \nabla R_{\text{RDP}}(x) = 0. \quad (4.30)$$

As total counts decrease, the likelihood term becomes flatter and noisier, so the relative influence of $\beta \nabla R_{\text{RDP}}(x)$ increases even when β is held fixed. Unlike the ML case, the factor K^T multiplies only the data-term gradient in (4.29) and does not act on $\nabla R_{\text{RDP}}(x)$. Therefore, for $\beta > 0$, matched and unmatched fixed points generally differ. In this prior-dominated regime, using an unmatched backprojector while correcting only the forward prediction can substantially alter the balance between data fidelity and regularisation, leading to marked divergence from the solution associated with the fully matched accurate model.

4.5.4 Empirical observations

For high-count, low-background phantom acquisitions (e.g. NEMA-style datasets), RC substantially improves agreement with accurate forward projections and remains stable. For low-count Bremsstrahlung acquisitions, however, we observe that RC can become unreliable when combined with strong RDP regularisation. These observations are consistent with the interpretation above: in prior-dominated regimes, mismatched adjoints change the effective weighting of the data term relative to the RDP, and RC does not in general restore a matched forward/adjoint pair.

4.6 Putting It All Together

Chapter 5

Anatomy-informed deconvolution for PET partial volume correction

5.1 Introduction and relationship to the thesis

Chapter 3 established the penalised-likelihood framework for emission tomography, reviewed representative methods for injecting prior knowledge into the optimisation problem (including anatomically guided and kernelised models), and motivated optimisation strategies that remain stable under regularisation. The core of this thesis concerns regularised and synergistic reconstruction under the full Poisson forward model $\bar{y}(x) = Ax + b$, where the ill-conditioning of A and the low-count regime make prior design, preconditioning, and choice of optimisation algorithm critical.

In this chapter we consider a deliberately simplified, image-domain inverse problem: post-reconstruction partial volume correction (PVC) via deconvolution with an estimated point spread function (PSF). Although this is not a substitute for incorporating resolution modelling within reconstruction, it provides a controlled sandbox in which the effects of anatomical guidance and regularisation can be studied without confounding factors such as subset noise, projector modelling error, and complex system geometry. In particular, the chapter serves two main purposes within the broader narrative:

1. it connects the priors reviewed in Chapter 3 (kernel methods, directional regularisation) to a tractable inverse problem where their qualitative behaviour

is easily interpreted;

2. it provides a practical PVC baseline that is straightforward to deploy on retrospectively reconstructed images.

Richardson–Lucy (RL) deconvolution is a purely image-domain analogue of MLEM under the Poisson model $\bar{y}(x) = Cx + b$: it is the same EM update, with the system matrix A replaced by the PSF convolution operator C [25, 71]. In this setting the “sensitivity image” simplifies substantially, since for a unit-sum PSF one has $C^\top \mathbf{1} \approx \mathbf{1}$, so the usual MLEM normalisation by $A^\top \mathbf{1}$ becomes approximately trivial. We extend Richardson–Lucy (RL) to a penalised deconvolution regime by incorporating a directional total variation prior (Equation 3.25), and we construct kernelised deconvolution updates by substituting the kernel image model $x = K\alpha$, exactly as in KEM/HKEM. The resulting anatomy-aware deconvolution variants are referred to throughout as Kernelised Richardson–Lucy (KRL), Hybrid Kernelised Richardson–Lucy (HKRL), and Richardson–Lucy with Directional Total Variation (RL–DTV). Their basic update equations are derived in the following sections.

5.2 Image-domain inverse problem and objective

Let $x \in \mathbb{R}_+^J$ denote the unknown (sharper) PET activity image on a voxel grid, and let $y \in \mathbb{R}_+^J$ denote an observed blurred image. Although reconstructed PET images are not independent raw count data, their local noise is often treated as approximately Poisson in post-reconstruction restoration models [71]. We therefore use the following image-domain approximation:

$$y_j \sim \text{Poisson}(\bar{y}_j(x)), \quad \bar{y}(x) = Cx + b, \quad j = 1, \dots, J, \quad (5.1)$$

where $C \in \mathbb{R}_+^{J \times J}$ is a discretised convolution operator representing a space-invariant PSF, and $b \in \mathbb{R}_+^J$ is an optional additive background term (set to 0 in the experiments unless stated otherwise). The operator C plays the same role as the system matrix A in (3.3), but in the image domain.

Within this approximation, assuming independence across voxels in the image-

domain observation model, the Poisson negative log-likelihood is

$$\mathcal{L}(x) = \sum_{j=1}^J (\bar{y}_j(x) - y_j \log \bar{y}_j(x)), \quad \bar{y}(x) = Cx + b, \quad (5.2)$$

dropping constants independent of x as in Chapter 3.

A MAP (penalised-likelihood) deconvolution estimate is then

$$\hat{x} \in \arg \min_{x \in \mathcal{C}} \{\mathcal{L}(x) + \beta \mathcal{R}(x)\}, \quad \mathcal{C} = \mathbb{R}_+^J, \quad (5.3)$$

with $\beta > 0$ and an appropriate regulariser $\mathcal{R}$.

Throughout, C is non-negative and (approximately) intensity preserving: for a unit-sum PSF kernel and negligible boundary effects, $C\mathbf{1} \approx \mathbf{1}$ and $C^\top \mathbf{1} \approx \mathbf{1}$. The adjoint $C^\top$ corresponds to convolution with the flipped PSF. For a symmetric PSFs common in PET, $C^\top \approx C$.

5.3 Baseline: Richardson–Lucy deconvolution

The RL algorithm is the EM algorithm for the Poisson model (5.1) with $b = 0$ [25]. Introducing elementwise multiplication and division, denoted $\odot$ and $\oslash$, the unit-sensitivity RL update is

$$x^{(k+1)} = x^{(k)} \odot C^\top \left(y \oslash (Cx^{(k)}) \right). \quad (5.4)$$

More generally, the right-hand side of (5.4) contains an elementwise division by $C^\top \mathbf{1}$; in this chapter this denominator is neglected because the PSF kernels are normalised and boundary effects are small. As with MLEM/OSEM under the full forward model, RL iterations tend to sharpen edges rapidly but amplify high-frequency noise as iterations increase. Consequently, unregularised RL is typically used with early stopping, which acts as an implicit regulariser but yields iteration-dependent bias.

5.3.1 Scaled-gradient interpretation and EM-type preconditioning

To connect with Chapter 3, we rewrite (5.4) as an additive scaled-gradient step. The gradient of the Poisson negative log-likelihood (5.2) (with $b = 0$) is

$$D_x \mathcal{L}(x) = C^\top (\mathbf{1} - y \odot (Cx)). \quad (5.5)$$

Using the same rearrangement as for MLEM in Chapter 3, and again using $C^\top \mathbf{1} \approx \mathbf{1}$, we obtain

$$x^{(k+1)} = x^{(k)} \odot \left(\mathbf{1} - D_x \mathcal{L}(x^{(k)}) \right). \quad (5.6)$$

This corresponds to an EM-type diagonal preconditioner analogous to $P_{EM}^{(k)} = \text{diag}(x^{(k)} / (A^\top \mathbf{1}))$ in the full reconstruction problem.

5.4 Kernelised Richardson–Lucy deconvolution

Kernel methods were reviewed in Chapter 3 as an implicit regularisation strategy in which prior information enters through the image model $x = K\alpha$ rather than through an explicit penalty [32, 33]. The same idea can be applied to image-domain deconvolution.

5.4.1 Kernel image model

Let u be a co-registered anatomical guidance image (magnetic resonance (MR) or CT), resampled to the PET grid. We represent

$$x = K\alpha, \quad \alpha \in \mathbb{R}_+^J, \quad K \in \mathbb{R}_+^{J \times J}, \quad (5.7)$$

where K is a sparse, row-normalised kernel matrix constructed from features extracted from u (and optionally spatial coordinates), as in Chapter 3. Row normalisation ($K\mathbf{1} = \mathbf{1}$) ensures that constant coefficient images are mapped to constant activity images, preserving the intended intensity scale of the kernel representation.

Substituting into (5.1) gives the kernelised forward model

$$\bar{y}(\alpha) = C(K\alpha) + b. \quad (5.8)$$

5.4.2 KRL update

Maximising the Poisson likelihood with respect to α under (5.8) yields an EM-type multiplicative update analogous to KEM:

$$\alpha^{(k+1)} = \alpha^{(k)} \odot \frac{K^T C^T (y \oslash (CK\alpha^{(k)} + b))}{K^T C^T \mathbf{1}}, \quad x^{(k+1)} = K\alpha^{(k+1)}. \quad (5.9)$$

As in KEM, the denominator can be precomputed once, and the update preserves non-negativity given $\alpha^{(0)} \geq 0$. Under the same unit-sensitivity approximation used for vanilla RL, $C^T \mathbf{1} \approx \mathbf{1}$, the denominator reduces to $K^T \mathbf{1}$; this term is generally not equal to $\mathbf{1}$ and cannot be ignored.

5.4.3 Hybrid KRL (HKRL)

A known limitation of purely anatomy-derived kernels is suppression of PET-unique structures when they are not supported by the guidance image [44]. Hybrid kernels mitigate this by incorporating features derived from the evolving PET estimate [44].

Because iteration-dependent kernels can propagate noise back into the coupling structure at late iterations, we adopt a frozen-kernel strategy consistent with the motivation in HKEM, freezing after 5 iterations (although this proved to be unnecessary for the sphere phantom as the best iteration was usually before this point).

5.5 Directional-TV regularised RL deconvolution

directional total variation (DTV) was introduced in Chapter 3 as an anatomically guided regulariser that promotes edge alignment between the emission image and a structural guidance image [43]. Here we use DTV as an explicit penalty for deconvolution.

5.5.1 Penalised objective and update

The penalised-likelihood deconvolution objective is

$$\hat{x} \in \arg \min_{x \in \mathbb{R}_+^J} \{ \mathcal{L}(x) + \beta R_{\text{sDTV}}(x; u) \}, \quad \mathcal{L}(x) = \sum_j ((Cx + b)_j - y_j \log(Cx + b)_j). \quad (5.10)$$

Here R_{sDTV} is a smoothed DTV penalty. A natural baseline solver is the scaled-gradient RL–DTV iteration

$$x^{(k+1)} = \left\{ x^{(k)} \odot \left(\mathbf{1} - \varsigma_k \left[D_x \mathcal{L}(x^{(k)}) + \beta \nabla_x R_{\text{sDTV}}(x^{(k)}; u) \right] \right) \right\}_{\geq 0}, \quad (5.11)$$

where $\{\}_{\geq 0}$ denotes projection onto non-negativity. Step sizes $\{\varsigma_k\}$ may be taken constant (with conservative tuning) or diminishing to improve stability. This update is the image-domain analogue of the scaled-gradient MAP schemes discussed in Chapter 3 (compare with BSREM-type forms), but without subsets.

In contrast to kernelised RL variants, RL–DTV is typically numerically stable once β is chosen, and the objective in (5.10) is smooth with simple bound constraints. This makes it well suited to a quasi-Newton solver: in the experiments of this chapter we therefore accelerate convergence by optimising (5.10) with the L-BFGS-B algorithm reviewed in Section 3.3, using bound constraints to enforce non-negativity. In this setting, L-BFGS-B replaces the hand-tuned step-size sequence $\{\varsigma_k\}$ by an automatic line search and curvature-informed scaling, often reaching a stationary point of (5.10) in substantially fewer iterations than the scaled-gradient RL–DTV update (5.11).

For RL and KRL/HKRL, early stopping is essential and is treated as part of the regularisation. In experiments we report results across iteration counts and select a stopping point based on NRMSE (Equation 5.15) for the phantom and qualitative assessment for the brain example. For RL–DTV, iteration sensitivity is reduced, but a stopping point is still selected to avoid unnecessary computation once the relative

objective decrease satisfies

$$\frac{\mathcal{O}(x^{(k)}) - \mathcal{O}(x^{(k+1)})}{\max\{|\mathcal{O}(x^{(k)})|, |\mathcal{O}(x^{(k+1)})|, 1\}} \leq f_{\text{tol}}, \quad \mathcal{O}(x) = \mathcal{L}(x) + \beta R_{\text{sDTV}}(x; u). \quad (5.12)$$

Here f_{tol} is chosen to be 10^{-6} . In L-BFGS-B implementations that expose the historical `factr` parameter, the equivalent objective tolerance is $f_{\text{tol}} = \text{factr} \varepsilon_{\text{mp}}$, where ε_{mp} is machine precision.

5.6 Data and experimental setup

Two datasets are used to characterise behaviour under controlled and realistic conditions.

1. **Spherical phantom:** a hot-sphere phantom enabling assessment of edge recovery and spill-over (see Figure 5.1).
2. **τ -PET brain with MRI guidance:** a representative in vivo dataset (see Figure 5.2).




Figure 5.1: Left, ground truth phantom, centre blurred and noisy reconstruction and right MRI guidance image for the simulated sphere phantom. Units are arbitrary.

In both cases, the anatomical image u is registered to the PET grid and resampled to the reconstruction voxel size. The registered anatomy is used either to define the kernel K (KRL/HKRL) or the directional field $\{\xi_j\}$ (RL-DTV).

To obtain a controlled test case with known truth while retaining a realistic resolution/noise regime, the sphere dataset is simulated from a digital activity distribution by forward projecting with an isotropic Gaussian resolution model of




Figure 5.2: Left, reconstructed τ -PET brain, right registered MRI guidance image. Units are arbitrary.

FWHM = 6 mm embedded in the forward projector, adding Poisson noise to the resulting sinograms, and then reconstructing using an operator without additional blurring. This deliberately introduces a resolution-modelling mismatch between simulation and reconstruction. The deconvolution operator used in the subsequent image-domain correction was the same 6 mm Gaussian model, implemented with the same boundary handling as the simulation blur.

Let x^* denote the digital ground-truth activity image and $\hat{x}$ the reconstructed image. For each sphere s , let $\mathcal{S}_s$ be the set of voxel indices belonging to that sphere, and let $\mathcal{S} = \bigcup_s \mathcal{S}_s$ denote the union over all spheres. We report an aggregate recovery coefficient computed by pooling all sphere voxels,

$$\text{RC} = \frac{\frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} \hat{x}_i}{\frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} x_i^*}. \quad (5.13)$$

To quantify background noise, let $\mathcal{B}$ denote the set of voxel indices in a warm background region. We compute the coefficient of variation (CoV) within $\mathcal{B}$ as

$$\text{CoV}_{\mathcal{B}} = \frac{\sigma_{\mathcal{B}}}{\mu_{\mathcal{B}}}, \quad \mu_{\mathcal{B}} = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \hat{x}_i, \quad \sigma_{\mathcal{B}} = \sqrt{\frac{1}{|\mathcal{B}| - 1} \sum_{i \in \mathcal{B}} (\hat{x}_i - \mu_{\mathcal{B}})^2}. \quad (5.14)$$

Finally, overall fidelity to the truth is summarised by the normalised root mean

squared error (NRMSE) over the full image grid Ω ,

$$\text{NRMSE} = \frac{\sqrt{\frac{1}{|\Omega|} \sum_{i \in \Omega} (\hat{x}_i - x_i^*)^2}}{\sqrt{\frac{1}{|\Omega|} \sum_{i \in \Omega} (x_i^*)^2}}. \quad (5.15)$$

The in vivo dataset is a τ -PET scan acquired on a healthy volunteer using the tracer ^{18}F -MK-6240, with a co-registered anatomical MRI serving as guidance. This dataset is used to qualitatively illustrate how the different priors behave in the presence of heterogeneous uptake patterns and fine anatomical structure, where over-regularisation can lead to anatomical “imprinting”. The FWHM of the deconvolution operator was estimated from a Hoffman brain phantom scan acquired on the same Siemens Vision scanner. The phantom reconstruction used PSF modelling; after removing extra-cerebral activity, the reconstructed image was co-registered to a Hoffman template, and the template was smoothed over a grid of Gaussian FWHM values, independently in the in-plane and axial directions. The smoothing that minimised the sum-of-squared differences to the reconstructed phantom image was selected, giving $\text{FWHM} = (4.6 \text{ mm}, 4.6 \text{ mm}, 6.4 \text{ mm})$.

5.6.1 Hyperparameter tuning

Hyperparameters were tuned by grid search.

For DTV reconstructions, we swept the regularisation weight

$$\beta \in \{0.01, 0.02, 0.05, 0.1, 0.2, 0.5, 1.0, 2.0, 5.0\}.$$

Kernel methods used a fixed number of neighbours (48-NN) within a cubic $9 \times 9 \times 9$ neighbourhood [33]. Feature channels were normalised before computing Gaussian similarities, as in the kernel construction described in Chapter 3; the reported σ_a and σ_f values are therefore interpreted in these normalised feature units. For KRL we swept the anatomical weighting σ_a . For HKRL we additionally swept the functional weighting σ_f . Because we used a k-NN neighbourhood strategy, distance weights (σ_{df}, σ_{da}) were not used.

The requirement for early stopping in RL, KRL and HKRL means that the

Table 5.1: Hyperparameter sweeps for kernel-based deconvolution

Method	Hyperparameter	Values swept
KRL/HKRL	σ_a	{0.1, 0.2, 0.5, 1.0, 2.0, 5.0}
HKRL	σ_f	{0.1, 0.2, 0.5, 1.0, 2.0, 5.0}

initialising image plays a pivotal role in regularisation. We investigate the effect that initialisation could have on image quality by running identical sweeps over the sphere dataset initialising with a uniform image of ones and the blurred PET image.

τ -PET brain data reconstructions were initialised using a uniform image of ones for RL/KRL/HKRL and the blurred PET image for RL-DTV.

5.7 Results

5.7.1 Spheres

The results for the best reconstructions are summarised in Table 5.2 and visualised in Figure 5.3.




Figure 5.3: Figures showing central slice of reconstructions with best NRMSE. Top, initialised by the blurred PET image and bottom, initialised by a uniform image.

Across both initialisation strategies, RL-DTV achieved the lowest NRMSE (0.0177), the highest aggregate recovery coefficient (approximately 0.889), and the lowest background CoV (approximately 0.123). KRL and HKRL substantially

Uniform initialisation							
Method	Min NRMSE	Iter.	RC	Bkg CoV	σ_a	σ_f	β
RL	0.0327	2	0.7645	0.2106	–	–	–
KRL	0.0190	3	0.8781	0.1546	0.5	–	–
HKRL	0.0191	3	0.8776	0.1546	0.5	5	–
RL-DTV	0.0177	143	0.8890	0.1233	–	–	2
Blurred initialisation							
Method	Min NRMSE	Iter.	RC	Bkg CoV	σ_a	σ_f	β
RL	0.0338	0	0.7307	0.2278	–	–	–
KRL	0.0208	2	0.8841	0.1805	5	–	–
HKRL	0.0212	2	0.8805	0.1805	5	5	–
RL-DTV	0.0177	164	0.8893	0.1233	–	–	2

Table 5.2: Tables summarising the hyperparameters and metric values of the best (measured by NRMSE) reconstructions of the various methods for (top) a uniform initialisation of ones and (bottom) the blurred PET image.

improved on unregularised RL, but were almost indistinguishable from one another in this simple sphere phantom. With blurred initialisation, the best RL result occurred at iteration 0, indicating that unconstrained RL deconvolution immediately increased the NRMSE once applied to this noisy reconstruction.

The effect of initialisation on the quality of reconstructions can be seen in Figures 5.4 and 5.5.

5.7.2 τ -PET brain

τ -PET brain data were only qualitatively analysed. Optimal images were chosen by visual quality, which is an obvious limitation of the study design. The chosen images are shown in Figure 5.6.

In the brain example, RL produced sharp but visibly noisy cortical structure, while KRL and HKRL suppressed noise but showed stronger dependence on the anatomical guidance image. RL-DTV gave the best qualitative compromise in this dataset: cortical definition was improved relative to the blurred image and the kernel methods, while noise amplification was lower than for unregularised RL. This visual assessment supports the quantitative sphere result, but should be interpreted




Figure 5.4: Figures showing iteration versus global NRMSE for the best performing reconstructions. Top, initialised by the blurred PET image and bottom, initialised by a uniform image.

cautiously because it is based on a single healthy volunteer scan and does not include a reference truth.

5.8 Discussion

Deblurring is inherently ill-posed: the PSF attenuates high frequencies, so naive inversion magnifies noise. Anatomical guidance provides a structured way to preferentially preserve edges that are plausible given u while damping noise-like oscillations, either by restricting the solution space (kernel methods) or by penalising anatomically inconsistent gradients (DTV).

All guidance-based methods risk bias when the guidance image does not reflect tracer uptake boundaries. KRL is the most restrictive because it enforces $x \in \text{range}(K)$; RL-DTV is less restrictive but can still encourage alignment where the data term is weak (low counts, low contrast). HKRL partially mitigates mismatch




Figure 5.5: Figures showing RC versus CoV for all reconstructions. Top, initialised by the blurred PET image and bottom, initialised by a uniform image. KRL and HKRL are visually indistinguishable in these plots.




Figure 5.6: Figures showing visually best reconstructions for the different reconstructions.

by allowing the coupling structure to adapt to PET features, but must be controlled (e.g. freezing) to avoid noise-driven feedback.

The behaviours observed here anticipate challenges in the full reconstruction setting: strong priors induce spatially varying curvature and require effective scaling (preconditioning), and guidance mechanisms must be designed to avoid suppressing modality-unique information. These considerations directly motivate the synergy-aware prior and optimisation strategy developed later in the thesis.

The present study should therefore be viewed as a proof-of-principle rather than a definitive validation. A stronger evaluation would require a larger set of patient scans and more demanding controlled simulations, for example BrainWeb-style digital brain phantoms with inserted PET-unique lesions or tumours. Such work is ongoing, supported by additional clinical datasets, and would allow the qualitative observations here to be tested with population-level variability and known ground truth.

5.9 Summary

This chapter formulated post-reconstruction PET PVC as an approximate image-domain Poisson inverse problem $\bar{y}(x) = Cx + b$ and studied RL deconvolution and three anatomy-aware regularised variants. KRL and HKRL implement implicit regularisation through kernelised image models, consistent with the kernel framework reviewed in Chapter 3 [32, 33, 44]. RL-DTV implements explicit anatomically guided regularisation via a DTV penalty [43] using the same directional notation as Chapter 3. Together, these methods provide a practical PVC baseline and a controlled setting in which to diagnose guidance-driven bias and noise amplification, informing the design choices for the synergistic priors and optimisation algorithms developed in subsequent chapters.

Chapter 6

Introduction to Weighted Directional Total Nuclear Variation

In this chapter, we introduce a smoothed weighted directional total nuclear variation (w -dTNV) prior for synergistic emission tomography. The chapter focuses on the regulariser itself: its relationship to total nuclear variation, its differentiability after singular-value smoothing, the use of anatomical information through directional operators, and the modality-wise weighting needed when coupling heterogeneous emission images.

The optimisation machinery used to apply this prior in large-scale PET/SPECT reconstruction, including prior-aware preconditioning, subset selection, and multi-bed PET handling, is developed separately in Chapter 7.

6.1 Total Nuclear Variation Revisited

This section builds on the definitions in Section 3.2.3. For convenience, we restate the key definitions here:

- Registered images: $\mathbf{z} := \{z_m\}_{m=1}^M$, where $z_m := W_m x_m$ and $W_m : \mathbb{R}_+^{J_m} \rightarrow \mathbb{R}_+^N$.
- Finite-difference operator: for $\ell = 1, \dots, d$,

$$\left((Dx)_n\right)^{(\ell)} := \frac{x_{n+\delta_\ell} - x_n}{\Delta_\ell}.$$

- Voxel-wise Jacobian:

$$\mathcal{G}_n(\mathbf{z}) := \begin{bmatrix} Dz_{1,n} & Dz_{2,n} & \cdots & Dz_{M,n} \end{bmatrix} \in \mathbb{R}^{d \times M}.$$

- Nuclear norm: for a matrix Y with singular values $\{\sigma_k(Y)\}_{k=1}^r$ and $r = \text{rank}(Y)$,

$$\|Y\|_* := \sum_{k=1}^r \sigma_k(Y).$$

- Total nuclear variation:

$$\mathcal{R}_{\text{TNV}}(\mathbf{z}) := \sum_{n=1}^N \|\mathcal{G}_n(\mathbf{z})\|_*.$$

6.1.1 Mathematical properties and differentiability

Briefly revisiting total variation, optimisation of $\mathcal{R}_{\text{TV}}$ (cf. (3.17)) is inconvenient in gradient-based methods because these functionals are not globally smooth: the TV gradient is undefined when $(Dx_m)_j = 0$. A standard remedy is to introduce a small smoothing parameter $\varepsilon > 0$.

If we define

$$\tilde{\mathcal{R}}_{\text{TV}}(x_m) := \sum_{j=1}^{J_m} \|(Dx_m)_j\|_{2,\varepsilon}, \quad \|v\|_{2,\varepsilon} := \sqrt{\|v\|_2^2 + \varepsilon^2}, \quad (6.1)$$

then the image-space gradient is

$$\nabla_{x_m} \tilde{\mathcal{R}}_{\text{TV}}(x_m) = -\text{div} \left(\frac{Dx_m}{\|Dx_m\|_{2,\varepsilon}} \right), \quad (6.2)$$

where the division is understood pointwise over voxels.

For nuclear-norm VTV, the unsmoothed singular-value map is not differentiable at zero singular values, and hence the corresponding nuclear norm is not C^1 at rank-deficient points. We therefore apply the Charbonnier function to the singular

values:

$$g(s) = \sqrt{s^2 + \varepsilon^2} - \varepsilon, \quad g'(s) = \frac{s}{\sqrt{s^2 + \varepsilon^2}}, \quad g''(s) = \frac{\varepsilon^2}{(s^2 + \varepsilon^2)^{3/2}}. \quad (6.3)$$

The constant $-\varepsilon$ matches the implemented objective value and does not affect gradients or optimisation.

The nuclear norm is orthogonally invariant: for any orthogonal $U \in \mathbb{R}^{d \times d}$ and $V \in \mathbb{R}^{M \times M}$,

$$\|Y\|_* = \|UYV^\top\|_*. \quad (6.4)$$

More generally, any function of the form

$$\phi(Y) := f(\sigma(Y)), \quad f(s) = \sum_{\ell=1}^r g(s_\ell), \quad r = \min\{d, M\}, \quad (6.5)$$

is orthogonally invariant and depends on Y only through its singular values $\sigma(Y) = (\sigma_1(Y), \dots, \sigma_r(Y))$. Since $g \in C^\infty([0, \infty))$ for $\varepsilon > 0$, the symmetric function f is smooth and hence ϕ is differentiable, with gradient given by the standard formula for orthogonally invariant matrix functions [72, 73].

Jacobian-space gradient. Let $\phi(\mathcal{G}) = \sum_{\ell=1}^{\min(d, M)} g(\sigma_\ell(\mathcal{G}))$, where $\sigma_\ell(\mathcal{G})$ are the singular values of the voxel-wise Jacobian matrix $\mathcal{G} \in \mathbb{R}^{d \times M}$. A thin SVD

$$\mathcal{G} = U \Sigma V^\top, \quad U \in \mathbb{R}^{d \times r}, \Sigma = \text{diag}(\sigma) \in \mathbb{R}^{r \times r}, V \in \mathbb{R}^{M \times r}, \quad (6.6)$$

with $r = \min\{d, M\}$ gives

$$\nabla_{\mathcal{G}} \phi(\mathcal{G}) = U \text{diag}(g'(\sigma)) V^\top, \quad g'(\sigma_\ell) = \frac{\sigma_\ell}{\sqrt{\sigma_\ell^2 + \varepsilon^2}}. \quad (6.7)$$

In practice, it is unnecessary to form U if we exploit the small Gram matrix

$$G := \mathcal{G}^\top \mathcal{G} \in \mathbb{R}^{M \times M}, \quad G \succeq 0. \quad (6.8)$$

Let $G = V \Lambda V^\top$ with $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_M)$. Then $\lambda_\ell = \sigma_\ell^2$ for the nonzero spectrum

and V coincides with the right singular vectors of $\mathcal{G}$. On the range of $\mathcal{G}$ we have $U = \mathcal{G}V\Sigma^{-1}$, so substituting into (6.7) yields

$$\nabla_{\mathcal{G}}\phi(\mathcal{G}) = \mathcal{G}V \operatorname{diag}\left(\frac{g'(\sigma_\ell)}{\sigma_\ell}\right) V^\top, \quad (6.9)$$

with the continuous extension at $\sigma_\ell = 0$. For Charbonnier smoothing,

$$\frac{g'(\sigma_\ell)}{\sigma_\ell} = \frac{1}{\sqrt{\sigma_\ell^2 + \varepsilon^2}} = \frac{1}{\sqrt{\lambda_\ell + \varepsilon^2}}, \quad (6.10)$$

hence

$$\nabla_{\mathcal{G}}\phi(\mathcal{G}) = \mathcal{G}V \operatorname{diag}\left(\frac{1}{\sqrt{\lambda_\ell + \varepsilon^2}}\right) V^\top = \mathcal{G}(G + \varepsilon^2 I)^{-1/2}. \quad (6.11)$$

Equation (6.11) shows that the Jacobian-space gradient is obtained by applying the inverse square-root of the $M \times M$ SPD matrix $G + \varepsilon^2 I$ to the columns of $\mathcal{G}$.

For $M = 2$ modalities, even an eigendecomposition of G can be avoided. Writing

$$G = \begin{pmatrix} a & c \\ c & b \end{pmatrix}, \quad a = \|\mathcal{G}_{:,1}\|_2^2, \quad b = \|\mathcal{G}_{:,2}\|_2^2, \quad c = \langle \mathcal{G}_{:,1}, \mathcal{G}_{:,2} \rangle, \quad (6.12)$$

define $S := G + \varepsilon^2 I = \begin{pmatrix} s_{11} & s_{12} \\ s_{12} & s_{22} \end{pmatrix}$ with $s_{11} = a + \varepsilon^2$, $s_{22} = b + \varepsilon^2$, $s_{12} = c$, and the invariants

$$t := \operatorname{tr}(S) = s_{11} + s_{22}, \quad s := \sqrt{\det(S)} = \sqrt{s_{11}s_{22} - s_{12}^2}. \quad (6.13)$$

Since $\varepsilon > 0$ implies $S \succ 0$, the following identity holds for 2×2 SPD matrices:

$$S^{-1/2} = \sqrt{t+2s}(S+sI)^{-1} = \frac{\sqrt{t+2s}}{2\det(S)+st} \begin{pmatrix} s_{22}+s & -s_{12} \\ -s_{12} & s_{11}+s \end{pmatrix}. \quad (6.14)$$

Therefore, for two modalities the Jacobian-space gradient can be evaluated via

$$\nabla_{\mathcal{G}}\phi(\mathcal{G}) = \mathcal{G}S^{-1/2}, \quad S = G + \varepsilon^2 I, \quad (6.15)$$

using only dot products and scalar square-roots per voxel.

Image-space gradient. Define voxel-wise

$$Q_n := \nabla_{\mathcal{G}} \phi(\mathcal{G}_n(\mathbf{z})) \in \mathbb{R}^{d \times M}, \quad (6.16)$$

and let $Q = \{Q_n\}_{n=1}^N$. Applying the chain rule through the linear map $\mathbf{z} \mapsto D\mathbf{z}$ yields

$$\delta \tilde{\mathcal{R}}_{\text{TNV}}(\mathbf{z})[\delta \mathbf{z}] = \sum_{n=1}^N \langle Q_n, (D\delta \mathbf{z})_n \rangle_F = \langle Q, D\delta \mathbf{z} \rangle. \quad (6.17)$$

By definition of the adjoint D^* ,

$$\langle Q, D\delta \mathbf{z} \rangle = \langle D^*Q, \delta \mathbf{z} \rangle. \quad (6.18)$$

Defining the discrete divergence as $\text{div} := -D^*$, we obtain

$$\nabla_{\mathbf{z}} \tilde{\mathcal{R}}_{\text{TNV}}(\mathbf{z}) = -\text{div } Q, \quad (6.19)$$

where div acts column-wise on Q .

6.2 Weighted Directional Total Nuclear Variation (w-dTNV)

This section extends TNV [35, 34] in two orthogonal ways:

- high-resolution anatomical information is incorporated directionally rather than as an additional coupled intensity channel;
- voxel-wise, modality-wise diagonal weighting compensates for heterogeneous scaling and spatially varying reliability across modalities.

The starting point is the standard TNV construction: penalising the nuclear norm of a multi-modality Jacobian promotes joint sparsity and alignment of edges, since a low nuclear norm favours a locally low-rank Jacobian.

6.2.1 Directional operators from guidance

To incorporate guidance information, such as a high-resolution anatomical image, we borrow the operator from directional total variation (Equation 3.24). The key design choice is that the guidance image is used only to steer the penalty, not to add an extra column to the joint Jacobian. This avoids directly coupling an anatomical channel, with distinct units, resolution, and artefacts, to the emission channels through the singular spectrum.

If a guidance image were appended as an additional modality, then a single global scaling would have data-dependent and spatially non-uniform consequences. Where emission gradients are weak, even modest guidance gradients can dominate the local Jacobian and hence the singular values; where emission gradients are strong, the same scaling may become negligible. Since the nuclear norm couples modalities through the eigenstructure of the local Gram matrix, such reweighting can also transfer structured guidance artefacts into the coupled reconstruction. By contrast, a directional operator modifies which gradient components are penalised, without forcing the presence of edges where the emission data do not support them.

Applying $\Pi_{m,n}$ to each modality gradient defines the directional multi-modality Jacobian

$$\mathcal{G}_n^{\text{dir}}(\mathbf{z}) = [\Pi_{1,n}(Dz_1)_n \mid \cdots \mid \Pi_{M,n}(Dz_M)_n] \in \mathbb{R}^{d \times M}. \quad (6.20)$$

Here D denotes the chosen discrete directional-difference operator and each $\Pi_{m,n} \in \mathbb{R}^{d \times d}$ is constructed from the guidance field at voxel n . In the implementation used in this work, a guidance image a defines

$$\xi_n = \frac{(Da)_n}{\sqrt{\|(Da)_n\|_2^2 + \eta^2}}, \quad \Pi_{m,n} = I_d - \gamma_m \xi_n \xi_n^\top, \quad \gamma_m \in [0, 1]. \quad (6.21)$$

When η is not specified explicitly it is set from a robust scale estimate of $\|(Da)_n\|_2$, namely 0.01 times the difference between its 99th and 1st percentiles. In homogeneous regions of the guidance image, ξ_n is small and $\Pi_{m,n}$ is close to the identity, so the regulariser reduces to an approximately isotropic TNV-type coupling.

Using CT only to construct the directional operators decouples anatomical

steering from emission-to-emission weighting. Modality weights determine how strongly emission gradients are coupled to one another through the nuclear norm, while $\Pi_{m,n}$ determines how strongly the penalty prefers guidance-consistent edge directions. Retaining a two-modality Jacobian also preserves the efficient 2×2 Gram-matrix structure derived in Section 6.1.

6.2.2 Modality-wise weighting and the full w-dTNV prior

PET and SPECT images of the same administered Y-90 activity differ substantially in dynamic range: the effective count level, sensitivity, and noise characteristics vary by orders of magnitude between the two modalities, and in clinical imaging the two reconstructions may also operate at different spatial resolutions. If the directional emission gradients $\Pi_{m,n}(Dz_m)_n$ are used directly without normalisation, the nuclear-norm coupling is dominated by the modality with the larger gradient magnitudes.

To compensate, we introduce a voxel-wise, modality-wise diagonal weight matrix

$$B_n := \text{diag}(\beta_{1,n}, \beta_{2,n}) \in \mathbb{R}^{2 \times 2}, \quad \beta_{m,n} > 0. \quad (6.22)$$

The weighted directional Jacobian is

$$Y_n(\mathbf{z}) := \mathcal{G}_n^{\text{dir}}(\mathbf{z}) B_n = [\Pi_{1,n}(Dz_1)_n \mid \Pi_{2,n}(Dz_2)_n] \text{diag}(\beta_{1,n}, \beta_{2,n}) \in \mathbb{R}^{d \times 2}. \quad (6.23)$$

The columns of Y_n are $y_{m,n} := \beta_{m,n} \Pi_{m,n}(Dz_m)_n \in \mathbb{R}^d$.

The *weighted directional Total Nuclear Variation* (w-dTNV) prior is then

$$\mathcal{R}_{\text{w-dTNV}}(\mathbf{z}) := V_{\text{vox}} \sum_{n=1}^N \phi(Y_n(\mathbf{z})), \quad \phi(Y) := \sum_{\ell=1}^{\min(d,M)} g(\sigma_\ell(Y)), \quad (6.24)$$

where V_{vox} is the voxel volume and g is the Charbonnier smoothing from (6.3). Choosing equal weights $\beta_{m,n} \equiv 1$ recovers unweighted d TNV. Choosing $\beta_{m,n}$ proportional to the inverse of a characteristic modality scale brings both channels to a comparable range before the nuclear-norm coupling is evaluated. The PET/SPECT normalisation used in the application study is specified in Chapter 8.

Effect on the singular spectrum. Right-multiplying by B_n scales the m -th column of Y_n by $\beta_{m,n}$. This scales the singular values of Y_n in a non-trivial way, since the singular values of YB are not simply $\sigma_\ell(Y)\beta_m$. The intended effect is instead on the Gram matrix: $Y_n^\top Y_n$ contains modality self-terms and cross-terms after scaling, so suitable $\beta_{m,n}$ values allow both modalities to contribute comparably to the coupled singular spectrum. The 2×2 Gram-matrix algebra in Section 6.1 applies directly to this weighted case.

6.3 Conclusion

This chapter introduced a smoothed, weighted, and directional Total Nuclear Variation (w -dTNV) prior for synergistic reconstruction. Charbonnier smoothing of the singular values yields a differentiable spectral functional whose Jacobian-space gradient can be written in Gram-matrix form. For two coupled emission modalities this reduces the key spectral operation to a 2×2 inverse square-root, allowing the prior and its gradient to be evaluated robustly without a general per-voxel SVD.

The chapter also separated two roles that are easily confounded in multimodality reconstruction. Anatomical information is used directionally through $\Pi_{m,n}$, steering the penalty without adding CT as another coupled intensity channel. Modality scaling is handled separately through the diagonal weight matrix B_n , allowing PET and SPECT gradients to enter the nuclear-norm coupling on comparable scales. Chapter 7 uses these definitions to construct practical preconditioners and subset-selection strategies for large-scale w -dTNV optimisation.

Chapter 7

Preconditioning and Subset Selection for Synergistic Reconstruction

Chapter 6 defined the smoothed weighted directional total nuclear variation prior. This chapter describes the optimisation machinery used to apply that prior in large-scale PET/SPECT reconstruction experiments. The emphasis is methodological: the experiments described here compare preconditioners and subset-selection strategies, but no numerical results are reported in this chapter.

The implementation inspected for this chapter is the current reconstruction codebase. This is important because the active preconditioner-combination path has changed relative to the application study in Chapter 8: the paper reconstructions use a local diagonal w -dT_{NTV} curvature approximation combined with a floor-regularised Lehmer mean, whereas the current optimisation experiments use Hessian-sum inversion for the prior-aware preconditioners.

7.1 Stochastic MAP Optimisation

Let $\mathbf{z} = (z_1, z_2)$ denote the PET and SPECT images represented on a common reconstruction grid. The objective has the form

$$\mathcal{O}(\mathbf{z}) = \sum_{i=1}^{N_d} \mathcal{D}_i(\mathbf{z}) + \mathcal{R}_{w\text{-dT}_{\text{NTV}}}(\mathbf{z}) + \mathcal{I}_{\mathbb{R}_+}(\mathbf{z}), \quad (7.1)$$

where $\mathcal{D}_i$ are subset data-fidelity terms, $\mathcal{R}_{w\text{-dTNV}}$ is the prior from Chapter 6, and $\iota_{\mathbb{R}_+}$ enforces non-negativity. Each $\mathcal{D}_i$ is a Poisson negative log-likelihood term composed with the relevant projector, image-space registration operator, and optional resolution model.

The stochastic optimisation experiments use SVRG [74, 39, 38]. For stochastic function list $\{\mathcal{O}_i\}_{i=1}^{N_s}$ and anchor point $\tilde{\mathbf{z}}$, the estimator is

$$\mathbf{g}^{(k)} = \nabla \mathcal{O}(\tilde{\mathbf{z}}) + \frac{1}{p_{i_k}} \left[\nabla \mathcal{O}_{i_k}(\mathbf{z}^{(k)}) - \nabla \mathcal{O}_{i_k}(\tilde{\mathbf{z}}) \right], \quad i_k \sim p, \quad (7.2)$$

with probabilities p_i determined by the selected subset strategy. In the implemented experiments the sampler draws with replacement. The full-gradient snapshot is refreshed once per epoch by setting the SVRG snapshot interval equal to the epoch length.

The image update is a projected, preconditioned gradient step,

$$\mathbf{z}^{(k+1)} = \left\{ \mathbf{z}^{(k)} - \varsigma^{(k)} P^{(k)} \mathbf{g}^{(k)} \right\}_{\geq 0}, \quad (7.3)$$

where $P^{(k)}$ is either a scalar diagonal preconditioner per modality or a voxel-wise 2×2 block preconditioner coupling PET and SPECT update components. The step size follows a linear decay rule,

$$\varsigma^{(k)} = \frac{\varsigma^{(0)}}{1 + \eta_{\text{relax}} k}, \quad \varsigma^{(0)} = 1, \quad \eta_{\text{relax}} = 0.002 \quad (7.4)$$

for the preconditioner and subset-selection experiments described below.

7.2 Prior-Aware Preconditioning

Preconditioning is critical for subset-based MAP reconstruction because the data term and the edge-preserving prior can have very different local curvature scales [75, 37, 76, 38]. The current experiments compare data-only BSREM scaling with several prior-aware curvature surrogates for $w\text{-dTNV}$.

7.2.1 Data Curvature

For modality m , let s_m denote the sensitivity image on the shared grid,

$$s_m = W_m A_m^\top \mathbf{1}, \quad (7.5)$$

where A_m is the native-grid system matrix and W_m maps native images to the shared prior grid. The EM/BSREM data preconditioner is

$$P_{\text{BSREM},m} = (z_m + \epsilon_x) \oslash s_m, \quad (7.6)$$

with element-wise division. Equivalently, the corresponding diagonal data Hessian surrogate is

$$\hat{H}_{\text{data},m} = s_m \oslash (z_m + \epsilon_x). \quad (7.7)$$

The implementation optionally smooths the numerator image used in the EM preconditioner with a Gaussian filter. In the experiments below the smoothing FWHM is 5 mm in each direction.

7.2.2 Current w -dTNN Prior Curvature

The current prior-aware preconditioners use the majorisation–minimisation (MM)/iteratively reweighted least squares (IRLS) curvature implied by the smoothed spectral penalty. At voxel n , let

$$Y_n = [\beta_{1,n} \Pi_{1,n}(Dz_1)_n \mid \beta_{2,n} \Pi_{2,n}(Dz_2)_n] \in \mathbb{R}^{d \times 2} \quad (7.8)$$

be the weighted directional Jacobian from Chapter 6. The active code forms the modality-space MM weight

$$C_n = \left(Y_n^\top Y_n + \epsilon^2 I_2 \right)^{-1/2} = \begin{pmatrix} c_{11,n} & c_{12,n} \\ c_{12,n} & c_{22,n} \end{pmatrix}. \quad (7.9)$$

This matrix is the same 2×2 inverse square-root used in the gradient calculation, now interpreted as the local quadratic-surrogate weight.

For diagonal preconditioning, the finite-difference stencil and directional projectors are collapsed to a per-voxel participation factor $\tau_{m,n}$. This factor includes the number of valid forward/backward stencil incidences at voxel n , the squared finite-difference scale factors, and the diagonal entries of $\Pi_{m,n}^\top \Pi_{m,n}$. The local diagonal w -dTENV curvature approximation is

$$\hat{H}_{\mathcal{R},m,n}^{\text{loc-diag}} = V_{\text{vox}} \beta_{m,n}^2 c_{mm,n} \tau_{m,n}. \quad (7.10)$$

This is the diagonal extraction of the local MM curvature surrogate, and is therefore an approximation rather than a strict Loewner majoriser in general. The Gershgorin diagonal majoriser inflates the modality-space diagonal by the absolute off-diagonal coupling,

$$\hat{H}_{\mathcal{R},1,n}^{\text{diag-G}} = V_{\text{vox}} \beta_{1,n}^2 (c_{11,n} + |c_{12,n}|) \tau_{1,n}, \quad \hat{H}_{\mathcal{R},2,n}^{\text{diag-G}} = V_{\text{vox}} \beta_{2,n}^2 (c_{22,n} + |c_{12,n}|) \tau_{2,n}. \quad (7.11)$$

For block preconditioning, the PET and SPECT update components at each voxel are coupled by a 2×2 block. Let $\mathcal{N}(n)$ be the set of target voxels whose finite-difference stencil depends on source voxel n . For target voxel $r \in \mathcal{N}(n)$, let $c_{r \leftarrow n} \in \mathbb{R}^d$ be the finite-difference impulse vector produced in the directional channels when source voxel n is perturbed. The local voxel-block curvature approximation uses

$$\alpha_{r \leftarrow n}^{mm'} = c_{r \leftarrow n}^\top \Pi_{m,r}^\top \Pi_{m',r} c_{r \leftarrow n}, \quad (7.12)$$

and

$$\hat{H}_{\mathcal{R},n}^{\text{loc-block}} = V_{\text{vox}} \sum_{r \in \mathcal{N}(n)} [B_r C_r B_r] \circ A_{r \leftarrow n}, \quad A_{r \leftarrow n} = \begin{pmatrix} \alpha_{r \leftarrow n}^{11} & \alpha_{r \leftarrow n}^{12} \\ \alpha_{r \leftarrow n}^{21} & \alpha_{r \leftarrow n}^{22} \end{pmatrix}, \quad (7.13)$$

where $B_r = \text{diag}(\beta_{1,r}, \beta_{2,r})$ and $\circ$ denotes the Hadamard product. The long reference reconstructions in the subset-selection study use a more conservative voxel-block curvature majoriser.

7.2.3 Combining Data and Prior Curvature

The current experiments invert the summed data and prior Hessian surrogates. For diagonal methods,

$$P_{m,n}^{\text{diag}} = \left(\hat{H}_{\text{data},m,n} + \hat{H}_{\mathcal{R},m,n} \right)^{-1}, \quad (7.14)$$

with numerical floors and optional maximum-value caps. For block methods,

$$P_n^{\text{block}} = \left(\begin{pmatrix} \hat{H}_{\text{data},1,n} & 0 \\ 0 & \hat{H}_{\text{data},2,n} \end{pmatrix} + \hat{H}_{\mathcal{R},n} \right)^{-1}. \quad (7.15)$$

This differs from the application implementation in Chapter 8, where the same local diagonal w -dTNNV curvature approximation was combined with the data preconditioner using a Lehmer mean.

The preconditioner is updated once per epoch in the experiments unless an explicit update interval is supplied. Between updates the curvature is therefore stale, which is a pragmatic compromise between optimisation stability and the computational cost of evaluating $\hat{H}_{\mathcal{R}}$.

Table 7.1: Preconditioner settings used in the current preconditioner experiments.

Preconditioner setting	Prior curvature model
Data-only EM/BSREM scaling	none
Local diagonal w -dTNNV curvature approximation	diagonal extraction of local MM curvature
Gershgorin diagonal w -dTNNV curvature majoriser	diagonal majoriser including cross-modal
Local voxel-block w -dTNNV curvature approximation	voxel-wise 2×2 local MM curvature
Conservative voxel-block w -dTNNV curvature majoriser	voxel-wise 2×2 conservative majoriser

7.2.4 Preconditioner Comparison Methods

Four preconditioner settings are compared: data-only BSREM, the local diagonal w -dTNNV curvature approximation, the Gershgorin diagonal w -dTNNV curvature majoriser, and the local voxel-block w -dTNNV curvature approximation. These settings are summarised in Table 7.1.

All sweep reconstructions use the same objective, prior settings, and initialisation within each dataset. The fixed sweep parameters are $\alpha = \beta = 1$, $\gamma_{\text{TNNV}} = 0.02$,

initial step size $\varsigma^{(0)} = 1$, relaxation parameter $\eta_{\text{relax}} = 0.002$, and 25 epochs. Time-of-flight PET is disabled in these optimisation experiments. Sensitivity-derived support masks are enabled with relative threshold 10^{-3} , and both EM and prior-aware preconditioner maximum-value caps are disabled. The EM component of the preconditioner is smoothed with a 5 mm Gaussian FWHM.

Experiments are configured for a single-bed anthropomorphic phantom dataset and a two-bed clinical PET/SPECT dataset. For the one-bed sweep, the reconstruction is routed through the subset-selection runner with separate PET/SPECT subset functions and the prior evaluated deterministically every iteration. The two-bed sweep uses the corresponding two-bed reconstruction path. Objective values, saved images, and convergence metrics are recorded for subsequent comparison against long-run reference reconstructions. The reference reconstructions use the conservative block majoriser and extended epoch counts, but their numerical results are not reported here.

7.3 Subset Selection

Subset methods accelerate early iterations in emission tomography by replacing full data passes with stochastic or ordered blocks of projection data [77]. In synergistic PET/SPECT reconstruction there is an additional design choice: PET, SPECT, and the coupled prior can be grouped together or sampled separately. Previous work on synergistic subset selection [[porter'optimising'2024](#)] **TODO: Zenodo-managed bibliography entry pending for [porter_optimising_2024](#); IEEE Xplore document [10654962](#)**. motivated a more systematic comparison of these choices under the current SVRG and preconditioning implementation.

7.3.1 Data-Subset Organisation

Two data-subset organisations are tested.

Separate subsets. PET and SPECT subsets are treated as distinct stochastic functions. For a one-bed acquisition with $N_{\text{PET}} = 18$ and $N_{\text{SPECT}} = 18$, this gives

$$\{\mathcal{D}_i\}_{i=1}^{36} = \{\mathcal{D}_{\text{PET},1}, \dots, \mathcal{D}_{\text{PET},18}, \mathcal{D}_{\text{SPECT},1}, \dots, \mathcal{D}_{\text{SPECT},18}\}. \quad (7.16)$$

For a two-bed PET acquisition with 9 PET subsets per bed and 9 SPECT subsets, the separate list contains 27 data functions: 18 PET bed-specific functions and 9 SPECT functions.

Paired subsets. PET and SPECT subsets with matching indices are summed into one stochastic function. For one bed this gives

$$\mathcal{D}_j^{\text{pair}} = \mathcal{D}_{\text{PET},j} + \mathcal{D}_{\text{SPECT},j}, \quad j = 1, \dots, 18. \quad (7.17)$$

For two PET beds, each paired function contains both PET bed-position subsets and the corresponding SPECT subset,

$$\mathcal{D}_j^{\text{pair}} = \mathcal{D}_{\text{PET},1,j} + \mathcal{D}_{\text{PET},2,j} + \mathcal{D}_{\text{SPECT},j}, \quad j = 1, \dots, 9. \quad (7.18)$$

The two-bed paired mode therefore requires equal PET and SPECT subset counts after partitioning.

7.3.2 Prior Sampling Strategy

The prior is handled in two ways. In *always* mode, the prior gradient is evaluated deterministically on every stochastic iteration and is not part of the sampler. In *subset* mode, the prior is appended as an additional stochastic function and is sampled with a prescribed probability.

Let N_d be the number of stochastic data functions after the selected data grouping, and let u be the target number of prior updates per data epoch. The implemented prior-sampling probability is

$$p_{\mathcal{R}} = \frac{u}{N_d + u}, \quad (7.19)$$

and each data function probability is scaled by $1 - p_{\mathcal{R}}$. The epoch length is adjusted to

$$K = \left\lceil \frac{N_d}{1 - p_{\mathcal{R}}} \right\rceil. \quad (7.20)$$

For paired mode the default is $u = N_d$, so the prior and data aggregate each receive probability $1/2$. For separate mode the default is $u = N_d/2$, so the prior receives probability $1/3$ and the data aggregate receives probability $2/3$. These defaults keep

the number of expected prior updates comparable to the paired setting while avoiding a prior update after every separate PET or SPECT subset.

7.3.3 Subset-Selection Experiment Methods

The main subset-selection sweep crosses two subset organisations (separate, paired), two prior modes (always, subset), two preconditioner settings (data-only BSREM and the local diagonal w -dTNN curvature approximation), and the configured directional strength $\gamma_{\text{TNV}} = 0.02$. The one-bed configuration uses 18 PET and 18 SPECT subsets. The two-bed configuration uses 9 PET subsets per bed position and 9 SPECT subsets.

All subset-selection runs use SVRG with random-with-replacement sampling, a snapshot interval equal to one adjusted epoch, $\alpha = \beta = 1$, initial step size 1, relaxation parameter 0.002, and 25 epochs. The prior is the directional TNV variant with a six-direction stencil, Neumann boundary conditions, Charbonnier smoothing, and no additional modality-specific priors. Sensitivity-derived support masks are enabled, and the EM preconditioner component is smoothed with a 5 mm Gaussian FWHM. The random seed is fixed at 42 for reproducibility.

Long-run reference reconstructions are generated using separate subset organisation, deterministic prior evaluation every iteration, the conservative voxel-block w -dTNN curvature majoriser, and 1000 epochs. The planned analysis compares objective histories and image-domain errors against these references at matched epoch counts. No subset-selection results are included in this methods-only chapter.

7.4 Multi-Bed PET Acquisition

The subset-selection and application experiments include PET acquisitions spanning multiple partially overlapping bed positions. This material is part of the optimisation chapter because it changes the data-fidelity organisation and the preconditioner sensitivity image, while the prior remains the same shared-grid w -dTNN prior.

7.4.1 Notation

Let x be the PET activity image on the combined reconstruction grid. For bed position $b = 1, \dots, B$, let R_b restrict the combined image to the axial support reconstructed by

bed b . The expected PET data for bed b are

$$\bar{y}_b(x) = A_b R_b x + r_b, \quad (7.21)$$

where A_b is the bed-specific PET system matrix and r_b collects additive corrections. The joint PET data term is

$$\mathcal{D}_{\text{PET}}^{\text{multi}}(x) = \sum_{b=1}^B \mathcal{D}_b(y_b \mid A_b R_b x + r_b), \quad (7.22)$$

with gradient

$$\nabla \mathcal{D}_{\text{PET}}^{\text{multi}}(x) = \sum_{b=1}^B R_b^\top A_b^\top \nabla_{\bar{y}_b} \mathcal{D}_b(y_b \mid A_b R_b x + r_b). \quad (7.23)$$

The corresponding PET sensitivity image on the combined grid is

$$s_{\text{PET}}^{\text{multi}} = \sum_{b=1}^B R_b^\top A_b^\top \mathbf{1}. \quad (7.24)$$

This sensitivity is used in the EM/BSREM data curvature term in exactly the same way as the single-bed sensitivity in Equation (7.7).

7.4.2 Why Not Stitch Independent Reconstructions?

A common practical alternative is to reconstruct each PET bed independently and fuse the resulting images using sensitivity-weighted stitching in the overlap region. This is benign for some early-stopped or weakly regularised reconstructions, but it is problematic with edge-preserving priors. The prior is then evaluated separately on each bed image, so noise patterns and anatomical edge decisions in the overlap are made independently before fusion. This can create abrupt intensity changes or inconsistent edge placement across the bed boundary [porter'simultaneous'2025]**TODO:** Zenodo-managed bibliography entry pending for porter_simultaneous_2025; The PET is Wonderful multi-bed MAP reconstruction abstract..

The joint formulation in Equation (7.22) avoids this by estimating a single combined PET image. Each bed contributes its own projection data and sensitivity,

while the image-domain prior is evaluated once over the combined field of view.

7.4.3 Use in the Subset Experiments

In the two-bed subset-selection experiments, each bed position is partitioned into the same number of PET subsets. In separate mode, PET bed subsets remain separate stochastic functions. In paired mode, subset index j groups PET bed 1, PET bed 2, and SPECT subset j into a single stochastic function as in Equation (7.18). The preconditioner receives the combined-bed PET sensitivity from Equation (7.24), so overlap voxels naturally reflect the contribution of both bed positions.

7.5 Conclusion

This chapter has separated the optimisation machinery from the definition of the w -dTNV prior itself. The current preconditioner experiments compare BSREM scaling with MM-derived prior-aware diagonal and block curvature surrogates, combined through Hessian-sum inversion in the active reconstruction code. The subset-selection experiments independently vary modality grouping and prior sampling, including a two-bed PET setting where the data term is evaluated jointly over bed positions. These methods provide the framework into which numerical convergence and image-quality results can be inserted once available.

Chapter 8

Weighted Directional Total Nuclear Variation for Joint ^{90}Y PET/SPECT Reconstruction with CTAC-derived Guidance

This chapter applies the weighted directional total nuclear variation (w -dTNV) prior developed in Chapter 6 to the clinically relevant problem of post-treatment ^{90}Y activity imaging after SIRT. We consider joint reconstruction of low-count ^{90}Y PET and bremsstrahlung SPECT, using computed tomography attenuation correction (CTAC) information directionally (via the CTAC-derived attenuation map μ) to stabilise edge geometry without enforcing intensity correlation.

Relative to Chapter 6, the focus here is validation and comparison with other methods: we specify modality-dependent forward models, co-registration operators, practical normalisation for heterogeneous PET/SPECT scales, and a reproducible evaluation protocol on a National Electrical Manufacturers Association (NEMA) International Electrotechnical Commission (IEC) phantom with PET bootstrapping and on a post-SIRT patient cohort.

8.1 Motivation and clinical context

Quantitative post-treatment imaging is required for personalised dosimetry following ^{90}Y SIRT [78]. Pre-treatment planning based on $^{99\text{m}}\text{Tc}$ -macroaggregated albumin (MAA) often mis-predicts microsphere deposition, motivating post-treatment verification of delivered activity distributions [79].

Two ^{90}Y modalities are clinically available:

- **^{90}Y PET:** detects the rare internal pair-production branch, providing comparatively high spatial resolution but extremely low count statistics [21].
- **Bremsstrahlung SPECT:** has substantially higher corrected counts (about $15\times$ for the data studied here) but suffers from severe blur, septal penetration, and scatter due to the continuous emission spectrum [58, 80].

Both modalities are typically acquired alongside CT for attenuation correction and anatomical localisation. Since PET and SPECT image the same underlying microsphere distribution while CTAC provides structural rather than activity information, independent reconstruction discards both a strong physical coupling and useful anatomical context.

The sequential hybrid kernel expectation–maximisation (SHKEM) method of Deidda *et al.* improves lesion contrast by reconstructing one modality to guide another [45, 81]. However, the sequential construction remains sensitive to iteration choice and noise amplification. This motivates a truly joint MAP formulation with a coupled regulariser that promotes shared geometry while allowing modality-specific contrast.

8.2 Joint MAP formulation for PET/SPECT

We use the joint MAP framework introduced in Section 3.2.3. For the present application we set $M = 2$ modalities, with $m = 1$ denoting PET and $m = 2$ denoting SPECT. Let $z_m \in \mathbb{R}_+^N$ denote the image for modality m represented on a shared grid and let $y_m \in \mathbb{N}^{I_m}$ denote the corresponding projection data.

The native image grid for modality m may differ from the shared grid. Once registration is fixed, let $W_m : \mathbb{R}_+^{J_m} \rightarrow \mathbb{R}_+^N$ denote the discrete resampling operator from

the native grid to the shared grid, and let $W_m^* : \mathbb{R}_+^N \rightarrow \mathbb{R}_+^{J_m}$ denote its discrete adjoint with respect to the chosen image-space inner products. The adjoint maps shared-grid variables back to the native grid before application of the system matrix. For each modality we therefore use the affine Poisson model

$$\bar{y}_m(z_m) = A_m W_m^* z_m + b_m, \quad (8.1)$$

where A_m is the system matrix and b_m collects additive corrections. For example, randoms and scatter for PET and simulated scatter/penetration terms for SPECT (see Sections 8.5.1 and 8.5.2).

The joint objective is

$$\hat{\mathbf{z}} = \arg \min_{\mathbf{z} \geq 0} \left\{ \sum_{m=1}^2 \mathcal{D}_m(y_m | A_m W_m^* z_m + b_m) + \mathcal{R}(\mathbf{z}) \right\}, \quad (8.2)$$

where $\mathcal{D}_m$ is the Poisson negative log-likelihood (up to an additive constant),

$$\mathcal{D}_m(y_m | \bar{y}_m) = \sum_{i=1}^{I_m} ((\bar{y}_m)_i - (y_m)_i \log(\bar{y}_m)_i), \quad (8.3)$$

and $\mathcal{R}$ is the coupled prior. In this chapter, $\mathcal{R}$ is the smoothed *w-dTNV* prior $\tilde{\mathcal{R}}_{\text{w-dTNV}}$ defined in Section 6.2 (Chapter 6), specialised to PET/SPECT with CTAC guidance and practical global modality normalisation. The overall regularisation strength is carried by the modality weights introduced below, so no additional scalar λ is used in the implementation.

8.3 *w-dTNV instantiated for ^{90}Y PET/SPECT with CTAC guidance*

8.3.1 Directional guidance from the CTAC-derived μ -map

As in Section 6.2.1, we use CT information directionally rather than appending CT as a third coupled channel. Here the guidance image is the CTAC-derived attenuation map μ (resampled to the PET grid). We pre-smooth μ with a small blurring kernel

(FWHM = 2 mm) to reduce high-frequency CT noise while retaining high guidance edge resolution. We then compute a stabilised direction field in the discrete gradient space,

$$\xi_n = \frac{(D\mu)_n}{\|(D\mu)_n\|_2 + \eta}, \quad \eta = 0.05 \cdot \max_n \|(D\mu)_n\|_2, \quad (8.4)$$

where D is the chosen directional finite-difference operator (specified in Section 8.3.3). The stabiliser η prevents amplification of small-magnitude μ differences in homogeneous regions.

We can then form the modality-wise directional operators $\Pi_{m,n} \in \mathbb{R}^{d \times d}$ using the construction already adopted in Chapter 6:

$$\Pi_{m,n} = I_d - \gamma_m \xi_n \xi_n^\top, \quad \gamma_m \in [0, 1]. \quad (8.5)$$

Unless stated otherwise, we set $\gamma_1 = \gamma_2 = 1$. Since ξ_n is stabilised and generally has norm less than one, this should be interpreted as anisotropic shrinkage of gradient components parallel to the local μ -edge normal direction rather than as a hard projection.

8.3.2 Global modality normalisation as a constant diagonal weighting

Chapter 6 introduced a general voxel-wise diagonal modality weighting $B_n = \text{diag}(\beta_{1,n}, \dots, \beta_{M,n})$. In the present study we use a spatially uniform normalisation to mitigate the dominant scale disparity between ^{90}Y PET and bremsstrahlung SPECT:

$$B_n \equiv B := \text{diag}(\rho_1, \rho_2), \quad \rho_1, \rho_2 > 0. \quad (8.6)$$

Weights are computed from early-stopped unregularised reconstructions represented on the shared grid, z_m^{init} , via robust dynamic-range (percentile) scaling over a CTAC-derived foreground mask Ω' :

$$\rho_m = \frac{\beta_m}{q_{0.99}(z_m^{\text{init}}|\Omega') - q_{0.01}(z_m^{\text{init}}|\Omega')}, \quad m \in \{1, 2\}, \quad (8.7)$$

where $q_p(\cdot|\Omega')$ denotes the p th percentile over voxels in Ω' , and β_m are user-controlled regularisation weights used for operating-point selection (Section 8.7.1). Percentile scaling provides a robust dynamic-range estimate and reduces sensitivity to outliers. This matches the intent of Section 6.2.2 while avoiding an additional spatially varying weight for this investigation [82].

8.3.3 Directional finite-difference operator for PET/SPECT coupling

We use the generic finite-difference definition from Chapter 6 (Section 6.1) but specify the set of offsets to better accommodate SPECT blur and modest registration error. Let $\Delta = (\Delta_x, \Delta_y, \Delta_z)$ denote the voxel sizes in SI units. We define a nine-direction half-stencil

$$\mathcal{E}_1 = \{(1, 0, 0), (0, 1, 0), (0, 0, 1)\} \cup \{(\pm 1, 1, 0), (\pm 1, 0, 1), (0, 1, \pm 1)\}, \quad (8.8)$$

and use both one- and two-voxel steps, $\mathcal{E} = \mathcal{E}_1 \cup \{2e : e \in \mathcal{E}_1\}$, yielding $d = 18$ directions. For a set of offsets $\mathcal{E} = \{e_\ell\}_{\ell=1}^d$, define the d -directional forward difference operator D by

$$((Dz)_n)^{(\ell)} := \frac{z_{n+e_\ell} - z_n}{\|e_\ell\|_\Delta}, \quad \|e\|_\Delta := \sqrt{(e_x\Delta_x)^2 + (e_y\Delta_y)^2 + (e_z\Delta_z)^2}, \quad (8.9)$$

with Neumann boundary conditions. This follows the “extended Jacobian” motivation for TNV-type coupling [34] and empirically improved stability in the present PET/SPECT setting.

8.3.4 The w - d TNV Jacobian and regulariser (two modalities)

Using the notation of Chapter 6, we form the guided gradients $\Pi_{m,n}(Dz_m)_n \in \mathbb{R}^d$ and define the weighted directional Jacobian

$$Y_n := Y_n^{\text{w-dTNV}}(\mathbf{z}) = \left[\Pi_{1,n}(Dz_1)_n \mid \Pi_{2,n}(Dz_2)_n \right] B \in \mathbb{R}^{d \times 2}. \quad (8.10)$$

The regulariser is the smoothed spectral potential introduced in Chapter 6, weighted by the voxel volume V_{vox} ,

$$\tilde{\mathcal{R}}_{w\text{-dTNV}}(\mathbf{z}) = V_{\text{vox}} \sum_{n=1}^N \phi(Y_n), \quad \phi(Y) = \sum_{\ell=1}^2 g(\sigma_{\ell}(Y)), \quad (8.11)$$

where g is the Charbonnier smoothing (Equation 6.3) and $\sigma_{\ell}(Y)$ are singular values. We set the smoothing parameter ε relative to initial joint-gradient magnitudes:

$$\varepsilon = 0.05 \cdot \text{median}_{n \in \Omega'} \sigma_1(Y_n^{\text{init}}), \quad Y_n^{\text{init}} := Y_n^{w\text{-dTNV}}(\mathbf{z}^{\text{init}}), \quad (8.12)$$

with $\mathbf{z}^{\text{init}} = (W_1 x_1^{\text{init}}, W_2 x_2^{\text{init}})$.

If $\Pi_{2,n}(Dz_2)_n \approx 0$ at some voxel (e.g. SPECT locally uniform after guidance), then Y_n has an effectively single nonzero column and $\phi(Y_n)$ reduces to a TV-type penalty on the remaining modality column. This is the same “no forced edges” behaviour discussed in Chapter 6.

8.3.5 Ablation comparators

To separate the contributions of PET/SPECT coupling and CTAC-derived directional guidance, we compare w -dTNV with two variational ablations. The first is weighted TNV (w -TNV), which keeps the two-modality Jacobian and the same modality normalisation but removes directional CTAC guidance by setting $\Pi_{m,n} = I_d$ for both modalities:

$$Y_n^{w\text{-TNV}}(\mathbf{z}) = \left[(Dz_1)_n \mid (Dz_2)_n \right] B. \quad (8.13)$$

This tests PET/SPECT coupling without anatomical anisotropy.

The second is directional TV (dTV), which keeps CTAC-derived directional guidance but removes PET/SPECT coupling. It is obtained by applying the same directional operator only to the PET gradient,

$$\tilde{\mathcal{R}}_{\text{dTV}}(z_1) = V_{\text{vox}} \sum_{n=1}^N \phi(\rho_1 \Pi_{1,n}(Dz_1)_n), \quad (8.14)$$

which reduces to a smoothed single-channel directional-TV penalty. This tests

anatomical anisotropy without using SPECT as a coupled functional channel. All variational methods use the same forward models, registration, foreground mask, dynamic-range normalisation, and optimisation framework; only the regulariser structure is changed.

8.4 Optimisation and implementation

We minimise the joint objective (8.2) using a subset stochastic variance-reduced gradient (SVRG) method [74], consistent with the stochastic framework studied in Chapter 7. We decompose the objective into S subset terms,

$$\mathcal{O}(\mathbf{z}) = \sum_{s=1}^S \mathcal{O}_s(\mathbf{z}), \quad \mathcal{O}_s(\mathbf{z}) = \left[\mathcal{D}_s(\mathbf{z}) + \frac{1}{S} \tilde{\mathcal{R}}_{\text{w-dTNV}}(\mathbf{z}) \right], \quad (8.15)$$

so that each subset gradient includes a consistent $(1/S)\nabla\tilde{\mathcal{R}}_{\text{w-dTNV}}$ contribution. Subsets are formed using staggered, view-interleaved partitioning. For the joint reconstructions, subsets are split by modality as well as by view so that each subset term contains either PET or SPECT data, matching the subset-selection analysis in Chapter 7.

At iteration k , with anchor point $\tilde{\mathbf{z}}$ refreshed once per epoch, we use the SVRG estimator

$$\mathbf{g}^{(k)} = \nabla\mathcal{O}(\tilde{\mathbf{z}}) + S \left[\nabla\mathcal{O}_{s_k}(\mathbf{z}^{(k)}) - \nabla\mathcal{O}_{s_k}(\tilde{\mathbf{z}}) \right], \quad s_k \sim \text{Unif}\{1, \dots, S\}, \quad (8.16)$$

and apply a projected diagonally preconditioned update

$$\mathbf{z}^{(k+1)} = \left\{ \mathbf{z}^{(k)} - \varsigma^{(k)} P^{(k)} \mathbf{g}^{(k)} \right\}_{\geq 0}, \quad (8.17)$$

with a decaying step size

$$\varsigma^{(k)} = \frac{\varsigma^{(0)}}{1 + \delta k}, \quad \varsigma^{(0)} = 1, \quad \delta = 0.005. \quad (8.18)$$

Reconstructions were run for 250 epochs, corresponding to approximately 500 full data passes because the full-gradient snapshot is recomputed once per epoch. This

fixed stopping point was used for all methods rather than a case-specific convergence criterion.

8.4.1 Diagonal preconditioning

We use a diagonal preconditioner combining:

- an EM-type data scaling [76, 37],

$$P_{\text{data},m} = z_m \oslash \left(W_m A_m^\top \mathbf{1} \right), \quad (8.19)$$

represented on the shared grid for each modality, and

- the local diagonal w -dTNN curvature approximation for $\tilde{\mathcal{R}}_{w\text{-dTNN}}$ described in Chapter 7 (Section 7.2).

In practice, $P_{\mathcal{R}}$ is updated once per epoch (and held fixed within the epoch) to limit overhead. Since $M = 2$, the per-voxel spectral quantities required by both the w -dTNN gradient and the diagonal surrogates can be evaluated efficiently using the 2×2 Gram-matrix structure (Equation 6.14 and the discussion around Equation 6.15).

Inspired by [38], we combine P_{data} and $P_{\mathcal{R}}$ via a floor-regularised Lehmer mean of order $q = 0.1$. This describes the application implementation used for the reconstructions in this chapter; the later optimisation experiments in Chapter 7 replace this combination rule with Hessian-sum inversion.

$$\tilde{P}_{\text{data}} = \max(P_{\text{data}}, p_{\min}), \quad \tilde{P}_{\mathcal{R}} = \max(P_{\mathcal{R}}, p_{\min}), \quad p_{\min} = 10^{-6}, \quad (8.20)$$

$$P = \frac{\tilde{P}_{\text{data}}^q + \tilde{P}_{\mathcal{R}}^q}{\tilde{P}_{\text{data}}^{q-1} + \tilde{P}_{\mathcal{R}}^{q-1}}. \quad (8.21)$$

For $q = 0$ this reduces to harmonic averaging. The choice $q = 0.1$ gives a slightly less severe, but still conservative, compromise between data and regulariser scaling. The floor prevents numerical zeros in either constituent from freezing the update.

8.4.2 Comparator: sequential hybrid kernel EM

We compare against the sequential hybrid kernel expectation–maximisation (SHKEM) method of Deidda *et al.* [45], using the reconstruction order and

hyperparameter settings reported as optimal in that work. The sequential order is CT-guided SPECT followed by CT/SPECT-guided PET: SPECT is reconstructed first using CT-derived anatomical features and hybrid kernel updates, and the resulting SPECT reconstruction is then used together with CT-derived features and the current PET iterate to construct the PET hybrid kernel [81]. To reduce noise amplification we use frozen kernels, halting kernel updates after 27 PET subiterations and 72 SPECT subiterations, matching the comparator configuration in [45]. All kernel parameters followed the original implementation from Deidda *et al.* For clinical PET data, the SHKEM implementation is applied in the same joint multi-bed setting as the variational methods.

8.4.3 Software implementation

All reconstructions were implemented in the Synergistic Image Reconstruction Framework (SIRF) [83] using Software for Tomographic Image Reconstruction (STIR) projector backends [84, 85, 86] and the Core Imaging Library [87], with custom code for the proposed regulariser. The SHKEM implementation follows [45], extended here to the same joint multi-bed PET reconstruction setting used by the variational methods.

8.5 Forward models, corrections, and co-registration

8.5.1 PET forward model

PET list-mode data were processed using the vendor toolchains to generate normalisation, attenuation, and additive corrections (randoms and scatter), which were held fixed during reconstruction [88, 89]. Reconstructions use a voxel size of $3.27 \times 2.13 \times 2.13 \text{ mm}^3$ (z,y,x) on a grid of $83 \times 255 \times 255$.

Patient PET data were acquired across two partially overlapping bed positions. For all methods, these bed positions were reconstructed jointly in a single combined-bed image space, with separate forward- and back-projections for each bed position; for the variational methods, the image-domain prior is evaluated over the combined volume. This is the multi-bed formulation derived in Chapter 7, and avoids the overlap-region artefacts that can arise when edge-preserving priors are applied

independently to each bed and stitched afterwards. Bed positions overlapped by 11 voxels. Phantom PET data were reconstructed using a single bed position.

Projection data are partitioned into staggered, view-interleaved subsets. For SHKEM we use $S = 9$ PET subsets, following the comparator study. For SVRG-based reconstructions we use 18 PET subsets for phantom data and 18 subsets per PET bed position for clinical data, giving 36 PET subsets in total for two-bed patient reconstructions.

8.5.2 SPECT forward model

For bremsstrahlung SPECT we use a geometric projector with an analytic collimator–detector response and a simulated additive correction term. SPECT images are discretised on a $128 \times 128 \times 128$ grid with isotropic voxel size 4.18 mm.

The collimator response is modelled as a depth-dependent Gaussian with linearly varying width, obtained by fitting an analytic collimator model [90]. Bremsstrahlung blur is modelled as an isotropic Gaussian with energy-dependent width fitted to ^{90}Y range simulations [18]; where multiple energy windows are used, effective blur parameters are obtained by count-weighted averaging across windows.

Scatter and septal penetration are incorporated via an additive term estimated by an outer reconstruct–simulate loop using Simulation of Imaging Nuclear Detectors (SIMIND) [91, 58]. We initialise this additive term using a learned estimate from a neural network [67]. The additive estimate is then refined by alternating reconstruction and SIMIND simulation for 10 outer iterations, using a damped update

$$b^{(k+1)} = \nu \hat{b}^{(k+1)} + (1 - \nu)b^{(k)}, \quad \nu = 0.2, \quad (8.22)$$

where $\hat{b}^{(k+1)}$ is the newly simulated additive estimate and $b^{(k)}$ is the previous estimate. The final additive estimate is fixed during the stochastic optimisation.

SPECT projection data are partitioned into staggered subsets, with $S = 12$ for both SVRG-based reconstructions and SHKEM.

8.5.3 Co-registration and warp operators

CTAC provides the anatomical reference for co-registration and for guidance via the μ -map. The shared reconstruction grid is the PET grid. The PET CTAC is therefore downsampled to the PET reconstruction grid and used both to define the attenuation map μ and as the anatomical reference space for the coupled prior, so W_1 is the identity on the shared grid.

The PET and SPECT CTAC volumes were retained in their native acquisition grids before registration. Since patient arm positioning differed between modalities (arms down in PET, arms up in SPECT), registration was restricted to the trunk. A trunk mask was obtained automatically from each CTAC using TotalSegmentator [TODO'totalsegmentator'2023] **TODO: Zenodo-managed bibliography entry pending for TotalSegmentator; replace TODO_totalsegmentator_2023 once imported.**, and the native-grid SPECT CTAC was registered to the PET CTAC using a rigid transformation followed by deformable registration initialised from the rigid solution. The resulting native-SPECT-to-PET deformation defines W_2 . The corresponding operator W_2^* in the data-fidelity term is the discrete adjoint of this resampling operator, not merely a matrix transpose; linear interpolation was used so that the forward and adjoint warp pair were implemented consistently.

8.6 Data and evaluation protocol

All quantitative image analysis is performed on PET reconstructions to preserve direct comparability with SHKEM, where PET is the target modality and SPECT provides guidance. SPECT reconstructions are shown for completeness but are not included in the quantitative endpoints. Reconstruction hyperparameters are tuned once using the clinical training patient (SIRT3) and then fixed for all phantom and clinical reconstructions. This single-patient tuning strategy avoids repeated adjustment on the evaluation cohort, but it does not constitute a robustness study of the selected weights. Throughout, dTV, w -TNV, and SHKEM are compared against the reference method w -dTNV.

8.6.1 Phantom data

A NEMA IEC Body Phantom with a sphere-to-background ratio of 8 was scanned at The Christie Hospital using a Siemens Biograph mCT PET/CT system and a GE Discovery 670 SPECT/CT system. PET was acquired over 15 h and SPECT over 15 min. From the 15 h PET acquisition, $B = 20$ statistically independent PET datasets equivalent to 15 min were generated by sinogram bootstrapping [92]. The SPECT energy window was 39.84–128.16 keV.

We quantify PET reproducibility across bootstraps via voxel-wise coefficient of variation (CoV), and PET quantification via recovery coefficients (RCs) for spheres of diameters 10 mm, 13 mm, 22 mm, 28 mm, and 37 mm. Three volume of interests (VOIs) were defined (Figure 8.1): a high-count (HC) VOI comprising five filled spheres (one sphere was excluded due to filling issues), a warm-background (BG) VOI comprising the phantom volume excluding the spheres plus a 10 mm margin and excluding the lung insert plus a 5 mm margin, and the cold lung insert (LI).




Figure 8.1: NEMA phantom analysis regions. Left: warm-background (BG), lung insert (LI), and high-count (HC) VOIs. Right: CTAC-derived attenuation map showing sphere edges, the lung insert, and the under-filled sphere excluded from the HC VOI.

Let $x_{b,j}^{(r)}$ be the reconstructed PET intensity at voxel j for method r in bootstrap realisation $b \in \{1, \dots, B\}$. Define the inter-bootstrap mean and standard deviation

$$\mu_j^{(r)} := \frac{1}{B} \sum_{b=1}^B x_{b,j}^{(r)}, \quad s_j^{(r)} := \sqrt{\frac{\sum_{b=1}^B \left(x_{b,j}^{(r)} - \mu_j^{(r)} \right)^2}{B-1}}, \quad (8.23)$$

and voxel-wise

$$\text{CoV}_j^{(r)} = \frac{s_j^{(r)}}{\mu_j^{(r)}}. \quad (8.24)$$

For a VOI $\mathcal{V}$ we report the VOI-averaged statistic

$$T_{\mathcal{V}}^{(r)} = \frac{1}{|\mathcal{V}|} \sum_{j \in \mathcal{V}} \text{CoV}_j^{(r)}. \quad (8.25)$$

For sphere s , let $\bar{C}_{b,s}^{(r)}$ denote the mean reconstructed intensity within the sphere and $\bar{C}_{b,\text{bkg}}^{(r)}$ the mean within BG for the same (r, b) . With true ratio $R_{\text{true}} = 8$, define

$$\text{RC}_{b,s}^{(r)} = \frac{\bar{C}_{b,s}^{(r)}}{\bar{C}_{b,\text{bkg}}^{(r)} R_{\text{true}}}. \quad (8.26)$$

To assess statistical significance of inter-bootstrap CoV, we compare each method r with the reference $r_0 = w\text{-dTNV}$ using a paired row-swap permutation test on $\log(T_{\mathcal{V}}^{(r)})$. Let $X_{\mathcal{V}}^{(r)} \in \mathbb{R}^{B \times |\mathcal{V}|}$ be the bootstrap-by-voxel stack restricted to VOI $\mathcal{V}$, and define the observed difference $\Delta = \log(T_{\mathcal{V}}^{(r_0)}) - \log(T_{\mathcal{V}}^{(r)})$. Under the null of exchangeability, each bootstrap index is independently swapped between the paired stacks with probability $1/2$ to form permuted differences Δ_{perm} . We use $N_{\text{perm}} = 5000$ permutations and compute the two-sided p-value

$$p = \frac{1 + \sum_{i=1}^{N_{\text{perm}}} \mathbf{1}(|\Delta_{\text{perm},i}| \geq |\Delta|)}{N_{\text{perm}} + 1}, \quad (8.27)$$

with Benjamini–Hochberg (BH) adjustment across methods within each VOI. No thresholding was applied to exclude near-zero voxels, so CoV estimates in the cold lung insert should be interpreted as a noise-realisation stability measure rather than a conventional percentage variability around a substantial mean signal.

Inference for RC uses a linear mixed model on

$$y_{b,r,s} := \log\left(\text{RC}_{b,s}^{(r)}\right), \quad b = 1, \dots, B, \quad (8.28)$$

with a random intercept for bootstrap replicate and fixed effects for method, sphere,

and their interaction:

$$y_{b,r,s} = \beta_0 + \alpha_r + \gamma_s + (\alpha\gamma)_{r,s} + u_b + \varepsilon_{b,r,s},$$

$$u_b \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \tau_b^2), \quad \varepsilon_{b,r,s} \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \tau^2). \quad (8.29)$$

For each sphere size, many-to-one contrasts compare each method to the reference $r_0 = w\text{-dTNV}$, with Benjamini–Hochberg adjustment across planned comparisons within that sphere. Since the model is fitted on the log scale, $\exp(\Delta_{r,s}) < 1$ indicates lower recovery than $w\text{-dTNV}$ for sphere s .

8.6.2 Clinical data

Clinical data were acquired at Oxford University Hospitals from 10 patients undergoing ^{90}Y SIRT: one patient was used for hyperparameter selection and nine for testing. PET was acquired on a GE Discovery 710 PET/CT system and SPECT on a GE Discovery NM/CT 670 system. Informed consent is not necessary for retrospective image reviews of this nature at Oxford University Hospitals. PET was acquired with two bed positions (15 min each) and SPECT with 15 min. Energy windows were 50–150 keV.

Tumour VOIs were delineated once per patient on the vendor Q.Clear $\beta\text{-4000}$ PET reconstruction using an automatic gradient-based method [93] and propagated unchanged across reconstructions so that differences reflected reconstruction effects rather than VOI definition. A homogeneous liver background VOI was drawn manually, avoiding motion and visible artefacts. The training-patient lesion and background VOIs are shown in Figure 8.2.

For patient p , lesion l , and method r , define tumour-to-background ratio (TBR)

$$\text{TBR}_{p,l,r} = \frac{\bar{C}_{p,l,r}}{\bar{C}_{p,\text{bkg},r}}, \quad (8.30)$$

and background noise (percentage CoV)

$$\text{CoV}_{p,r}(\%) = 100 \frac{s_{p,r}}{\bar{C}_{p,\text{bkg},r}}, \quad (8.31)$$




Figure 8.2: Training-patient VOIs. The five identified lesions and the background reference region are shown on the vendor PET reconstruction with registered attenuation maps. Top row: lesions 1–3. Bottom row: lesions 4–5 and background.

where $s_{p,r}$ is the within-background-VOI standard deviation.

We analyse $\log(\text{TBR})$ using a linear mixed-effects model with fixed method effects and random intercepts for patient and lesion-within-patient:

$$\begin{aligned} y_{p,l,r} &:= \log(\text{TBR}_{p,l,r}) = \beta_0 + \alpha_r + u_p + v_{p,l} + \varepsilon_{p,l,r}, \\ u_p &\stackrel{\text{iid}}{\sim} \mathcal{N}(0, \tau_p^2), \quad v_{p,l} \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \tau_l^2), \quad \varepsilon_{p,l,r} \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \tau^2). \end{aligned} \quad (8.32)$$

Background noise is analysed analogously at the patient level using

$$\begin{aligned} z_{p,r} &:= \log(\text{CoV}_{p,r}(\%)) = \beta_0 + \alpha_r + u_p + \varepsilon_{p,r}, \\ u_p &\stackrel{\text{iid}}{\sim} \mathcal{N}(0, \tau_{p,\text{CoV}}^2), \quad \varepsilon_{p,r} \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \psi^2). \end{aligned} \quad (8.33)$$

Planned many-to-one contrasts compare dTV, w -TNV, and SHKEM to w -dTNV, with BH adjustment across the three comparisons within each endpoint. Results are reported as ratios $\exp(\Delta)$ on the original scale, where values below one indicate lower TBR or lower background CoV than w -dTNV.

8.7 Results

8.7.1 Hyperparameter selection

Regularisation weights β_{PET} and β_{SPECT} in (8.7) were tuned on a single training patient by sweeping the TBR–CoV_{bkg} trade-off and selecting an operating point on the Pareto frontier [94] with background CoV approximately matched to the vendor reconstruction. The selected weights were then fixed for all phantom and clinical experiments:

Table 8.1: Selected regularisation weights (tuned on a single training patient) used for all phantom and clinical experiments.

Method	β_{PET}	β_{SPECT}
w -dTNV	0.01	0.02
w -TNV	0.0075	0.02
dTV	0.02	–

For SHKEM, hyperparameters followed Deidda *et al.* [45]; the PET subiteration displayed and analysed here (27) was chosen to approximately match the background CoV of the vendor reconstruction.

Figure 8.3 shows the training-patient trade-off curves for two representative lesions. Lesion 1 shows only a modest initial benefit from synergistic information and is largely insensitive to further increases in SPECT weighting until over-regularisation introduces negative bias. Lesion 2 shows a clearer synergistic benefit: increasing SPECT weighting shifts the TBR–CoV_{bkg} curve towards higher TBR at comparable background CoV up to a distinct elbow region, after which excessive coupling reduces TBR and/or increases background CoV. For lesion 2, the PET image obtained with w -dTNV also exhibited features consistent with a CTAC artefact; this is discussed in Section 8.8.

8.7.2 Phantom results

A representative reconstruction from a single bootstrapped PET realisation is shown in Figure 8.4 (PET and corresponding SPECT; for SHKEM, SPECT is the separately reconstructed guidance image; dTV does not reconstruct SPECT).




Figure 8.3: Training-patient hyperparameter sweep. Impact of varying PET weighting while holding SPECT regularisation constant (top row) and varying SPECT weighting while holding PET regularisation constant (bottom row) on TBR and CoV_{bkg} for lesions 1 and 2. The first and third columns show w -dTNV; the second and fourth show w -TNV.




Figure 8.4: Visual comparison of NEMA phantom reconstructions for a randomly chosen bootstrap. From left to right: w -dTNV, dTV, w -TNV, SHKEM (subiteration 27). PET reconstructions are shown on the top row and SPECT on the bottom row (SPECT guidance for SHKEM). Images are displayed in reconstruction intensity units, proportional to activity concentration.

Figure 8.5 and Table 8.2 summarise sphere recovery. Quantitative analysis demonstrated that w -dTNV achieved significantly higher recovery than dTV and w -TNV for all sphere sizes ≤ 22 mm. w -dTNV showed similar recovery to SHKEM in the larger spheres but significantly improved recovery for the two smallest spheres at the matched operating point.




Figure 8.5: Recovery coefficient (RC) for NEMA phantom spheres for w -dTNV, dTV, w -TNV and SHKEM (subiteration 27). Also shown are the equivalent RCs for a high-count, long-OSEM reconstruction. The grey line indicates perfect recovery. Statistical significance is reported for paired comparisons versus w -dTNV (Benjamini–Hochberg adjusted): * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Method	s	$\exp(\Delta_{r,s})$	95% CI	p_{BH}
dTV	37 mm	0.9113	[0.8001, 1.0380]	0.2472
	28 mm	0.8543	[0.7500, 0.9730]	0.0267
	22 mm	0.8165	[0.7168, 0.9300]	0.0034
	13 mm	0.7819	[0.6865, 0.8906]	0.00043
	10 mm	0.7706	[0.6766, 0.8777]	$< 10^{-4}$
w -TNV	37 mm	0.9108	[0.7997, 1.0374]	0.2472
	28 mm	0.8544	[0.7502, 0.9732]	0.0267
	22 mm	0.8090	[0.7103, 0.9215]	0.0034
	13 mm	0.7598	[0.6671, 0.8654]	0.0001
	10 mm	0.7384	[0.6483, 0.8410]	$< 10^{-4}$
SHKEM (27)	37 mm	0.9931	[0.8720, 1.1312]	0.9175
	28 mm	0.9739	[0.8551, 1.1093]	0.6906
	22 mm	0.9616	[0.8443, 1.0953]	0.5559
	13 mm	0.8723	[0.7659, 0.9936]	0.0397
	10 mm	0.7381	[0.6480, 0.8407]	$< 10^{-4}$

Table 8.2: Effect ratios $\exp(\Delta_{r,s})$ (comparator vs. baseline w -dTNV) with 95% confidence intervals and BH-adjusted p -values, reported by sphere diameter s . Values below one indicate lower recovery coefficient than w -dTNV.

Figure 8.6 summarises voxel-wise reproducibility across bootstraps. At the matched operating point, dTV and w -TNV yield lower voxel-wise CoV than

w -dTNV in the evaluated VOIs, consistent with stronger effective smoothing. SHKEM exhibits higher variability in background and high-count regions but improved stability in the cold lung insert. Sphere-level CoV trends were consistent with the voxel-wise analysis, with w -TNV and dTV generally providing the lowest variability across the evaluated sphere diameters and SHKEM the highest.




Figure 8.6: Mean voxel-wise coefficient of variation (CoV) across 20 bootstrapped PET realisations within each VOI. Statistical significance is reported for paired comparisons versus w -dTNV (BH-adjusted): * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

8.7.3 Clinical results

Figure 8.7 shows representative PET and SPECT reconstructions for the training patient. These images illustrate the operating-point behaviour used to select the fixed regularisation weights: w -dTNV improves lesion contrast where the SPECT and CTAC-derived guidance contain useful structural information, while the response is weaker where there is no clear anatomical correlate.

Across the test cohort, w -dTNV yields higher lesion TBR than dTV, w -TNV, and SHKEM at the matched operating point, while background CoV differences are smaller and method-dependent. SHKEM reconstructions appeared visually smoother and, as in the phantom, appeared more effective at suppressing cold-background noise:




Figure 8.7: Visual comparison of reconstruction methods for training patient SIRT3. Columns show w -dTNV, dTV, w -TNV and SHKEM (subiteration 27). The upper two quadrants show PET reconstructions; the lower two show SPECT reconstructions or guidance images. Images are displayed in reconstruction intensity units, proportional to activity concentration.




Figure 8.8: Clinical cohort summary: lesion-wise mean TBR (left) and patient-wise background CoV (right). Statistical significance is reported for comparisons versus w -dTNV(BH-adjusted): * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Method	$\exp(\Delta)$	95% CI	p_{BH}
dTV	0.802	[0.758, 0.849]	$< 10^{-4}$
w -TNV	0.915	[0.865, 0.968]	0.004
SHKEM (27)	0.935	[0.884, 0.990]	0.020

Table 8.3: Effect ratios $\exp(\Delta)$ with 95% confidence intervals and BH-adjusted p -values for mean tumour-to-background ratio (TBR), comparing each method to w -dTNV. Values below one indicate lower TBR than w -dTNV.

Method	$\exp(\Delta)$	95% CI	p_{BH}
dTV	1.183	[0.903, 1.549]	0.35
w -TNV	1.007	[0.833, 1.167]	0.984
SHKEM (27)	1.339	[0.933, 1.913]	0.093

Table 8.4: Effect ratios $\exp(\Delta)$ with 95% confidence intervals and BH-adjusted p -values for background CoV, comparing each method to w -dTNV. Values above one indicate higher background CoV than w -dTNV.

8.8 Discussion

The phantom results indicate that functional coupling alone (w -TNV) is insufficient to recover small-sphere boundaries at the matched operating point, and that anatomical guidance (explicit in w -dTNV, implicit in SHKEM through CT-informed kernel features) is beneficial for boundary localisation. The recovery gain for w -dTNV relative to w -TNV and dTV is most pronounced for the smallest spheres, consistent with guidance stabilising edge geometry when PET statistics are poor and SPECT resolution is limited.

The reproducibility analysis highlights a trade-off: dTV and w -TNV provide lower voxel-wise CoV than w -dTNV at the selected operating point, consistent with stronger effective smoothing. All three variational methods outperform SHKEM in background and high-count repeatability, whereas SHKEM is most stable in the cold lung insert. A plausible mechanism is that SHKEM kernel weights in near-zero regions favour local averaging among consistently low voxels, so “zeros reinforce zeros” under positivity constraints, improving cold-region stability while not necessarily improving reproducibility elsewhere.

Notably, dTV and w -TNV produced broadly similar PET appearances in the phantom, despite dTV including μ -directional guidance and w -TNV including

PET/SPECT coupling. This suggests that, in this experiment, neither CTAC-derived directional guidance alone nor isotropic PET/SPECT coupling alone was sufficient to recover the smallest structures. The improved recovery observed with w -dTNN is therefore consistent with a complementary mechanism: the higher-count SPECT image provides additional functional support for the location of activity transitions, while the μ -derived directional field reduces penalisation of gradients aligned with anatomical boundaries.

Clinically, lesion responses to synergy and directional guidance are heterogeneous. Lesions with clear anatomical correlates (e.g. calcification or liver capsule boundaries visible in the μ -map) tend to benefit more from μ -guided anisotropy and PET/SPECT coupling, while lesions lacking such correlates may show limited benefit. This heterogeneity is consistent with the design goal of w -dTNN: CTAC guidance steers *orientation* preferences, but cannot invent edges unsupported by the functional data.

For example, lesion 2 in the training patient shows substantial improvement in the TBR–CoV trade-off as SPECT weighting is increased (Figure 8.3). Post-hoc examination indicates that this lesion lies adjacent to a calcified region visible in the μ -map, providing a sharp anatomical edge that coincides with the functional boundary. Lesion 3 also improves under both w -dTNN and w -TNN; this lesion abuts the liver on three sides, where the soft-tissue interface generates reasonably strong anatomical contrast. For w -dTNN, the μ -map liver boundary provides explicit directional guidance aligned with the functional uptake margin, while for w -TNN the SPECT boundary may be sufficiently well defined to support PET/SPECT coupling even without directional CTAC guidance.

Conversely, lesion 1 shows little benefit from increasing SPECT weighting until excessive weighting introduces negative bias. Visual inspection shows no obvious anatomical correlate at the tumour margin, and the PET and SPECT images appear already well aligned. The reason for this lesion’s limited response to synergistic reconstruction remains unclear.

A practical limitation is sensitivity to CTAC artefacts. Since ξ_n in (8.4) depends




Figure 8.9: Training-patient CTAC-derived guidance. The CTAC-derived attenuation map μ (top row; cm^{-1}) and magnitude of the w -dTNV directional weighting field $\|\xi\|$ (bottom row) are shown for lesions 1 and 2. CTAC streak artefacts are visible in μ and are reproduced in $\|\xi\|$, providing a pathway for artefact transfer into the emission reconstruction.

on $D\mu$, structured CT artefacts (e.g. arms-down streaking) can imprint into $\|\xi\|$ and therefore into the directional weighting field, providing a pathway for artefact transfer into the emission reconstructions. In the training patient, the low-dose PET CTAC was acquired with arms alongside the torso, introducing streaking and beam-hardening artefacts in the μ -map. The stabiliser η mitigates amplification of small gradients but cannot reliably suppress coherent high-contrast artefacts. SHKEM anatomical guidance was derived from the SPECT attenuation map, acquired with arms raised, rather than from the low-dose PET CTAC used to construct the w -dTNV directional weights; this difference in guidance input may contribute to the absence of streak-related artefacts in the SHKEM reconstructions.

The contrast with the NEMA phantom is instructive (Figure 8.10). In the phantom, the CTAC is largely free of coherent streak artefacts and the resulting $\|\xi\|$ follows the true boundaries and insert edges. In this setting, directional smoothing attenuates small high-frequency fluctuations without introducing structured bias.

These observations motivate future work on robust direction-field estimation and artefact-aware guidance strategies. From a practical perspective, if a low-dose CTAC is used to provide directional information, acquisition protocols that minimise artefacts, such as arms raised where feasible, are preferable.




Figure 8.10: NEMA phantom, central axial slice. CTAC-derived attenuation map μ (left; cm^{-1}) and magnitude of the w -dTNV directional weighting field $\|\xi\|$ (right). With a high-quality CTAC free of prominent streak artefacts, $\|\xi\|$ concentrates around true object boundaries and insert edges, enabling stable μ -map-guided directional smoothing.

To preserve direct comparability with SHKEM, regularisation weights were selected once on the clinical training case and quantitative evaluation focused on PET. This means the robustness of the chosen weights across acquisition protocols and lesion presentations was not fully characterised. The jointly reconstructed SPECT images were assessed qualitatively only. Although they appeared visually less noisy while maintaining lesion conspicuity, a rigorous assessment of SPECT quantification would require modality-specific lesion delineation and evaluation protocols. The sensitivity of w -dTNV to residual PET/SPECT/CT registration error also remains an important topic for future work.

More broadly, this chapter focuses on the favourable case in which PET and SPECT image the same underlying ^{90}Y microsphere distribution. The same framework may be extendable to settings with non-identical functional distributions but shared anatomy, such as planning–treatment image pairs ($^{99\text{m}}\text{Tc}$ -MAA versus post-treatment ^{90}Y), multi-nuclide imaging, or multi-timepoint studies. In such cases, the directional/anatomical component remains well motivated, but the inter-modality coupling strength may need to be relaxed or made spatially adaptive to avoid transferring modality-specific or temporally evolving uptake features.

8.9 Conclusion

This chapter validated the w - d TNV prior of Chapter 6 in joint ^{90}Y PET/SPECT reconstruction with CTAC-derived directional guidance. The key modelling choice is to use CTAC information directionally (via $\Pi_{m,n}$) while preserving a two-modality Jacobian, retaining the computational advantages of the 2×2 Gram structure (Chapter 6) and avoiding unstable intensity coupling to CT.

On a bootstrapped NEMA IEC phantom, w - d TNV improves small-object recovery relative to w -TNV, d TV, and a sequential hybrid kernel baseline at a matched noise operating point, with significant gains over the isotropic and non-synergistic variational baselines. Repeatability improves relative to SHKEM in the background and high-count regions, although it is reduced relative to the more strongly smoothing w -TNV and d TV baselines and SHKEM remains most stable in the cold lung insert. In a post-SIRT patient cohort of 45 lesions in nine patients, w - d TNV yields significantly higher lesion TBR than the comparators at approximately matched background CoV. Lesion responses are heterogeneous: lesions with clear anatomical correlates, such as calcification or liver boundaries, benefit most from synergy and μ -map-informed edge alignment, whereas lesions without corresponding anatomical structure show more limited improvement. These findings support CTAC-guided synergistic variational coupling as a practical route towards more reliable post-SIRT ^{90}Y activity imaging for personalised dosimetry, while highlighting the need for careful handling of CT artefacts and PET/SPECT/CT registration in clinical translation.

Appendix A

`simind-python-connector`: Python Workflows for SIMIND Simulations

This appendix documents the `simind-python-connector` package used to support reproducible SIMIND-based simulation workflows in this thesis. The inspected package version was 1.0.1, released on 16 April 2026 and archived on Zenodo with DOI 10.5281/zenodo.19605413. The public repository is available at GitHub.

TODO: Replace the DOI URL above with the Zenodo-managed citation key for DOI 10.5281/zenodo.19605413 once the automatic bibliography import has added it.

A.1 Purpose and Scope

SIMIND is a Monte Carlo code for simulating clinical SPECT scintillation-camera imaging and related system effects [64, 91]. It is especially useful in this thesis because ^{90}Y bremsstrahlung SPECT requires careful modelling of scatter, attenuation, collimator response, and septal penetration (Chapters 4 and 8). The role of `simind-python-connector` is not to replace SIMIND, but to make SIMIND easier to drive from Python and easier to connect to reconstruction software.

The package is independent of SIMIND and is not affiliated with, endorsed by, or maintained by the SIMIND project or Lund University. SIMIND itself is not distributed with the Python package; a separately licensed SIMIND installation must be available on the local machine or inside the configured container environment.

This separation is useful for reproducibility because the open Python workflow can be version controlled while the licensed simulation executable remains an external dependency.

The package provides two main capabilities. First, `SimindPythonConnector` offers a direct Python interface for configuring voxel phantoms, writing SIMIND input files, setting runtime switches, running SIMIND, and loading projection outputs as NumPy arrays. Second, adaptor classes translate the resulting data into reconstruction-library-native representations for Software for Tomographic Image Reconstruction (STIR), Synergistic Image Reconstruction Framework (SIRF), and PyTomography workflows [84, 95]. The package therefore sits between simulation and reconstruction: SIMIND produces projection data, while reconstruction is still performed by the target reconstruction framework.

TODO: Add Zenodo-managed BibTeX entry for PyTomography if the PyTomography adaptor is retained in the final cited text.

TODO: Add Zenodo-managed BibTeX entry for the Schneider CT-density conversion paper if this utility is discussed beyond software documentation.

A.2 Connector Workflow

The core workflow is deliberately explicit. A user constructs a connector from a base SIMIND configuration, output directory, and output prefix; supplies a source image and attenuation or density information; specifies energy windows; applies runtime switches such as radionuclide, collimator, photon multiplier, and random seed; and then calls `run()`. The connector writes SIMIND-compatible source and density files, saves the generated configuration, executes SIMIND, and loads the resulting Interfile projection data.

For simple Python use, the package can be installed from PyPI with `pip install simind-python-connector`; optional plotting dependencies are provided through `simind-python-connector[examples]`. The base package depends only on standard Python scientific and medical-image utilities such as NumPy, pydicom, and PyYAML. STIR, SIRF, PyTomography, and SIMIND are optional

Table A.1: Main `simind-python-connector` interfaces relevant to thesis workflows.

Interface	Role
<code>SimindPythonConnector</code>	Backend-agnostic connector that runs SIMIND and returns projection outputs as NumPy arrays with header and data-file paths.
<code>RuntimeOperator</code>	Container for per-run SIMIND runtime switches and optional orbit-file selection.
<code>StirSimindAdaptor</code>	Converts connector inputs and outputs for STIR-based reconstruction workflows.
<code>SirfSimindAdaptor</code>	Provides the analogous adaptor path for SIRF, which builds on STIR for PET/SPECT reconstruction.
<code>PyTomographySimindAdaptor</code>	Accepts torch tensors, executes the SIMIND connector path, and returns PyTomography-compatible projection tensors and headers.
<code>hu_to_density_schneider</code>	Converts CT Hounsfield units to material density using the packaged Schneider-style tissue lookup.

or external dependencies and are needed only for the corresponding adaptor or simulation paths.

A.3 Geometry and Data Conventions

A central difficulty in connecting simulation and reconstruction software is that array-order, projection-axis, and unit conventions differ between packages. The connector documentation therefore makes these conventions explicit. STIR and SIRF image arrays are handled as (z, y, x) . PyTomography object-space tensors are handled as (x, y, z) , while SIMIND input image files are written internally in (z, y, x) order. Projection arrays may also differ: raw Interfile loading is exposed as $(views, bins, axial)$ when no singleton time-of-flight axis is present, whereas PyTomography SIMIND projections are used in $(theta, r, z)$ order.

The connector also handles the common unit conversion between reconstruction libraries and SIMIND. SIMIND geometry parameters are configured in centimetres, whereas STIR/SIRF and many Python reconstruction workflows express voxel sizes in millimetres. The connector therefore derives the SIMIND voxel-size switch from the supplied voxel size and writes corresponding source and density geometry values

into the generated configuration.

These details matter for the work in this thesis because an apparently small axis flip or millimetre-to-centimetre error would change the simulated collimator response, attenuation path lengths, or reconstructed orientation. The package examples and tests therefore include explicit geometry guardrails: fixed toy configurations, centre-of-mass checks, view-index checks, and backend-specific adaptor examples.

A.4 Validation and Reproducibility

The package separates lightweight tests from tests that need heavyweight external dependencies. Pure Python unit tests cover configuration handling, connector behaviour, Interfile parsing, density conversion, and optional-backend guards. Backend-dependent tests are marked separately for STIR, SIRT, PyTomography, and SIMIND, and are skipped automatically when the corresponding dependency is unavailable. This is important for continuous integration because most public CI runners cannot include a licensed SIMIND executable.

Dedicated Docker environments are supplied for STIR, SIRT, and PyTomography. These containers isolate reconstruction-library dependencies and can run either connector-only validation or SIMIND-dependent validation when a local SIMIND executable is mounted under the expected repository layout. The repository helper scripts look for the executable at `./simind/simind` and the SIMIND data directory at `./simind/smc_dir/`. This mirrors the thesis requirement that the simulation code remains external while the workflow around it remains inspectable.

The included examples cover basic NumPy simulations, runtime-switch comparisons, multi-window simulations, SCATTWIN/PENETRATE scoring comparisons, Schneider-style density conversion, adaptor-based OSEM reconstructions, and DICOM-driven setup paths. Together these examples define the intended reproducibility contract: the connector should make SIMIND runs scriptable, recordable, and portable across reconstruction frameworks, while leaving reconstruction algorithm choices to the software stack used downstream.

Appendix B

Colophon

This is a description of the tools you used to make your thesis. It helps people make future documents, reminds you, and looks good.

(example) This document was set in the Times Roman typeface using L^AT_EX and BibT_EX, composed with a text editor.

References

- [1] Francis X. Sundram and John R. Buscombe. “Selective internal radiation therapy for liver tumours”. In: *Clinical Medicine* 17.5 (Oct. 2017), p. 449. ISSN: 14734893. DOI: 10.7861/CLINMEDICINE.17-5-449. URL: <https://pubmed.ncbi.nlm.nih.gov/articles/PMC6301920/> (visited on 08/21/2025).
- [2] Lidia Strigari et al. “The evidence base for the use of internal dosimetry in the clinical practice of molecular radiotherapy”. en. In: *European Journal of Nuclear Medicine and Molecular Imaging* 41.10 (2014), pp. 1976–1988. ISSN: 1619-7089. DOI: 10.1007/s00259-014-2824-5. URL: <https://doi.org/10.1007/s00259-014-2824-5>.
- [3] Riad Salem et al. “Clinical and dosimetric considerations for Y90: recommendations from an international multidisciplinary working group”. en. In: *European Journal of Nuclear Medicine and Molecular Imaging* 46.8 (July 2019), pp. 1695–1704. ISSN: 1619-7089. DOI: 10.1007/s00259-019-04340-5. URL: <https://doi.org/10.1007/s00259-019-04340-5> (visited on 02/01/2026).
- [4] Marlene Reincke et al. “Hepatic decompensation after transarterial radioembolization: A retrospective analysis of risk factors and outcome in patients with hepatocellular carcinoma”. en. In: *Hepatology Communications* 6.11 (2022). eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1002/hep4.2072>, pp. 3223–3233. ISSN: 2471-254X. DOI: 10.1002/hep4.2072. URL: <https://onlinelibrary.wiley.com/doi/abs/10.1002/hep4.2072> (visited on 02/01/2026).

- [5] S. Cheenu Kappadath et al. “Quantitative Evaluation of 90Y-PET/CT and 90Y-SPECT/CT-based Dosimetry following Yttrium-90 Radioembolization”. In: *Medical physics* 51.9 (Sept. 2024), pp. 6061–6074. ISSN: 0094-2405. DOI: 10.1002/mp.17175. URL: <https://pmc.ncbi.nlm.nih.gov/articles/PMC11731408/> (visited on 02/01/2026).
- [6] Jonathan Gear et al. “EANM enabling guide: how to improve the accessibility of clinical dosimetry”. en. In: *European Journal of Nuclear Medicine and Molecular Imaging* 50.7 (June 2023), pp. 1861–1868. ISSN: 1619-7089. DOI: 10.1007/s00259-023-06226-z. URL: <https://doi.org/10.1007/s00259-023-06226-z> (visited on 02/01/2026).
- [7] Bing Bai, James Bading, and Peter S Conti. “Tumor Quantification in Clinical Positron Emission Tomography”. In: *Theranostics* 3.10 (2013), pp. 787–801. ISSN: 1838-7640. DOI: 10.7150/thno.5629. URL: <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC3840412/>.
- [8] LauraMay Davis et al. “Personalisation of Molecular Radiotherapy through Optimisation of Theragnostics”. en. In: *Journal of Personalized Medicine* 10.4 (2020), p. 174. ISSN: 2075-4426. DOI: 10.3390/jpm10040174. URL: <https://www.mdpi.com/2075-4426/10/4/174>.
- [9] Simon R. Arridge, Matthias J. Ehrhardt, and Kris Thielemans. “(An overview of) Synergistic reconstruction for multimodality/multichannel imaging methods”. In: *Philosophical Transactions of the Royal Society A* 379.2200 (June 2021). ISSN: 1364503X. DOI: 10.1098/RSTA.2020.0205. URL: <https://royalsocietypublishing.org/doi/10.1098/rsta.2020.0205> (visited on 08/21/2025).
- [10] Ronald Boellaard et al. “FDG PET/CT: EANM procedure guidelines for tumour imaging: version 2.0”. en. In: *European Journal of Nuclear Medicine and Molecular Imaging* 42.2 (Feb. 2015), pp. 328–354. ISSN: 1619-7089. DOI: 10.1007/s00259-014-2961-x. URL: <https://doi.org/10.1007/s00259-014-2961-x> (visited on 01/26/2026).

- [11] Simon R. Cherry et al. *Physics in nuclear medicine*. Vol. 3. Saunders, 2003. URL: <https://jnm.snmjournals.org/content/jnumed/54/7/1168.2.full-text.pdf> (visited on 01/16/2026).
- [12] SD Bass. “Colloquium: Positronium physics and biomedical applications — Rev. Mod. Phys.” In: *Reviews of Modern Physics* 95.2 (2023), p. 021002. URL: <https://journals.aps.org/rmp/abstract/10.1103/RevModPhys.95.021002> (visited on 01/16/2026).
- [13] Kengo Shibuya et al. “Annihilation photon acollinearity in PET: volunteer and phantom FDG studies”. en. In: *Physics in Medicine & Biology* 52.17 (Aug. 2007), p. 5249. ISSN: 0031-9155. DOI: 10.1088/0031-9155/52/17/010. URL: <https://doi.org/10.1088/0031-9155/52/17/010> (visited on 01/26/2026).
- [14] Hal O. Anger. “Scintillation camera”. In: *Review of scientific instruments* 29.1 (1958). ISBN: 0034-6748, pp. 27–33.
- [15] National Institute for Health and Care Excellence (NICE). *Selective internal radiation therapies for treating hepatocellular carcinoma [TA688]*. Tech. rep. Mar. 2021. URL: <https://www.nice.org.uk/guidance/ta688> (visited on 08/21/2025).
- [16] National Institute for Health and Care Excellence (NICE). “Selective internal radiation therapy for neuroendocrine tumours that have metastasised to the liver [IPG786]”. In: (May 2024). URL: <https://www.nice.org.uk/guidance/ipg786> (visited on 08/21/2025).
- [17] Douglas J. Simpkin, S. James Cullom, and T. Rockwell Mackie. “The spatial and energy dependence of bremsstrahlung production about beta point sources in H₂O”. In: *Medical Physics* 19.1 (Jan. 1992), pp. 105–114. ISSN: 2473-4209. DOI: 10.1118/1.596884. URL: <https://onlinelibrary.wiley.com/doi/full/10.1118/1.596884> (visited on 08/21/2025).

- [18] Erwann Rault et al. “Fast simulation of yttrium-90 bremsstrahlung photons with GATE”. en. In: *Medical Physics* 37.6Part1 (2010), pp. 2943–2950. ISSN: 2473-4209. DOI: 10.1118/1.3431998. URL: <https://onlinelibrary.wiley.com/doi/abs/10.1118/1.3431998>.
- [19] S Jan et al. “GATE: a simulation toolkit for PET and SPECT”. en. In: *Physics in Medicine & Biology* 49.19 (2004), p. 4543. ISSN: 0031-9155. DOI: 10.1088/0031-9155/49/19/007. URL: <https://dx.doi.org/10.1088/0031-9155/49/19/007>.
- [20] Sophia Pells et al. “Quantitative validation of Monte Carlo SPECT simulation: application to a Mediso AnyScan GATE simulation”. In: *EJNMMI Physics* 10.1 (Dec. 2023), pp. 1–25. ISSN: 21977364. DOI: 10.1186/s40658-023-00581-4/FIGURES/11. URL: <https://link.springer.com/articles/10.1186/s40658-023-00581-4> (visited on 07/30/2025).
- [21] R. G. Selwyn et al. “A new internal pair production branching ratio of ^{90}Y : The development of a non-destructive assay for ^{90}Y and ^{90}Sr ”. In: *Applied Radiation and Isotopes* 65.3 (Mar. 2007), pp. 318–327. ISSN: 0969-8043. DOI: 10.1016/J.APRADISO.2006.08.009. (Visited on 08/21/2025).
- [22] Bhupinder Singh et al. “Attenuation artifact, attenuation correction, and the future of myocardial perfusion SPECT”. en. In: *Journal of Nuclear Cardiology* 14.2 (Mar. 2007), pp. 153–164. ISSN: 1532-6551. DOI: 10.1016/j.nuclcard.2007.01.037. URL: <https://doi.org/10.1016/j.nuclcard.2007.01.037> (visited on 02/01/2026).
- [23] James A. Patton and Timothy G. Turkington. “SPECT/CT Physical Principles and Attenuation Correction”. en. In: *Journal of Nuclear Medicine Technology* 36.1 (Mar. 2008), pp. 1–10. ISSN: 0091-4916, 1535-5675. DOI: 10.2967/jnmt.107.046839. URL: <https://tech.snmjournals.org/content/36/1/1> (visited on 02/01/2026).
- [24] A P Dempster, N M Laird, and D B Rubin. “Maximum Likelihood from Incomplete Data Via the EM Algorithm”. en. In: *Journal of the Royal Sta-*

- tistical Society: Series B (Methodological)* 39.1 (1977), pp. 1–22. ISSN: 00359246. DOI: 10.1111/j.2517-6161.1977.tb01600.x. URL: <https://onlinelibrary.wiley.com/doi/10.1111/j.2517-6161.1977.tb01600.x>.
- [25] L. A. Shepp and Y. Vardi. “Maximum Likelihood Reconstruction for Emission Tomography”. In: *IEEE Transactions on Medical Imaging* 1.2 (Oct. 1982), pp. 113–122. ISSN: 1558-254X. DOI: 10.1109/TMI.1982.4307558. URL: <https://ieeexplore.ieee.org/document/4307558> (visited on 12/04/2025).
- [26] Eugene Veklerov and Jorge Llacer. “Stopping Rule for the MLE Algorithm Based on Statistical Hypothesis Testing”. In: *IEEE Transactions on Medical Imaging* 6.4 (Dec. 1987), pp. 313–319. ISSN: 1558-254X. DOI: 10.1109/TMI.1987.4307849. URL: <https://ieeexplore.ieee.org/document/4307849/> (visited on 12/04/2025).
- [27] C. F. Jeff Wu. “On the Convergence Properties of the EM Algorithm”. In: *The Annals of Statistics* 11.1 (1983), pp. 95–103. ISSN: 0090-5364. URL: <https://www.jstor.org/stable/2240463> (visited on 01/19/2026).
- [28] Leonid I. Rudin, Stanley Osher, and Emad Fatemi. “Nonlinear total variation based noise removal algorithms”. In: *Physica D: Nonlinear Phenomena* 60.1 (Nov. 1992), pp. 259–268. ISSN: 0167-2789. DOI: 10.1016/0167-2789(92)90242-F. URL: <https://www.sciencedirect.com/science/article/pii/016727899290242F> (visited on 12/04/2025).
- [29] J Nuyts et al. “A concave prior penalizing relative differences for maximum-a-posteriori reconstruction in emission tomography”. In: *IEEE Transactions on Nuclear Science* 49.1 (2002), pp. 56–60. ISSN: 1558-1578. DOI: 10.1109/TNS.2002.998681.
- [30] J.E. Bowsher et al. “Utilizing MRI information to estimate F18-FDG distributions in rat flank tumors”. In: *IEEE Symposium Conference Record Nuclear Science 2004*. Vol. 4. Oct. 2004, 2488–2492 Vol. 4. DOI: 10.1109/NSSMIC.

- 2004 . 1462760. URL: <https://ieeexplore.ieee.org/abstract/document/1462760> (visited on 12/04/2025).
- [31] Matthias Joachim Ehrhardt and Simon R Arridge. “Vector-Valued Image Processing by Parallel Level Sets”. In: *IEEE Transactions on Image Processing* 23.1 (2014), pp. 9–18. ISSN: 1941-0042. DOI: 10.1109/TIP.2013.2277775.
- [32] Guobao Wang and Jinyi Qi. “PET Image Reconstruction Using Kernel Method”. In: *IEEE transactions on medical imaging* 34.1 (2015), pp. 61–71. ISSN: 0278-0062. DOI: 10.1109/TMI.2014.2343916. URL: <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC4280333/>.
- [33] Will Hutchcroft et al. “Anatomically-aided PET reconstruction using the kernel method”. en. In: *Physics in Medicine and Biology* 61.18 (2016), pp. 6668–6683. ISSN: 0031-9155, 1361-6560. DOI: 10.1088/0031-9155/61/18/6668. URL: <https://iopscience.iop.org/article/10.1088/0031-9155/61/18/6668>.
- [34] Kevin M. Holt. “Total Nuclear Variation and Jacobian Extensions of Total Variation for Vector Fields”. In: *IEEE Transactions on Image Processing* 23.9 (Sept. 2014), pp. 3975–3989. ISSN: 10577149. DOI: 10.1109/TIP.2014.2332397. (Visited on 08/21/2025).
- [35] David S. Rigie and Patrick J. La Rivière. “Joint reconstruction of multi-channel, spectral CT data via constrained total nuclear variation minimization”. In: *Physics in Medicine and Biology* 60.5 (Feb. 2015). arXiv: 1409.7712, pp. 1741–1762. ISSN: 13616560. DOI: 10.1088/0031-9155/60/5/1741. (Visited on 08/21/2025).
- [36] Florian Knoll et al. “Joint MR-PET Reconstruction Using a Multi-Channel Image Regularizer”. In: *IEEE Transactions on Medical Imaging* 36.1 (Jan. 2017), pp. 1–16. ISSN: 1558-254X. DOI: 10.1109/TMI.2016.2564989. URL: <https://ieeexplore.ieee.org/document/7466848> (visited on 12/04/2025).

- [37] Sangtae Ahn and J A Fessler. “Globally convergent ordered subsets algorithms: application to tomography”. en. In: *2001 IEEE Nuclear Science Symposium Conference Record (Cat. No.01CH37310)*. IEEE, 2002, pp. 1064–1068. ISBN: 978-0-7803-7324-2. DOI: 10.1109/NSSMIC.2001.1009736. URL: <http://ieeexplore.ieee.org/document/1009736/>.
- [38] Matthias J. Ehrhardt, Zeljko Kereta, and Georg Schramm. “Fast PET reconstruction with variance reduction and prior-aware preconditioning”. English. In: *Frontiers in Nuclear Medicine* 5 (Sept. 2025). ISSN: 2673-8880. DOI: 10.3389/fnume.2025.1641215. URL: <https://www.frontiersin.org/journals/nuclear-medicine/articles/10.3389/fnume.2025.1641215/full> (visited on 01/16/2026).
- [39] Robert Twyman et al. “An Investigation of Stochastic Variance Reduction Algorithms for 3D Penalised PET Image Reconstruction”. en. In: (2022). DOI: 10.36227/techrxiv.19497680.v1. URL: https://www.techrxiv.org/articles/preprint/An_Investigation_of_Stochastic_Variance_Reduction_Algorithms_for_3D_Penalised_PET_Image_Reconstruction/19497680/1.
- [40] Stuart Geman and Donald Geman. “Stochastic Relaxation, Gibbs Distributions, and the Bayesian Restoration of Images”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* PAMI-6.6 (Nov. 1984), pp. 721–741. ISSN: 1939-3539. DOI: 10.1109/TPAMI.1984.4767596. URL: <https://ieeexplore.ieee.org/abstract/document/4767596> (visited on 12/29/2025).
- [41] Elias Jonsson, Sung-cheng Huang, and Tony Chan. “Total variation regularization in positron emission tomography”. In: (1998).
- [42] V.Y. Panin, G.L. Zeng, and G.T. Gullberg. “Total variation regulated EM algorithm [SPECT reconstruction]”. In: *IEEE Transactions on Nuclear Science* 46.6 (Dec. 1999), pp. 2202–2210. ISSN: 1558-1578. DOI: 10.1109/23.

819305. URL: <https://ieeexplore.ieee.org/abstract/document/819305> (visited on 12/04/2025).
- [43] Matthias J. Ehrhardt et al. “PET Reconstruction with an Anatomical MRI Prior Using Parallel Level Sets”. In: *IEEE Transactions on Medical Imaging* 35.9 (Sept. 2016), pp. 2189–2199. ISSN: 1558254X. DOI: 10.1109/TMI.2016.2549601. (Visited on 07/31/2025).
- [44] Daniel Deidda et al. “Hybrid PET/MR Kernelised Expectation Maximisation Reconstruction for Improved Image-Derived Estimation of the Input Function from the Aorta of Rabbits”. In: *Contrast Media & Molecular Imaging* 2019 (Jan. 2019), p. 3438093. ISSN: 1555-4309. DOI: 10.1155/2019/3438093. URL: <https://pmc.ncbi.nlm.nih.gov/articles/PMC6360049/> (visited on 10/26/2025).
- [45] Daniel Deidda et al. “Triple modality image reconstruction of PET data using SPECT, PET, CT information increases lesion uptake in images of patients treated with radioembolization with 90Y micro-spheres”. en. In: *EJNMMI Physics* 10.1 (May 2023), p. 30. ISSN: 2197-7364. DOI: 10.1186/s40658-023-00549-4. URL: <https://doi.org/10.1186/s40658-023-00549-4> (visited on 12/09/2024).
- [46] Peter Blomgren and Tony F. Chan. “Color TV: Total variation methods for restoration of vector-valued images”. In: *IEEE Transactions on Image Processing* 7.3 (1998), pp. 304–309. ISSN: 10577149. DOI: 10.1109/83.661180. (Visited on 08/21/2025).
- [47] P J Green. “Bayesian reconstructions from emission tomography data using a modified EM algorithm”. eng. In: *IEEE transactions on medical imaging* 9.1 (1990), pp. 84–93. ISSN: 0278-0062. DOI: 10.1109/42.52985. URL: <http://www.ncbi.nlm.nih.gov/pubmed/18222753>.
- [48] Larry Zeng. “Green’s one-step-late algorithm dose not work for SPECT with attenuation correction”. In: *15th International Meeting on Fully Three-Dimensional Image Reconstruction in Radiology and Nuclear Medicine*.

- Vol. 11072. SPIE, 2019, pp. 395–399. DOI: 10.1117/12.2534946. URL: <https://www.spiedigitallibrary.org/conference-proceedings-of-spie/11072/110722A/Greens-one-step-late-algorithm-dose-not-work-for-SPECT/10.1117/12.2534946.full>.
- [49] C. G. BROYDEN. “The Convergence of a Class of Double-rank Minimization Algorithms: 2. The New Algorithm”. In: *IMA Journal of Applied Mathematics* 6.3 (Sept. 1970), pp. 222–231. ISSN: 0272-4960. DOI: 10.1093/imamat/6.3.222. URL: <https://doi.org/10.1093/imamat/6.3.222> (visited on 02/04/2026).
- [50] R. Fletcher. “A new approach to variable metric algorithms”. In: *The Computer Journal* 13.3 (Jan. 1970), pp. 317–322. ISSN: 0010-4620. DOI: 10.1093/comjnl/13.3.317. URL: <https://doi.org/10.1093/comjnl/13.3.317> (visited on 02/04/2026).
- [51] Donald Goldfarb. “A family of variable-metric methods derived by variational means”. en. In: *Mathematics of Computation* 24.109 (1970), pp. 23–26. ISSN: 0025-5718, 1088-6842. DOI: 10.1090/S0025-5718-1970-0258249-6. URL: <https://www.ams.org/mcom/1970-24-109/S0025-5718-1970-0258249-6/> (visited on 02/04/2026).
- [52] D. F. Shanno. “Conditioning of quasi-Newton methods for function minimization”. en. In: *Mathematics of Computation* 24.111 (1970), pp. 647–656. ISSN: 0025-5718, 1088-6842. DOI: 10.1090/S0025-5718-1970-0274029-X. URL: <https://www.ams.org/mcom/1970-24-111/S0025-5718-1970-0274029-X/> (visited on 02/04/2026).
- [53] Dong C. Liu and Jorge Nocedal. “On the limited memory BFGS method for large scale optimization”. en. In: *Mathematical Programming* 45.1 (Aug. 1989), pp. 503–528. ISSN: 1436-4646. DOI: 10.1007/BF01589116. URL: <https://doi.org/10.1007/BF01589116> (visited on 02/04/2026).
- [54] Stephen Boyd. “Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers”. en. In: *Foundations and Trends®*

- in Machine Learning* 3.1 (2010), pp. 1–122. ISSN: 1935-8237, 1935-8245. DOI: 10.1561/22000000016. URL: <http://www.nowpublishers.com/article/Details/MAL-016>.
- [55] Yu-Jung Tsai et al. “Fast Quasi-Newton Algorithms for Penalized Reconstruction in Emission Tomography and Further Improvements via Preconditioning”. In: *IEEE Transactions on Medical Imaging* 37 (2018), pp. 1000–1010. DOI: 10.1109/TMI.2017.2786865. URL: https://www.researchgate.net/publication/322058547_Fast_Quasi-Newton_Algorithms_for_Penalized_Reconstruction_in_Emission_Tomography_and_Further_Improvements_via_Preconditioning.
- [56] Andrew J. Woerner and Guy E. Johnson. “Advances in Y-90 radioembolization for the treatment of hepatocellular carcinoma”. en. In: *Hepatoma Research* 8.0 (Jan. 2022), N/A–N/A. ISSN: ISSN 2454-2520 (Online); ISSN 2394-5079 (Print). DOI: 10.20517/2394-5079.2021.122. URL: <https://www.oaepublish.com/articles/2394-5079.2021.122> (visited on 01/20/2026).
- [57] Marie-Martine Bé and Bureau international des poids et mesures, eds. *Table of radionuclides*. en. Monographie BIPM 5. Sèvres: BIPM, 2004. ISBN: 978-92-822-2204-1.
- [58] Xing Rong et al. “Development and evaluation of an improved quantitative 90Y bremsstrahlung SPECT method”. In: *Medical Physics* 39.5 (May 2012), pp. 2346–2358. ISSN: 2473-4209. DOI: 10.1118/1.3700174. URL: <https://onlinelibrary.wiley.com/doi/full/10.1118/1.3700174> (visited on 08/27/2025).
- [59] Xing Rong and Eric C. Frey. “A collimator optimization method for quantitative imaging: application to Y-90 bremsstrahlung SPECT”. eng. In: *Medical Physics* 40.8 (Aug. 2013), p. 082504. ISSN: 2473-4209. DOI: 10.1118/1.4813297.

- [60] Payvand Taherparvar and Nazila Shahmari. "Evaluation the effect of different collimators and energy window on Y-90 bremsstrahlung SPECT imaging by SIMIND Monte Carlo program". en. In: *Nuclear Medicine Review* 22.2 (2019), pp. 45–55. ISSN: 1644-4345. DOI: 10.5603/NMR.a2019.0016. URL: https://journals.viamedica.pl/nuclear_medicine_review/article/view/NMR.a2019.0016 (visited on 01/21/2026).
- [61] Minho Kim et al. "Effects of collimator on imaging performance of Yttrium-90 Bremsstrahlung photons: Monte Carlo simulation". In: *Nuclear Engineering and Technology* 51.2 (Apr. 2019), pp. 539–545. ISSN: 1738-5733. DOI: 10.1016/j.net.2018.11.007. URL: <https://www.sciencedirect.com/science/article/pii/S1738573318305308> (visited on 01/21/2026).
- [62] Caner İnce et al. "Collimator and energy window optimization for YTTRIUM-90 bremsstrahlung SPECT imaging". In: *Applied Radiation and Isotopes* 167 (Jan. 2021), p. 109453. ISSN: 0969-8043. DOI: 10.1016/j.apradiso.2020.109453. URL: <https://www.sciencedirect.com/science/article/pii/S0969804320305972> (visited on 01/21/2026).
- [63] Yuni K. Dewaraja et al. "Improved quantitative 90Y bremsstrahlung SPECT/CT reconstruction with Monte Carlo scatter modeling". In: *Medical Physics* 44.12 (Dec. 2017), pp. 6364–6376. ISSN: 2473-4209. DOI: 10.1002/MP.12597. URL: <https://onlinelibrary.wiley.com/doi/full/10.1002/mp.12597> (visited on 07/30/2025).
- [64] Michael Ljungberg. "- The SIMIND Monte Carlo Program". In: *Monte Carlo Calculations in Nuclear Medicine*. CRC Press, 2012. ISBN: 978-0-429-15116-3.
- [65] Se Young Chun et al. "Algorithms and Analyses for Joint Spectral Image Reconstruction in Y-90 Bremsstrahlung SPECT". In: *IEEE Transactions on Medical Imaging* PP (Oct. 2019), pp. 1–1. DOI: 10.1109/TMI.2019.2949068.

- [66] Lin Fu and Jinyi Qi. “A residual correction method for high-resolution PET reconstruction with application to on-the-fly Monte Carlo based model of positron range”. In: *Medical Physics* 37.2 (Feb. 2010), pp. 704–713. ISSN: 2473-4209. DOI: 10.1118/1.3284980. URL: <https://onlinelibrary.wiley.com/doi/full/10.1118/1.3284980> (visited on 09/11/2025).
- [67] Haowei Xiang et al. “A deep neural network for fast and accurate scatter estimation in quantitative SPECT/CT under challenging scatter conditions”. en. In: *European Journal of Nuclear Medicine and Molecular Imaging* 47.13 (2020), pp. 2956–2967. ISSN: 1619-7089. DOI: 10.1007/s00259-020-04840-9. URL: <https://doi.org/10.1007/s00259-020-04840-9>.
- [68] Mattijs Elschot et al. “Quantitative Comparison of PET and Bremsstrahlung SPECT for Imaging the In Vivo Yttrium-90 Microsphere Distribution after Liver Radioembolization”. en. In: *PLOS ONE* 8.2 (2013), e55742. ISSN: 1932-6203. DOI: 10.1371/journal.pone.0055742. URL: <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0055742>.
- [69] “XCOM: Photon Cross Sections Database”. en. In: *NIST* (Sept. 2009). Last Modified: 2024-08-14T11:59:04:00. URL: <https://www.nist.gov/pml/xcom-photon-cross-sections-database> (visited on 11/25/2025).
- [70] Alexander S. Pasciak et al. “Radioembolization and the Dynamic Role of 90Y PET/CT”. English. In: *Frontiers in Oncology* 4 (Feb. 2014). ISSN: 2234-943X. DOI: 10.3389/fonc.2014.00038. URL: <https://www.frontiersin.org/journals/oncology/articles/10.3389/fonc.2014.00038/full> (visited on 01/08/2026).
- [71] Alex Sawatzky et al. “EM-TV Methods for Inverse Problems with Poisson Noise”. en. In: *Level Set and PDE Based Reconstruction Methods in Imaging: Cetraro, Italy 2008, Editors: Martin Burger, Stanley Osher*. Ed. by Martin Burger et al. Cham: Springer International Publishing, 2013, pp. 71–142. ISBN: 978-3-319-01712-9. URL: https://doi.org/10.1007/978-3-319-01712-9_2.

- [72] Adrian S. Lewis. “The convex analysis of unitarily invariant matrix functions”. In: *Journal of Convex Analysis* 2.1 (1995), pp. 173–183.
- [73] A S Lewis. “Derivatives of Spectral Functions”. In: *Mathematics of Operations Research* 21.3 (1996), pp. 576–588. ISSN: 0364-765X. DOI: 10.1287/moor.21.3.576. URL: <https://pubsonline.informs.org/doi/10.1287/moor.21.3.576>.
- [74] Rie Johnson and Tong Zhang. “Accelerating Stochastic Gradient Descent using Predictive Variance Reduction”. In: *Advances in Neural Information Processing Systems*. Vol. 26. Curran Associates, Inc., 2013. (Visited on 01/12/2026).
- [75] N H Clinthorne et al. “Preconditioning methods for improved convergence rates in iterative reconstructions”. en. In: *IEEE Transactions on Medical Imaging* 12.1 (1993), pp. 78–83. ISSN: 02780062. DOI: 10.1109/42.222670. URL: <http://ieeexplore.ieee.org/document/222670/>.
- [76] A R De Pierro and M E B Yamagishi. “Fast EM-like methods for maximum ”a posteriori” estimates in emission tomography”. en. In: *IEEE Transactions on Medical Imaging* 20.4 (2001), pp. 280–288. ISSN: 02780062. DOI: 10.1109/42.921477. URL: <http://ieeexplore.ieee.org/document/921477/>.
- [77] H.M. Hudson and R.S. Larkin. “Accelerated image reconstruction using ordered subsets of projection data”. In: *IEEE Transactions on Medical Imaging* 13.4 (Dec. 1994), pp. 601–609. ISSN: 1558-254X. DOI: 10.1109/42.363108. URL: <https://ieeexplore.ieee.org/document/363108> (visited on 12/04/2025).
- [78] Hugo Levillain et al. “International recommendations for personalised selective internal radiation therapy of primary and metastatic liver diseases with yttrium-90 resin microspheres”. en. In: *European Journal of Nuclear Medicine and Molecular Imaging* 48.5 (2021), pp. 1570–1584. ISSN: 1619-7089. DOI: 10.1007/s00259-020-05163-5. URL: <https://doi.org/10.1007/s00259-020-05163-5>.

- [79] Paul Haste et al. “Correlation of Technetium-99m Macroaggregated Albumin and Yttrium-90 Glass Microsphere Biodistribution in Hepatocellular Carcinoma: A Retrospective Review of Pretreatment Single Photon Emission CT and Posttreatment Positron Emission Tomography/CT”. In: *Journal of Vascular and Interventional Radiology* 28.5 (May 2017), 722–730.e1. ISSN: 1051-0443. DOI: 10.1016/J.JVIR.2016.12.1221. (Visited on 08/21/2025).
- [80] D. Minarik, K. Sjögreen Gleisner, and M. Ljungberg. “Evaluation of quantitative 90Y SPECT based on experimental phantom studies”. en. In: *Physics in Medicine & Biology* 53.20 (Sept. 2008), p. 5689. ISSN: 0031-9155. DOI: 10.1088/0031-9155/53/20/008. URL: <https://dx.doi.org/10.1088/0031-9155/53/20/008> (visited on 01/27/2023).
- [81] Daniel Deidda et al. “Hybrid kernelised expectation maximisation for Bremsstrahlung SPECT reconstruction in SIRT with 90Y micro-spheres”. en. In: *EJNMMI Physics* 9.1 (Apr. 2022), p. 25. ISSN: 2197-7364. DOI: 10.1186/s40658-022-00452-4. URL: <https://doi.org/10.1186/s40658-022-00452-4> (visited on 12/04/2025).
- [82] Yu Jung Tsai et al. “Benefits of Using a Spatially-Variant Penalty Strength with Anatomical Priors in PET Reconstruction”. In: *IEEE Transactions on Medical Imaging* 39.1 (Jan. 2020), pp. 11–22. ISSN: 1558254X. DOI: 10.1109/TMI.2019.2913889. (Visited on 08/21/2025).
- [83] Evgueni Ovtchinnikov et al. “SIRF: Synergistic Image Reconstruction Framework”. In: *2017 IEEE Nuclear Science Symposium and Medical Imaging Conference (NSS/MIC)*. IEEE, 2017, pp. 1–3. ISBN: 978-1-5386-2282-7. DOI: 10.1109/NSSMIC.2017.8532815. URL: <https://ieeexplore.ieee.org/document/8532815/>.
- [84] Kris Thielemans et al. “STIR: Software for tomographic image reconstruction release 2”. In: *Physics in medicine and biology* 57 (2012), pp. 867–883. DOI: 10.1088/0031-9155/57/4/867. URL: <https://www.researchgate.>

net/publication/221791839_STIR_Software_for_tomographic_image_reconstruction_release_2.

- [85] Georg Schramm and Kris Thielemans. “PARALLELPROJ—an open-source framework for fast calculation of projections in tomography”. English. In: *Frontiers in Nuclear Medicine* 3 (Jan. 2024). ISSN: 2673-8880. DOI: 10 . 3389 / fnume . 2023 . 1324562. URL: <https://www.frontiersin.org/journals/nuclear-medicine/articles/10.3389/fnume.2023.1324562/full> (visited on 01/26/2026).
- [86] Berta Marti Fuster et al. *Integration of advanced 3D SPECT modeling into the open-source STIR framework - Fuster - 2013 - Medical Physics - Wiley Online Library*. May 2013. URL: <https://aapm.onlinelibrary.wiley.com/doi/10.1118/1.4816676> (visited on 01/26/2026).
- [87] Jakob Sauer Jørgensen et al. *Core Imaging Library - Part I: a versatile Python framework for tomographic imaging*. 2021. URL: <https://github.com/TomographicImaging/CIL/blob/f42d88ddae06e032853239e84d28de39ab06e41f/CITATION.cff>.
- [88] C.W. Stearns et al. “Random coincidence estimation from single event rates on the Discovery ST PET/CT scanner”. In: *2003 IEEE Nuclear Science Symposium. Conference Record*. Vol. 5. Oct. 2003, 3067–3069 Vol.5. DOI: 10 . 1109 / NSSMIC . 2003 . 1352545. URL: <https://ieeexplore.ieee.org/document/1352545> (visited on 12/09/2025).
- [89] M. Iatrou et al. “3D implementation of Scatter Estimation in 3D PET”. In: *2006 IEEE Nuclear Science Symposium Conference Record*. Vol. 4. Oct. 2006, pp. 2142–2145. DOI: 10 . 1109 / NSSMIC . 2006 . 354338. URL: <https://ieeexplore.ieee.org/document/4179452> (visited on 12/09/2025).
- [90] Hal O. Anger. “Sensitivity, resolution, and linearity of the scintillation camera”. In: *IEEE Transactions on Nuclear Science* NS-13.3 (1966), pp. 380–392. ISSN: 15581578. DOI: 10 . 1109 / TNS . 1966 . 4324123. (Visited on 08/07/2025).

- [91] Michael Ljungberg. “SIMIND Monte Carlo based image reconstruction”. en. In: *Journal of Nuclear Medicine* 56.supplement 3 (2015), p. 43. ISSN: 0161-5505, 2159-662X. URL: https://jnm.snmjournals.org/content/56/supplement_3/43.
- [92] P. J. Markiewicz, A. J. Reader, and J. C. Matthews. “Assessment of bootstrap resampling performance for PET data”. In: *Physics in Medicine & Biology* 60.1 (Dec. 2014), p. 279. ISSN: 0031-9155. DOI: 10.1088/0031-9155/60/1/279. URL: <https://iopscience.iop.org/article/10.1088/0031-9155/60/1/279> (visited on 09/12/2025).
- [93] Justin K. Mikell et al. “Impact of 90Y PET gradient-based tumor segmentation on voxel-level dosimetry in liver radioembolization”. en. In: *EJNMMI Physics* 5.1 (Nov. 2018), p. 31. ISSN: 2197-7364. DOI: 10.1186/s40658-018-0230-y. URL: <https://doi.org/10.1186/s40658-018-0230-y> (visited on 01/03/2026).
- [94] Vilfredo Pareto. “The New Theories of Economics”. en. In: *Journal of Political Economy* (Sept. 1897). ISSN: 0022-3808. DOI: 10.1086/250454. URL: <https://www.journals.uchicago.edu/doi/10.1086/250454> (visited on 01/11/2026).
- [95] Evgueni Ovtchinnikov et al. “SIRF: Synergistic Image Reconstruction Framework”. In: *Computer Physics Communications* 249 (Apr. 2020), p. 107087. ISSN: 0010-4655. DOI: 10.1016/j.cpc.2019.107087. URL: <https://www.sciencedirect.com/science/article/pii/S0010465519303984> (visited on 12/09/2025).